CÆNL: Collapse-Based Active Neural Learning with Objective-Aligned Representation Control

Saad Alqithami

alqithami@gmail.com

23 October 2026

Abstract

Active learning changes the representations used to acquire subsequent labels. We study this interaction through Collapse-Based Active Neural Learning (CÆNL), which separates acquisition, Dynamic Collapse Regularization (DCR), and its coefficient policy. DCR penalizes a normalized within-class to between-class scatter residual above a frozen reference, with a separate class-separation safeguard. Fixed, scheduled, deterministic-feedback, and learned-feedback policies apply this objective. The analysis establishes scatter invariances, conditional local descent, and coefficient bounds, and shows why clean geometry alone cannot guarantee input robustness. Earlier classification studies are followed by a 206-run evaluation with pretraining exact-content filtering, independent control selection, ten-seed tests on each of two backbones, and common-label component ablations. At 60% annotations on the filtered ImageNet-100 inventory, entropy acquisition with DCR and MACC-Lite reaches $73.393 \pm 0.707\%$ accuracy on ResNet-50 and $73.151 \pm 0.557\%$ on ResNet-18 (mean $\pm$ seed SD), exceeding entropy alone by 1.900 and 1.612 percentage points. Both gains survive the prespecified 17-comparison Holm correction under paired-t and exact sign-flip tests. With acquired labels held constant, the scatter-ratio penalty alone improves accuracy by 2.120 points. Common-label ablations support the scatter intervention, while comparisons with tuned fixed regularization leave the additional value of feedback unresolved. Completed diagnostics evaluate 30 frozen classifiers at five AutoAttack radii and establish no robust-accuracy advantage at the prespecified 2/255 endpoint. A separate 58-run raw-source AudioCaps study adds development-selected fixed control and ten paired confirmation seeds; none of its five adjusted CIDEr-D comparisons rejects. Earlier spectral studies yield a small fixed-regularization C4 loss reduction and worse CIFAR-10 diffusion FID. Short cached-data timings quantify feedback overhead without establishing whole-job efficiency. The results support the evaluated scatter-based classification intervention and delineate the task dependence of geometry-informed regularization.

Keywords: neural collapse; active learning; representation geometry; feedback control; regularization; discriminative learning

1 Introduction

Active learning seeks to make effective use of a limited annotation budget by selecting examples from an unlabeled pool [1]. In deep learning, this selection depends on a representation that is itself updated after new labels are acquired. Uncertainty, coverage, and gradient diversity provide complementary criteria for choosing examples [2, 3], while the training objective determines how those examples reshape the representation. These two decisions interact: changing the learned features can alter subsequent predictions and acquisitions, even when the acquisition rule remains fixed. This interaction motivates studying representation regularization as an explicit part of the active-learning process.

Neural collapse offers a geometric motivation for such regularization. During the terminal phase of classifier training, features from the same class become concentrated, and class means and classifier weights approach a structured configuration [4]. Relative within-class compactness and between-class separation therefore provide measurable properties of a developing representation. In an active-learning setting, however, the labeled sample and its class support change across rounds. A geometry measurement must be interpreted with respect to that changing sample, and a desired geometric structure must be connected to an objective that can influence training. The relevant question is whether using this information improves prediction at a given annotation budget.

Turning geometry into a feedback signal introduces two design questions. First, the observed statistic should have an explicit relationship to the loss term whose influence the controller changes. Otherwise, a coefficient update may modulate an objective that does not directly penalize the measured deviation. Second, the benefit of the representation loss must be distinguished from the benefit of adapting its coefficient. Adaptive loss weighting has established precedents, including gradient balancing in multitask learning [5]; here the feedback concerns a representation statistic within a task. A comparison with unregularized training assesses the complete treatment, whereas tuned constant and preset-schedule controls test whether observing geometry adds predictive value to coefficient selection. These considerations define our central question: *under a common acquisition and optimization protocol, how does geometry-driven regularization affect*

prediction, and how do fixed, scheduled, and feedback-controlled strengths compare?

We investigate this question through Collapse-Based Active Neural Learning (CÆNL), a framework with separately specified acquisition, representation-loss, and coefficient-control components. Dynamic Collapse Regularization (DCR) uses a normalized within-class to between-class scatter statistic for classification. A one-sided penalty acts on deviations above targets calibrated at initialization, and a separate safeguard discourages loss of between-class separation. The same statistic enters the loss and the feedback state, implementing *objective–feedback alignment*. The Meta-Adaptive Collapse Controller (MACC) provides deterministic bounded feedback in MACC-Lite and learned coefficient actions in Full MACC. Acquisition can use predicted-class geometric deviation or an established criterion such as entropy. This modular formulation allows the acquisition rule and the training treatment to be compared separately. Here, "collapse-based" refers to regulating relative class dispersion; the scatter statistic does not by itself characterize every property of neural collapse.

The contributions are as follows:

1. **Objective-aligned representation control.** We formulate the normalized scatter objective, its fixed targets and separation safeguard, and the deterministic and learned coefficient updates within a budgeted active-learning loop. The formulation makes the measured quantity, the modulated loss, and the acquired-label information available to each component explicit.

2. **Supporting geometric and control analysis.** We characterize scatter decomposition and invariance, the local effect of an isolated regularizer step, and bounds on coefficient updates. An explicit full-update condition and a constructive robustness counterexample identify the additional assumptions needed to connect feature geometry to task behavior.

3. **Controlled evaluation across methods and tasks.** We compare acquisition rules and regularization controls in six-seed ImageNet-100 studies and a separate 104-run control study. A subsequent 206-run study filters exact-content duplicates before training and adds common-label component replay, a matched contrastive-feature comparator, and a full-label reference with matched optimizer steps. A completed 30-checkpoint, 150-suite study evaluates smaller-radius robustness and practical training-path cost. Separate spectral studies cover language, diffusion, and audio, with an additional 58-run raw-source AudioCaps study using independently selected fixed strength and ten confirmation seeds.

The principal classification finding is a reproducible benefit of the regularized treatment over entropy acquisition alone within the evaluated ImageNet-100 protocols. In the content-filtered study, the final-budget gains are 1.900 percentage points on ResNet-50 and 1.612 points on ResNet-18, with unchanged feedback settings across backbones. Common-label replay identifies a benefit from the scatter-ratio penalty conditional on the shared acquired-label sequence. The incremental effects of adding the separation safeguard or feedback remain inconclusive, as do the end-to-end comparisons of Lite with independently tuned fixed DCR. Thus the experiments distinguish evidence for the representation intervention from evidence for adaptive superiority. The earlier acquisition, controller, and robustness studies remain complementary; their designs and statistical families are distinguished in Section 6.1.

The other tasks assess a separate spectral objective rather than the supervised class-scatter loss. Fixed spectral regularization produces a small C4 language-modeling loss reduction, all evaluated spectral conditions worsen CIFAR-10 diffusion FID relative to unregularized fine-tuning, and the AudioCaps comparisons establish no captioning benefit. Classification robustness is also evaluated separately, with near-zero accuracy under the specified AutoAttack threat model. These results make task performance the criterion for judging a geometric intervention and delineate the scope of the framework beyond its classification gains.

Section 2 positions the approach in the literature. Sections 3–5 develop the formulation, controllers, and supporting analysis. Sections 6 and 7 present the protocols and results, followed by interpretation and limitations in Section 8 and conclusions in Section 9.

2 Background and related work

CÆNL sits at the intersection of representation geometry, label acquisition, and adaptive optimization. These areas address different decisions: what structure features should acquire, which examples should be annotated, and how strongly an auxiliary objective should influence training. Reviewing these decisions separately helps identify the role of collapse-based feedback within an active-learning system.

2.1 Neural collapse and its theoretical scope

Papyan et al. [4] identified a coordinated terminal-training geometry: within-class variability contracts, centered class means approach a simplex equiangular tight frame (ETF), classifier weights align with those means, and classification approaches a nearest-centroid rule. These properties jointly define the phenomenon. A scalar scatter ratio can measure relative class compactness, but cannot by itself verify ETF angles, classifier alignment, or agreement with nearest-centroid decisions. It is also distinct from the pseudoinverse-normalized NC1 statistic. We therefore use collapse as motivation for a measurable regularization target, rather than equating one favorable statistic with the full phenomenon.

Theoretical work explains when such geometry is compatible with an optimization objective. Zhu

et al. [6] analyze regularized cross-entropy in an unconstrained-feature model, where features are optimization variables rather than outputs constrained by a complete network. Under their assumptions, global minimizers have simplex-ETF structure and other critical points are strict saddles. This establishes a precise link between an objective and its optimal geometry, while leaving the feature-learning dynamics of a general finite network outside that model.

Subsequent results show why architecture and regularization must be specified. Súkeník et al. [7] find regimes of the deep unconstrained-feature model in which lower-rank solutions are preferred to deep neural collapse. Their later work [8] establishes approximate collapse at global optima of particular regularized ResNet and transformer models with LayerNorm, with the approximation improving as depth increases. These results concern different models and assumptions. Hong and Ling [9] further analyze shallow ReLU networks with explicit data dependence: the conditions for collapse and its generalization consequences depend on architecture and signal-to-noise structure. Collapse can coexist with poor generalization in their low-signal regime.

Recent work also broadens the geometry beyond a balanced simplex. Wang et al. [10] study feature–classifier space misalignment in long-tailed learning and introduce explicit alignment strategies. Alcala et al. [11], in a 2026 preprint, characterize alternative configurations in the orthoplex regime, where feature dimension d and class count C satisfy $d + 2 \leq C \leq 2d$. Together, these studies motivate measuring class separation as well as contraction. They also distinguish geometric alignment of features and classifiers from the objective–feedback alignment investigated here: consistency between the statistic penalized by a loss and the statistic observed by its coefficient controller.

2.2 Discriminative and spectral representation objectives

The use of within-class and between-class moments predates neural collapse. Fisher's discriminant criterion [12] seeks directions that separate class means relative to within-class variation. DeepLDA [13] brings a discriminant objective into learned representations. Center loss [14] instead penalizes feature distances to learned class centers alongside softmax training. Supervised contrastive learning [15] uses class labels to bring same-class embeddings together while separating different-class embeddings. These approaches establish compactness and separation as useful, longstanding representation objectives; the use of class moments alone is not the contribution of CÆNL.

The distinctions concern the quantity optimized and how it enters training. Center loss controls distances to learned prototypes, DeepLDA uses a discriminant criterion, and supervised contrastive learning operates through relations among embeddings. Our classification objective uses a normalized log ratio of equal-class-weighted scatter, penalizes only deviations above a frozen reference target, and includes a separate safeguard against loss of between-class separation. Its controller observes that same ratio. This construction permits direct comparisons between a constant coefficient, a predetermined schedule, and feedback without changing the underlying auxiliary objective. The present experiments evaluate those controls, rather than ranking all discriminative representation losses.

Class-conditional contraction should also be distinguished from degenerate representation collapse. In self-supervised learning, VICReg [16] preserves per-dimension variability and reduces covariance redundancy while encouraging agreement between views. A useful representation can therefore require preserving variation, not simply suppressing it. This distinction is particularly relevant to our spectral extensions: reducing spectral complexity without class labels need not preserve the distinctions required for language prediction, generation, or captioning. Those extensions use a separate objective and require their own task evaluations.

2.3 Active learning through uncertainty and coverage

Pool-based active learning allocates annotations using the current model and the remaining unlabeled data [1]. Predictive entropy is an accessible uncertainty score, but uncertainty about an output and uncertainty about model parameters need not be the same. Gal et al. [17] integrate approximate Bayesian neural-network uncertainty into image acquisition, including Bayesian active learning by disagreement (BALD), which selects points for their information about model parameters. Such criteria motivate careful interpretation of any geometric score: distance from a class center is not automatically an estimate of epistemic uncertainty.

Batch selection introduces a complementary coverage problem because individually uncertain points may be redundant. Core-set selection [2] addresses coverage in embedding space. BADGE [3] combines uncertainty and diversity by selecting from hallucinated last-layer gradient embeddings. Noise Stability [18] estimates uncertainty from output changes under parameter perturbations and relates the resulting selection to large, diverse gradients. These mechanisms use different information and incur different acquisition costs; an equal annotation count does not make their computational workloads identical.

Uncertainty Herding [19] formulates uncertainty coverage to connect the low- and high-label-budget regimes. Its greedy selection balances uncertainty with coverage in a fixed feature space, addressing the budget dependence of purely uncertainty- or representation-driven selection. Our numerical comparison implements its margin-weighted coverage mechanism using shared-initial task-model features under the common candidate-pass and warm-start conventions specified in Section 6.7.

Recent acquisition methods make this balance adaptive. Jung et al. [20] study combinations of acquisition

functions and their accuracy–computation trade-offs, including feedback-based aggregation. Cui et al. [21] use labeled-data feedback to construct a TrustSet and learn a batch acquisition policy. In CÆNL, the optional Mahalanobis score instead measures deviation from the predicted class using acquired-label statistics. It is a roundwise selection criterion. Our principal control study holds entropy acquisition fixed so that its comparisons address the training treatment; the acquisition alternatives are examined separately.

2.4 Adaptive weighting and representation feedback

Adapting the influence of loss terms is an established optimization idea. Kendall et al. [22] learn task weights through a likelihood formulation based on homoscedastic uncertainty. GradNorm [5] adjusts weights using gradient magnitudes and relative training rates to balance multitask learning. Both demonstrate that the choice of feedback signal is central to an adaptive weighting rule. Their task-balancing objectives differ from controlling a representation statistic within a single classification task, but they provide relevant precedents for learning or updating loss weights during training.

MACC-Lite uses a bounded deterministic update driven by deviations of the monitored scatter statistic from its frozen target. Full MACC learns coefficient actions through a policy-gradient mechanism [23]; its acquisition-policy head is disabled in the evaluated conditions. Thus the feedback acts on regularization strengths, whereas aggregation and TrustSet act on the selection of examples. Holding this distinction explicit avoids attributing a training improvement to a better query policy.

This literature motivates a specific empirical question: once a representation objective is fixed, does observing its state improve the choice of its coefficient? A constant, a predetermined schedule, and feedback can produce different training trajectories even when they share the same loss. Comparing only against unregularized training cannot distinguish the usefulness of the loss from the usefulness of adaptation. Our separately tuned controls and subsequent confirmation therefore test that distinction directly. Table 1 summarizes the relevant methodological relationships and identifies the comparisons actually executed.

2.5 Geometry, robustness, and predictive evaluation

Representation geometry, predictive confidence, and resistance to input perturbations describe different properties. Su et al. [24] study neural collapse in both standard and adversarially trained networks, showing that small adversarial perturbations can disrupt the clean simplex structure of nonrobust models. Their analysis also identifies aligned clean and perturbed geometry in robust models. These findings make robustness a separate empirical question, rather than a consequence of clean-feature compactness. AutoAttack [25] provides an ensemble evaluation designed to reduce dependence on a single attack; its interpretation still requires the input norm, radius, and preprocessing to be stated.

Calibration likewise concerns the relationship between reported confidence and correctness [26]. An auxiliary representation loss can change both accuracy and confidence, so we report calibration separately from the scatter statistic and clean accuracy. This evaluation perspective connects the literature to our experimental design: geometry motivates the intervention, while matched task outcomes assess its usefulness. The paper's focus is the specified coupling of representation loss and coefficient feedback, with acquisition, robustness, and cross-task effects evaluated as distinct questions.

3 Problem formulation and representation measurements

The classification formulation distinguishes three roles for representation geometry: a differentiable training statistic, an aggregate feedback signal, and a sample-level acquisition score. The first two use the same normalized scatter measurement; acquisition uses a separate raw-feature distance. The spectral formulation for the other tasks is given in Section 4.5.

3.1 Sets, budgets, and information access

Let $\mathcal{D} = (x_i)_{i\in\mathcal{I}}$ be a fixed training pool with index set $\mathcal{I} = \{1, \dots, N\}$ and class labels $y_i \in \{1, \dots, C\}$. All partitions below are sets of indices. After acquisition round r, $\mathcal{A}_r \subseteq \mathcal{I}$ contains the annotated indices and $\mathcal{U}_r = \mathcal{I} \setminus \mathcal{A}_r$ the remaining pool. A fixed feedback set $\mathcal{H} \subset \mathcal{A}_0$ is included in the annotation budget and excluded from model-gradient updates and acquisition statistics. Thus, $\mathcal{T}_r = \mathcal{A}_r \setminus \mathcal{H}$ supplies gradient-training examples. Feedback labels are available to the control and tuning procedures specified in Sections 4 and 6.4.

For total annotation budget B, round $r+1$ selects distinct indices $Q_{r+1} \subseteq \mathcal{U}_r$ with $|Q_{r+1}| = b_{r+1} \le \min\{|\mathcal{U}_r|, B - |\mathcal{A}_r|\}$. After their labels are obtained,

$$\mathcal{A}_{r+1} = \mathcal{A}_r \cup Q_{r+1}, \qquad \mathcal{U}_{r+1} = \mathcal{U}_r \setminus Q_{r+1}.$$

Consequently, $|\mathcal{A}_r| = |\mathcal{A}_0| + \sum_{j=1}^{r} b_j \le B$, including feedback labels. The initial stratified annotation protocol is specified in Section 6.2; subsequent query decisions do not use the hidden labels of $\mathcal{U}_r$.

The model produces logits $f_\theta(x) \in \mathbb{R}^C$ and pooled features $z_l(x) \in \mathbb{R}^{d_l}$ at monitored layers $l \in \mathcal{L}$. Training, feedback, acquisition, and evaluation transforms are specified separately. Official evaluation labels do not enter the acquisition or online controller. Partition disjointness in the earlier studies is by index; content duplicates and their sensitivity analysis are reported in Section 7.5. The later inventory is filtered for exact-content identities before training (Section 6.5). Performance is compared at fixed annotation budgets.

Table 1: Positioning relative to representative acquisition, representation, and loss-weighting methods. Numerical entries are ImageNet-100 clean accuracy at 60% annotations in the earlier six-seed studies (mean ± seed SD), not numbers pooled from different published benchmarks. A dash denotes no estimate in that earlier matrix; the later matched one-view contrastive variant is reported in Section 7.1.

Method	Principal strength or focus	Distinction or limitation for this study	Matched top-1 (%)
Center loss / DeepLDA [13, 14]	Direct discriminative representation objectives	No matched coefficient-control experiment here	—
Supervised contrastive learning [15]	Label-based attraction and separation of embeddings	One-view feature variant tested later; original SupCon recipe not reproduced	—
VICReg [16]	Preserves variance while reducing redundancy	Self-supervised objective; distinct from class-conditional contraction	—
GradNorm / uncertainty weighting [5, 22]	Adaptive balancing of multiple task losses	Different feedback signals and optimization goals; not executed here	—
BADGE [3]	Uncertainty and gradient diversity	Acquisition criterion; does not modulate the scatter objective	71.110 ± 0.467
Noise Stability [18]	Perturbation-based acquisition signal	Higher measured acquisition cost in this implementation	71.023 ± 0.249
Uncertainty Herding [19]	Uncertainty-weighted feature coverage	Local fixed-initial-R50 adaptation; published benchmark recipe not reproduced	71.747 ± 0.626
Aggregation / TrustSet [20, 21]	Adaptive acquisition using feedback	Acquisition actions differ from the regularization-only Full MACC studied here	—
Aligned DCR + Lite (entropy)	Shared statistic in loss and coefficient feedback	Specified-control comparison; later tuned-control results in Section 7.3	73.360 ± 0.208

Using 60% of the pool does not alone establish a 40% annotation saving at equal accuracy.

3.2 Class-balanced normalized scatter

Consider a labeled measurement set $\mathcal{S}$ with finite features. Let $\mathcal{C}_{\mathcal{S}} = \{y_i : i \in \mathcal{S}\}$ be its represented classes, $K = |\mathcal{C}_{\mathcal{S}}|$, and $\mathcal{S}_c = \{i \in \mathcal{S} : y_i = c\}$ with $n_c = |\mathcal{S}_c|$. Thus K may be smaller than the total class count C. For each monitored layer, define

$$h_{i,l} = \frac{z_l(x_i)}{\max\{\|z_l(x_i)\|_2, \epsilon_{\mathrm{n}}\}}, \qquad \epsilon_{\mathrm{n}} = 10^{-6}. \tag{1}$$

Let $\mu_{c,l} = n_c^{-1} \sum_{i \in \mathcal{S}_c} h_{i,l}$ and $\bar{\mu}_l = K^{-1} \sum_{c \in \mathcal{C}_{\mathcal{S}}} \mu_{c,l}$. The equal-class-weighted population moments and stabilized log ratio are

$$W_l = \frac{1}{K} \sum_{c \in \mathcal{C}_{\mathcal{S}}} \frac{1}{n_c} \sum_{i \in \mathcal{S}_c} \|h_{i,l} - \mu_{c,l}\|_2^2, \tag{2}$$

$$B_l = \frac{1}{K} \sum_{c \in \mathcal{C}_{\mathcal{S}}} \|\mu_{c,l} - \bar{\mu}_l\|_2^2, \tag{3}$$

$$q_l = \log(W_l + \epsilon_{\mathrm{s}}) - \log(B_l + \epsilon_{\mathrm{s}}). \tag{4}$$

Here $\epsilon_{\mathrm{s}} = 10^{-6}$ and $\kappa_l = \exp(q_l) = (W_l + \epsilon_{\mathrm{s}})/(B_l + \epsilon_{\mathrm{s}})$. The denominator n_c defines a population moment, rather than an unbiased sample-variance estimate. Each class receives weight $1/K$ regardless of its example count.

The implementation requires $K \geq 2$ and $n_c \geq 2$ for every represented class. Training further requires balanced batches: $K = 16$ classes with eight distinct examples per class. Assessment permits unequal counts and computes the moments once over the complete assessment set, using all represented classes. In contrast, feedback uses moving averages of training-batch q_l and B_l. A mean of batch log ratios is generally different from the log ratio computed over their pooled examples; class support and input transforms also differ across these measurements.

A smaller q_l indicates lower stabilized within-class dispersion relative to between-class dispersion. It does not determine every pairwise class distance, the full covariance spectrum, or the classifier boundary. In particular, identical features give $W_l = B_l = 0$ and $q_l = 0$; numerical finiteness does not certify useful class separation. Section 4 therefore specifies a separate separation penalty and its limitations. Normalization gives $\|h_{i,l}\|_2 \leq 1$, with equality above the norm floor. The moments are invariant to a common orthogonal transformation and, when the floor is inactive before and after scaling, to common positive feature rescaling (Proposition 1).

3.3 Sample-level geometry for acquisition

At an acquisition checkpoint, use *raw* pooled features from $\mathcal{T}_r$, with the same deterministic transform used for pool scoring. For each global class $c \in \{1, \ldots, C\}$, let $\mathcal{T}_{r,c} = \{i \in \mathcal{T}_r : y_i = c\}$ and $m_{r,c} = |\mathcal{T}_{r,c}|$. The reference statistics are

$$\widehat{\mu}_{c,l} = \frac{1}{m_{r,c}} \sum_{i \in \mathcal{T}_{r,c}} z_l(x_i),$$

$$\widehat{v}_{c,l,j} = \frac{1}{m_{r,c}} \sum_{i \in \mathcal{T}_{r,c}} (z_{l,j}(x_i) - \widehat{\mu}_{c,l,j})^2.$$

Round and checkpoint dependence are suppressed in this notation. The implementation requires $m_{r,c} \geq 2$ for every global class; there is no fallback estimate for an unsupported class. With predicted class $\widehat{c}(x) = \arg\max_c f_\theta(x)_c$, define

$$u_l(x) = \sum_{j=1}^{d_l} \frac{(z_{l,j}(x) - \widehat{\mu}_{\widehat{c}(x),l,j})^2}{\widehat{v}_{\widehat{c}(x),l,j} + \epsilon_a}, \tag{5}$$

$$s(x) = \sum_{l \in \mathcal{L}} \alpha_l u_l(x), \qquad \alpha_l \geq 0, \quad \sum_l \alpha_l = 1. \tag{6}$$

Here $\epsilon_a = 10^{-5}$ is added once to each variance, and layer3 and layer4 have weights 0.5 each. Thus u_l is a squared diagonal Mahalanobis distance; it uses neither a square root nor a full covariance inverse. Acquired true labels define the reference classes, while predictions assign unlabeled candidates to a reference class.

The checkpoint and reference statistics remain fixed during a round's query selection and are recomputed after subsequent training. The largest scores are selected within repeated candidate passes, with ties resolved by candidate order, until the round quota is filled (Section 4.4). This procedure does not select the global top-scoring subset of the entire pool. A high score indicates deviation from the predicted-class reference; confident outliers can also score highly. The raw-feature diagonal score does not inherit the general orthogonal or exact scale invariance of the normalized scatter statistic. Its usefulness for acquisition is therefore evaluated separately from the training and feedback effects.

4 Objective-aligned DCR and regularization control

Building on Section 3, the classification method uses the normalized scatter statistic to define both a representation loss and the feedback that adjusts its strength. The construction has two time scales: coefficients act within training phases, while label acquisition occurs between phases. We first specify the shared objective and frozen reference targets, then the deterministic and learned controllers, and finally their integration into active learning. The spectral task variants use the separate formulation in Section 4.5.

4.1 Frozen targets and classification objective

Each seed begins with a shared cross-entropy-only checkpoint trained on $\mathcal{T}_0$. Before the first acquisition, 16 fixed class-balanced anchor batches from $\mathcal{T}_0$ are evaluated using deterministic center crops. Let $\overline{q}_{l,0}$ and $\overline{B}_{l,0}$ be their mean log ratio and mean between-class dispersion. The targets are

$$q_l^\star = \overline{q}_{l,0} + \log(0.9), \qquad B_l^{\min} = 0.5\overline{B}_{l,0}. \tag{7}$$

Thus $\exp(q_l^\star)$ is 90% of the geometric mean of the anchor ratios. Calibration requires $\overline{B}_{l,0} > 10\epsilon_{\mathrm{s}}$, giving $B_l^{\min} > 5\epsilon_{\mathrm{s}} > 0$. These seed-specific targets are shared across methods and remain frozen as the annotated set grows. They define a reference intervention, rather than an optimal geometry or a threshold for the differently aggregated 100-class assessment statistic.

For a current gradient-training batch, $\mathcal{L}_{\mathrm{CE}}$ is its mean cross-entropy and q_l, B_l are the instantaneous moments defined in Section 3. Define the one-sided Huber penalty

$$H_\delta(u) = \begin{cases} 0, & u \leq 0, \\ u^2/(2\delta), & 0 < u < \delta, \\ u - \delta/2, & u \geq \delta. \end{cases} \tag{8}$$

The aligned Dynamic Collapse Regularization (DCR) objective is

$$R_l = H_\delta(q_l - q_l^\star), \tag{9}$$

$$F_l = H_\delta\big(\log B_l^{\min} - \log(B_l + \epsilon_{\mathrm{s}})\big), \tag{10}$$

$$\mathcal{L}_{\mathrm{AD}} = \mathcal{L}_{\mathrm{CE}} + \sum_{l \in \mathcal{L}} \lambda_l^{\mathrm{eff}} R_l + \beta_B \sum_{l \in \mathcal{L}} F_l. \tag{11}$$

We use $\delta = 0.25$, $\epsilon_{\mathrm{s}} = \epsilon_{\mathrm{n}} = 10^{-6}$, and $\beta_B = 0.01$. The ratio penalty vanishes when $q_l \leq q_l^\star$; it therefore penalizes excess dispersion without forcing equality to the target. The separation penalty vanishes when $B_l + \epsilon_{\mathrm{s}} \geq B_l^{\min}$. It discourages the degenerate geometry identified in Section 3, but does not exclude every degenerate stationary point or separate every class pair.

The nonnegative applied coefficients λ_l^{eff} distinguish fixed, scheduled, and feedback-controlled treatments. All regularized classification treatments retain the same separation penalty, including when Full MACC disables a ratio term. The unregularized reference omits both penalties. Coefficients and targets are held constant during each model-gradient update; gradients pass through the current batch moments. The dense class-incidence calculation costs $O(nKd_l)$ per layer for a batch of n examples, with $K = 16$, without forming a full covariance matrix.

For feedback, moving moments are initialized at the calibration means. After each optimizer step k, the observations from that step's forward pass update

$$\widetilde{m}_l^{(k)} = 0.95\widetilde{m}_l^{(k-1)} + 0.05 m_l^{(k)}, \qquad m \in \{q, B\}.$$

At a control boundary t, denote the latest values by $\widehat{q}_{l,t}$ and $\widehat{B}_{l,t}$. This aligns the loss and feedback at the level of the measured statistic, while retaining the distinction between instantaneous and smoothed measurements made in Section 3. Cross-entropy, augmentation, momentum, and interactions among layers still affect the resulting trajectory.

4.2 Deterministic coefficient control: MACC-Lite

The Meta-Adaptive Collapse Controller (MACC) adjusts regularization strength from observed training

state. MACC-Lite uses a deterministic rule. With residual $e_{l,t} = \widehat{q}_{l,t} - q_l^\star$ and $[e]_+ = \max\{e, 0\}$, its bounded leaky update is

$$u_{l,t} = \lambda_{\max} \tanh\left(\frac{[e_{l,t}]_+}{\delta}\right), \tag{12}$$

$$\lambda_{l,t+1} = (1-\rho)\lambda_{l,t} + \rho u_{l,t}. \tag{13}$$

Here $\lambda_{\max} = 0.10$, $\rho = 0.20$, and the coefficient state is initialized at $\lambda_{l,0} = 0.03$ before the first phase-start action. Lite applies $\lambda_l^{\text{eff}} = \lambda_l$ without activation flags. Positive residuals set a positive desired strength; when the residual is nonpositive, the coefficient decays toward zero. The convex update preserves its stated bounds (Proposition 4), without establishing closed-loop stability or target convergence.

Holdout cross-entropy is recorded, but does not enter the Lite action. The original fixed-strength comparison uses $\lambda_l = 0.03$. The later study independently selects a constant and a predetermined schedule while keeping the Lite rule unchanged (Section 6.4). These controls distinguish the effect of the representation objective from the additional value of responding to its observed state.

4.3 Learned coefficient control: Full MACC

Full MACC learns which coefficient increments and activation flags to apply. Its policy has two 128-unit hidden layers with GELU activations. At boundary t, the state for the next action contains each layer's smoothed log ratio, target deviation, temporal-change feature, current coefficient, and activation flag, followed by training progress and feedback cross-entropy on $\mathcal{H}$. The two-layer classification configuration therefore has 12 state entries. Progress is the fraction of planned control windows completed across the acquisition phases.

With $\widehat{\kappa}_{l,t} = \exp(\widehat{q}_{l,t})$ and $q_{l,t}^{\text{prev}}$ the reference from the preceding completed window, the temporal feature is

$$d_{l,t} = \frac{|\widehat{q}_{l,t} - q_{l,t}^{\text{prev}}|\,\widehat{\kappa}_{l,t}}{\max\{\widehat{\kappa}_{l,t}, 10^{-12}\}}. \tag{14}$$

Away from the numerical floor, this is the absolute log-ratio change, without additional standardization. The reference starts at the calibration mean and is refreshed after each completed window.

Each action chooses a layer-wise increment from $\{-0.01, -0.005, 0, 0.005, 0.01\}$ and a binary flag g_l. The stored coefficient is incremented and projected to $[0, 0.10]$; the applied coefficient is $\lambda_l^{\text{eff}} = g_l \lambda_l$. Stored coefficients begin at 0.03 and flags at one, before the first action. The selected coefficients remain fixed during the following window of at most 100 optimizer steps.

Index completed windows by t. Let ℓ_t^H be feedback cross-entropy at the end of window t, and $s_H = \max\{0.1, \ell_{\text{phase},0}^H\}$ the scale fixed at the start of its acquisition phase. Define the geometric potential from the smoothed moments as

$$\Psi_t = \sum_{l\in\mathcal{L}} H_\delta(\widehat{q}_{l,t} - q_l^\star) + \beta_B \sum_{l\in\mathcal{L}} H_\delta\Big(\log B_l^{\min} - \log(\widehat{B}_{l,t} + \epsilon_{\text{s}})\Big).$$

For the action applied during that window, the reward is

$$r_t = \text{clip}\left(\frac{\ell_{t-1}^H - \ell_t^H}{s_H}, -1, 1\right) + 0.1\,\text{clip}(\Psi_{t-1} - \Psi_t, -1, 1) - 0.01\frac{1}{|\mathcal{L}|}\sum_{l\in\mathcal{L}}\frac{\lambda_{l,t}^{\text{eff}}}{0.10}. \tag{15}$$

The three terms reward feedback-loss reduction and geometric progress, and penalize applied strength. The ratio terms in Ψ_t are not multiplied by the controlled coefficients, so switching a ratio term off cannot alone produce a geometric reward. At the first window of each phase, the preceding loss and potential are that phase's starting values. Official evaluation labels do not enter this reward.

Uniform exploration decreases linearly from probability 0.10 toward 0.02 over the planned actions. Otherwise, actions are sampled from the policy's categorical distributions. Only policy-sampled actions receive REINFORCE updates [23]; uniform exploratory actions still contribute to the moving reward baseline. The advantage uses the previous baseline, which is updated afterward with observation weight 0.05. The first reward initializes this baseline and has zero centered advantage, although the entropy term can still update the policy. Policy optimization uses Adam with learning rate 3×10^{-4} and entropy coefficient 0.001.

Algorithm 1 specifies the order of reward attribution, policy updating, and action selection. The terminal action receives its reward before the phase ends. The acquisition-policy head is disabled: query counts and candidate-pass settings remain those in Section 4.4. Full MACC consequently learns regularization control online within each run; it is not meta-trained across tasks.

4.4 Round-based active learning

Figure 1 connects the data partitions and measurements of Section 3 to the loss and controllers above. Algorithm 2 gives the classification execution order. Each method starts from the same seed-specific model and optimizer checkpoint. A round first acquires labels using that checkpoint, then trains on the enlarged $\mathcal{T}_r$.

During acquisition, the checkpoint is fixed and required pool predictions and features are cached. Candidate passes contain at most 10,000 examples, sampled without replacement within each pass. For a pass of size n, the quota is $\min\{b_{\text{rem}}, \lceil 0.10n \rceil\}$, where b_{rem} is the remaining round quota. Selected indices are removed; unselected candidates may recur in later passes. The Mahalanobis conditions rank Equation (6) within

Algorithm 1 Full-MACC update at the end of a control window

Input: Controller state C with policy π_ϕ and pending pair (S_t, A_t); model f_θ after at most 100 SGD steps under A_t; fixed holdout $\mathcal{H}$; reward baseline b_t; terminal flag d.
Output: Updated controller state C, baseline b_{t+1}, and pending pair $P \in \{(S_{t+1}, A_{t+1}), \varnothing\}$.

Evaluate the completed window
$\ell_t^H \leftarrow \textsc{CrossEntropy}(f_\theta, \mathcal{H})$; $\Psi_t \leftarrow \textsc{GeometricPotential}(C)$
$r_t \leftarrow \textsc{Reward}(\ell_{t-1}^H, \ell_t^H, \Psi_{t-1}, \Psi_t, s_H, \lambda_t^{\text{eff}})$ ▷ *Equation* (15)*; phase-fixed* s_H
Credit the action before selecting its successor
if b_t is uninitialized **then**
  $b_t \leftarrow r_t$ ▷ *The first centered advantage is zero*
end if
$a_t \leftarrow r_t - b_t$
$b_{t+1} \leftarrow 0.95 b_t + 0.05 r_t$
if A_t was sampled from π_ϕ **then**
  $\phi \leftarrow \textsc{ReinforceUpdate}(\phi, S_t, A_t, a_t, \eta_{\text{ent}} = 0.001)$ ▷ *Regularization heads and entropy term only*
end if
$P \leftarrow \varnothing$ ▷ *Uniform exploratory actions receive no policy-gradient update*
if $\neg d$ **then**
  $S_{t+1} \leftarrow \textsc{BuildState}(C, \ell_t^H)$
  $A_{t+1} \leftarrow \textsc{SampleAction}(\pi_\phi, S_{t+1})$ ▷ *Policy sampling or scheduled uniform exploration*
  $C \leftarrow \textsc{ApplyBoundedAction}(C, A_{t+1})$ ▷ *Project strengths to* $[0, 0.10]$ *and apply layer flags*
  $P \leftarrow (S_{t+1}, A_{t+1})$
end if
$C \leftarrow \textsc{RefreshReferences}(C, \ell_t^H, \Psi_t, \widehat{q}_t)$
return (C, b_{t+1}, P) ▷ *Terminal reward is retained; no unused next action is sampled*

each pass, breaking ties by candidate order. Random, entropy, BADGE, and Noise Stability share the annotation and candidate-pass budgets, with method-specific selection rules and computational costs.

After the labels are acquired, model parameters and optimizer momentum are retained. Controller state and moving geometry also carry across phases, while the learning-rate schedule restarts. Each phase measures its starting holdout loss and geometric potential and selects an initial action before gradient training begins. Subsequent control boundaries occur every 100 steps and at the final step. Lite selects a new strength at nonterminal boundaries; Full MACC credits the completed action and then selects its successor. At the terminal boundary, Full MACC credits the last action without sampling an unused successor. Reinitializing the next phase's loss and potential references prevents changes caused by acquisition and the learning-rate restart from being attributed to the preceding phase's final action.

4.5 Spectral formulation for language, diffusion, and audio

The other tasks retain the idea of controlling a representation penalty, but use the covariance-spectrum formulation discussed in Section 2.2. They do not use the supervised class partition or the scatter target of the classification method. Their objective is

$$\mathcal{L}_{\text{spec}} = \mathcal{L}_{\text{task}} + \sum_{l \in \mathcal{L}} \lambda_l \|p_l - p_l^\star\|_2^2, \tag{16}$$

where λ_l denotes the applied strength, including any activation flag. The penalty matches eigenvalue vectors rather than full covariance matrices.

For a centered feature matrix $Z_l \in \mathbb{R}^{n_l \times d_l}$, the implementation computes descending, nonnegative eigenvalues $\sigma_{l,j}$ of $Z_l^\top Z_l / n_l$, using the smaller Gram matrix $Z_l Z_l^\top / n_l$ when appropriate. With $k_l = \min(k, d_l)$, retain the first k_l eigenvalues and pad with zeros if fewer are available. The batch spectrum is normalized over these retained entries:

$$p_{l,j} = \frac{\sigma_{l,j}}{\sum_{a=1}^{k_l} \sigma_{l,a} + \epsilon_{\text{spec}}}, \quad j = 1, \ldots, k_l. \tag{17}$$

With $\epsilon_{\text{spec}} = 10^{-8}$, the total mass can be appreciably smaller than one when the retained spectral sum is very small. A rank-r flat target assigns mass $1/r$ to its first r entries and zero to the rest. The target rank is restricted by the feature dimension, retained spectrum size, and, when necessary, the centered sample-rank bound $n_l - 1$. Matching these eigenvalues leaves eigenvectors and class arrangements unspecified.

The spectral monitor uses the full moving covariance spectrum, whereas the loss is computed by eigenspectrum autodifferentiation from the current truncated batch spectrum. Smoothness is local to regions where the relevant eigenvalues remain separated. Ties crossing the retained-spectrum or target-rank boundary can make the penalty nonsmooth; autodifferentiation does not establish global smoothness. For an exactly normalized nonzero spectrum π, effective rank is $\exp(-\sum_j \pi_j \log \pi_j)$. Neither this spectral measurement nor the loss is the supervised trace ratio of Equation (4). The spectral controller rules and settings are task-specific; Table 3 and Section 6.9 distinguish them from the aligned classification controllers above.

Language modeling uses token hidden states with next-token cross-entropy. Diffusion uses spatially pooled U-Net mid-block activations with noise-prediction mean squared error. Audio captioning uses learned prefix embeddings and pooled audio-token-encoder features with next-token cross-entropy. These

Algorithm 2 CÆNL active learning in the original classification study

Input: Indexed pool $\mathcal{D}$ with indices $\mathcal{I}$; initial annotations $\mathcal{A}_0$; fixed feedback set $\mathcal{H} \subset \mathcal{A}_0$; official evaluation set $\mathcal{V}$; methods $\mathcal{M}$; layers $\mathcal{L}$; per-round budget b; $R = 5$ acquisition rounds.

Output: Scheduled checkpoints and clean/robust evaluation results for every method, using the same seed-specific initialization.

Shared initialization and calibration

$\mathcal{T}_0 \leftarrow \mathcal{A}_0 \setminus \mathcal{H}$; $\mathcal{U}_0 \leftarrow \mathcal{I} \setminus \mathcal{A}_0$

$(\theta_0, \Omega_0) \leftarrow \textsc{TrainCrossEntropy}(\mathcal{T}_0)$ ▷ *Ω denotes optimizer state, including momentum*

$\Gamma \leftarrow \textsc{CalibrateTargets}(f_{\theta_0}, \mathcal{T}_0, \mathcal{L})$ ▷ *Freeze all $q_l^\star$ and $B_l^{\min}$*

for each method $m \in \mathcal{M}$ **do**

  $(\theta, \Omega) \leftarrow \textsc{Restore}(\theta_0, \Omega_0)$; $C \leftarrow \textsc{InitializeControl}(m, \Gamma)$

  $\mathcal{A} \leftarrow \mathcal{A}_0$; $\mathcal{U} \leftarrow \mathcal{U}_0$

  for $r = 1, \ldots, R$ **do**

    Acquire labels using the current model

    $F \leftarrow \textsc{InferPool}(f_\theta, \mathcal{U})$; $G \leftarrow \textsc{FitRequiredStatistics}(f_\theta, \mathcal{A} \setminus \mathcal{H}, m)$

    $Q \leftarrow \textsc{CappedAcquisition}(m, F, G, b)$ ▷ *$|Q| = b$ distinct examples; hidden labels are inaccessible*

    $\textsc{QueryLabels}(Q)$; $\mathcal{A} \leftarrow \mathcal{A} \cup Q$; $\mathcal{U} \leftarrow \mathcal{U} \setminus Q$

    $\mathcal{T}_r \leftarrow \mathcal{A} \setminus \mathcal{H}$; $K_r \leftarrow \textsc{ScheduledSteps}(\mathcal{T}_r, 20 \text{ epochs})$

    $C \leftarrow \textsc{BeginPhase}(C, f_\theta, \mathcal{H})$ ▷ *Set phase references and first action; retain states*

    Train and update the designated controller

    for $k = 1, \ldots, K_r$ **do**

      $B_k \leftarrow \textsc{ClassBalancedBatch}(\mathcal{T}_r)$

      $(\theta, \Omega, \xi_k) \leftarrow \textsc{TrainingStep}(\theta, \Omega, B_k, m, C, \Gamma)$

      $C \leftarrow \textsc{UpdateMoments}(C, \xi_k)$

      if $k \bmod 100 = 0$ **or** $k = K_r$ **then**

        $C \leftarrow \textsc{UpdateControl}(C, f_\theta, \mathcal{H}, d = (k = K_r))$

      end if

    end for

    $\mathcal{K}_{m,r} \leftarrow \textsc{SaveCheckpoint}(\theta, \Omega, C)$; $E_{m,r} \leftarrow \textsc{EvaluateClean}(f_\theta, \mathcal{V})$

  end for

  $E_m^{\text{robust}} \leftarrow \textsc{AutoAttack}(\mathcal{K}_{m,R}, \mathcal{V})$ ▷ *Frozen full-set protocol; final scheduled checkpoint*

end for

return $\{(\mathcal{K}_{m,r}, E_{m,r}) : m \in \mathcal{M},\ 1 \le r \le R\}$, $\{E_m^{\text{robust}} : m \in \mathcal{M}\}$

TrainingStep applies the scheduled SGD update and returns training moments ξ_k; unregularized methods omit both penalties. UpdateControl leaves fixed strength unchanged, applies Equation (13) for Lite at nonterminal boundaries, or calls Algorithm 1 for Full MACC. Control operations are omitted for unregularized methods. Neither acquisition nor feedback uses official evaluation labels. The later control study follows this loop with the changes in Section 6.4: fixed/scheduled control selection, no official evaluation during tuning, and no additional attack evaluation.

experiments evaluate spectral regularization and coefficient control without learned label acquisition.

5 Analysis of representation geometry and coefficient control

Sections 3 and 4 define the geometry measurements, the aligned objective, and the coefficient updates. We now establish their mathematical properties and identify the conditions under which geometric quantities can inform prediction. The analysis proceeds from the measurements to an isolated regularizer step, the controller response, and conditional error and robustness bounds. The scatter and control results concern the aligned classification method; the spectral variants in Section 4.5 use different objectives and update rules.

5.1 Scatter geometry and acquisition scores

The equal-class weighting in Section 3 gives a variance decomposition even when class counts differ. All quantities in the following proposition refer to the same measurement set.

Proposition 1 (Scatter decomposition and invariance)**.** *For represented classes $\mathcal{C}_\mathcal{S}$ with positive counts, let*

$$T_l = \frac{1}{K} \sum_{c \in \mathcal{C}_\mathcal{S}} \frac{1}{n_c} \sum_{i \in \mathcal{S}_c} \|h_{i,l} - \bar{\mu}_l\|_2^2.$$

Then $T_l = W_l + B_l$. Under Equation (1), $0 \le W_l, B_l \le T_l \le 1$. These moments and the stabilized log ratio q_l are invariant under a common orthogonal transformation of the features. They are also invariant under common positive rescaling of the raw features when the normalization floor is inactive both before and after rescaling.

Proof. Expand $h_{i,l} - \bar{\mu}_l = (h_{i,l} - \mu_{c,l}) + (\mu_{c,l} - \bar{\mu}_l)$. The cross term sums to zero within each class, yielding $T_l = W_l + B_l$. Since $\|h_{i,l}\|_2 \le 1$,

$$T_l = \frac{1}{K} \sum_{c} \frac{1}{n_c} \sum_{i \in \mathcal{S}_c} \|h_{i,l}\|_2^2 - \|\bar{\mu}_l\|_2^2 \le 1.$$

Nonnegativity gives the remaining bounds. Orthogonal maps preserve norms and squared distances; normalization cancels a common positive scale whenever neither normalization activates the floor. □

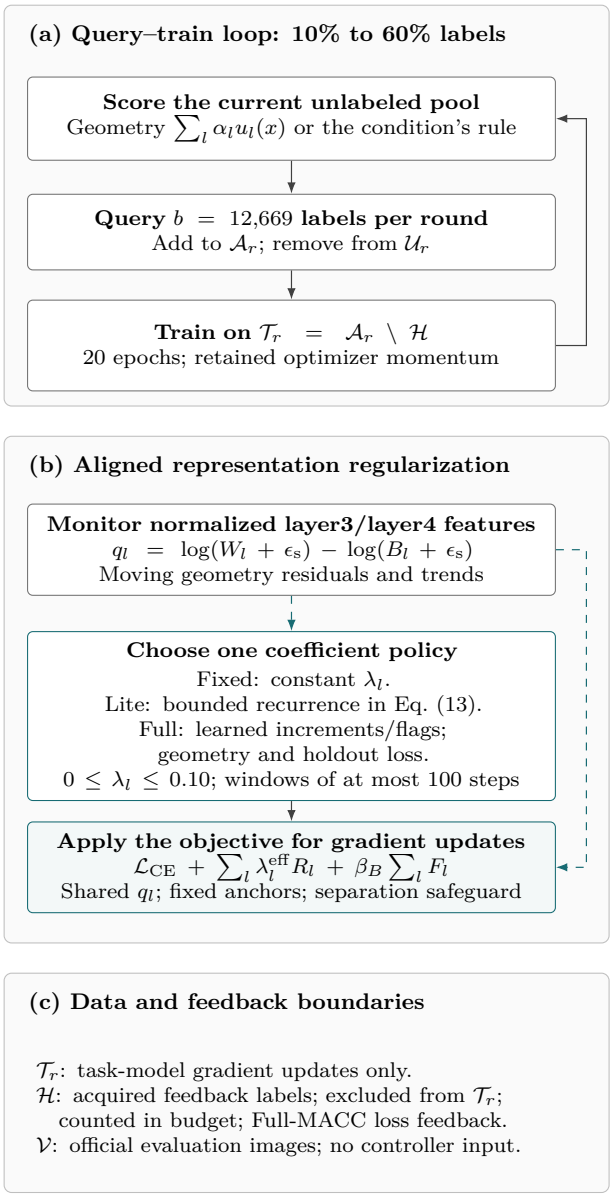


Figure 1: The classification workflow of CÆNL. (a) Acquisition expands the annotated set between training phases. (b) The regularizer and adaptive controllers use the same normalized scatter statistic, while the separation safeguard penalizes insufficient between-class separation. Fixed strength, deterministic feedback, and learned feedback are alternatives. The initial fixed strength is 0.03; the earlier 104-run study selects 0.10 and also tests preset schedules (Section 6.4). Full MACC credits the terminal reward and keeps its acquisition-policy head disabled. (c) The fixed holdout supplies Full-MACC feedback without task-model gradient updates; official evaluation is separate. Solid arrows indicate execution or applied control; dashed arrows indicate measured signals. Language, diffusion, and audio use their task-specific spectral formulations.

Between-class dispersion is an average separation:

$$B_l = \frac{1}{2K^2} \sum_{c,c' \in \mathcal{C}_S} \|\mu_{c,l} - \mu_{c',l}\|_2^2.$$

Consequently, the safeguard in Equation (11) penalizes a small aggregate dispersion, but permits some class means to coincide when others remain separated. This distinction matters for the prediction bound below.

The acquisition score uses a different geometry: raw features and a fitted diagonal metric. Its interpretation follows from the quadratic-form identity, with a centering term retained explicitly.

Lemma 2 (Expected quadratic score)**.** *Let $Z \in \mathbb{R}^d$ have mean μ, covariance Σ, and finite second moments. For a fixed symmetric matrix A and fixed center m,*

$$\mathbb{E}[(Z-m)^\top A(Z-m)] = \operatorname{tr}(A\Sigma) + (\mu - m)^\top A(\mu - m). \tag{18}$$

Proof. Write $Z - m = (Z - \mu) + (\mu - m)$. The mixed terms have zero expectation, while $\mathbb{E}[\operatorname{tr}(A(Z-\mu)(Z-\mu)^\top)] = \operatorname{tr}(A\Sigma)$. □

At the population center, Euclidean squared distance has expectation $\operatorname{tr}(\Sigma)$, whereas squared Mahalanobis distance has expectation d for $A = \Sigma^{-1}$ when Σ is nonsingular. For the population diagonal metric $A = \operatorname{diag}((\Sigma_{jj} + \epsilon_a)^{-1})$, the expectation instead equals $\sum_{j=1}^d \Sigma_{jj}/(\Sigma_{jj} + \epsilon_a)$. Thus the implemented score is variance-standardized deviation, rather than an estimate of the within-class trace. Equation (6) additionally uses estimated centers and a predicted class. Conditioning on fitted references gives the identity with the corresponding conditional moments; conditioning on a predicted-class event can change those moments again. Neither the identity nor a high score establishes the expected predictive value of acquiring that example's label.

5.2 Local effect of the aligned regularizer

Objective–feedback alignment has a precise local meaning: when the ratio target is exceeded, the ratio penalty points along the gradient of that same measured residual. The next result makes the required step-size and smoothness assumptions explicit.

Proposition 3 (Direct local alignment)**.** *Fix a measurement batch, a layer, and its target, and define $g(\theta) = q_l(\theta) - q_l^\star$. Suppose $g(\theta) > 0$ and g has an L_g-Lipschitz gradient, $L_g > 0$, on a convex neighborhood containing the step segment. Hold the applied coefficient $\lambda = \lambda_l^{\text{eff}} \ge 0$ fixed and let $a = H'_\delta(g(\theta)) = \min\{g(\theta)/\delta, 1\}$. For $\eta \ge 0$, the isolated ratio-penalty step $\theta' = \theta - \eta\lambda a \nabla g(\theta)$ satisfies*

$$g(\theta') \le g(\theta) - \eta\lambda a \left(1 - \frac{L_g \eta \lambda a}{2}\right) \|\nabla g(\theta)\|_2^2. \tag{19}$$

Hence g and $H_\delta(g)$ do not increase when $0 \le \eta\lambda a \le 2/L_g$. The residual decreases strictly if $0 < \eta\lambda a < 2/L_g$ and $\nabla g(\theta) \ne 0$.

Proof. Apply the smoothness inequality $g(\theta + v) \le g(\theta) + \nabla g(\theta)^\top v + (L_g/2)\|v\|_2^2$ with $v = -\eta\lambda a \nabla g(\theta)$. The penalty H_δ is nondecreasing in its argument. □

The result applies to a fixed or adaptively selected coefficient once its value is held fixed for the model update, as in Section 4.1. It establishes the local direction

of the isolated intervention, rather than an advantage of feedback over fixed strength. When $g \le 0$, the ratio term has zero gradient. For the complete training update, cross-entropy, the separation penalty, other layers, momentum, and changing batches can alter the direction. The smoothness condition is local and need not hold across activation or normalization-floor boundaries. For the actual optimizer direction v, defined by $\theta^+ = \theta - \eta v$, suppose the same local smoothness assumption holds on a convex neighborhood containing the actual update segment. It then yields

$$g(\theta^+) \le g(\theta) - \eta\langle \nabla g(\theta), v\rangle + \frac{L_g \eta^2}{2}\|v\|_2^2. \tag{20}$$

If $v \neq 0$ and $\langle \nabla g(\theta), v\rangle > 0$, strict descent follows for $0 < \eta < 2\langle \nabla g(\theta), v\rangle/(L_g\|v\|_2^2)$. For plain gradient descent, write $p = \nabla g(\theta)$ and $v = d + \lambda a p$, where d contains the other gradient terms. Then $\langle p, v\rangle = \langle p, d\rangle + \lambda a\|p\|_2^2$; the other losses may oppose the isolated intervention. With momentum, clipping, or weight decay, v must denote the resulting realized direction. The scalar residual and coefficient bounds alone do not ensure this directional condition or a monotone full-training trajectory.

5.3 Bounded coefficients and the feedback response

MACC-Lite combines a bounded desired strength with a leaky update. Its coefficient bounds hold for every finite residual sequence, including the first phase-start action described in Section 4.4.

Proposition 4 (Bounded feedback coefficients)**.** *Let $0 \le \lambda_{l,0} \le \lambda_{\max}$, $0 < \rho \le 1$, and $\delta > 0$. Equation (13) keeps every coefficient in $[0, \lambda_{\max}]$. For $t \ge 1$, its response to an arbitrary sequence of desired strengths is*

$$\lambda_{l,t} = (1-\rho)^t \lambda_{l,0} + \rho \sum_{j=0}^{t-1} (1-\rho)^{t-1-j} u_{l,j}.$$

If $u_{l,j} = u_l$ is constant, this reduces to

$$\lambda_{l,t} - u_l = (1-\rho)^t (\lambda_{l,0} - u_l). \tag{21}$$

Proof. The map $\lambda_{\max}\tanh([e]_+/\delta)$ lies in $[0, \lambda_{\max}]$. Each update is a convex combination of this value and the previous coefficient, proving invariance of the interval by induction. Unrolling the recurrence gives the arbitrary-input formula; subtracting a constant u_l gives Equation (21). □

With $\rho = 0.20$, the fixed-input coefficient error contracts by a factor of 0.8 per control event. If the residual remains nonpositive, $u_l = 0$ and the coefficient decays toward zero; this does not deactivate the separate separation penalty. Full MACC also has bounded applied coefficients because its projection and binary flags preserve $[0, \lambda_{\max}]$, but it does not follow the Lite recurrence. These are properties of the coefficient mechanism. In the coupled system, residuals depend on the evolving network and moving measurements, so coefficient boundedness alone establishes neither target convergence nor closed-loop stability.

5.4 Conditions linking geometry to prediction

Within-class concentration can control prediction error when class separation and the decision rule are specified. The following population statement identifies these additional conditions.

Proposition 5 (Conditional nearest-centroid error bound)**.** *Let H be a feature vector with finite second moments and equal class probabilities over $C \ge 2$ classes. Define population means $\mu_c = \mathbb{E}[H \mid Y = c]$, minimum separation $d_{\min} = \min_{c \neq c'} \|\mu_c - \mu_{c'}\|_2 > 0$, and*

$$W_{\text{pop}} = \frac{1}{C}\sum_c \mathbb{E}[\|H - \mu_c\|_2^2 \mid Y = c].$$

The classifier $\widehat{Y}(H) = \arg\min_c \|H - \mu_c\|_2$, with any fixed tie rule, satisfies

$$\Pr\{\widehat{Y}(H) \neq Y\} \le \min\{1, 4W_{\text{pop}}/d_{\min}^2\}. \tag{22}$$

Proof. If $\|H - \mu_Y\|_2 < d_{\min}/2$, the triangle inequality makes every other centroid farther away. An error therefore requires $\|H - \mu_Y\|_2^2 \ge d_{\min}^2/4$. Markov's inequality and averaging over the equally probable classes give the bound, capped by one. □

Here W_{pop} is a population moment, distinct from the finite-batch W_l used in training. The trained softmax head is not constrained to implement nearest-centroid prediction, and a positive aggregate B_l does not imply $d_{\min} > 0$. Consequently, reducing the batch ratio does not by itself tighten this population error bound.

Input-space robustness. A corresponding robustness condition concerns logit margins and input sensitivity, as in Lipschitz-based certification [48]. For a correctly classified input x with true class y, let $m_{y,j}(x) = f_\theta(x)_y - f_\theta(x)_j$. Suppose each margin is $L_{y,j}$-Lipschitz with respect to $\|\cdot\|_\infty$ throughout the admissible radius-ϵ neighborhood. Every admissible perturbation v with $\|v\|_\infty \le \epsilon$ then satisfies

$$m_{y,j}(x+v) \ge m_{y,j}(x) - \epsilon L_{y,j}.$$

Thus $m_{y,j}(x) > \epsilon L_{y,j}$ for every $j \neq y$ is sufficient to preserve the correct class. Neither the scatter nor the spectral objective supplies the required bounds on $L_{y,j}$. Section 6 therefore specifies separate evaluations of predictive performance, representation geometry, and adversarial robustness, together with matched controls for the value of coefficient adaptation.

The need for input-sensitivity control is not merely a missing step in the preceding sufficient condition. The following construction preserves the clean measurements while changing robustness.

Proposition 6 (Identical clean geometry with arbitrarily small robust radius)**.** *There is a four-example data set in* $[0,1]^2$*, with two distinct examples in each of two classes, and a family of finite ReLU feature maps with a fixed linear classifier such that every member has identical clean features, positive within-class and between-class scatter, identical clean logits and cross-entropy, and perfect clean accuracy. For every* $\epsilon > 0$*, a member admits an admissible misclassifying perturbation of* L_∞ *norm less than* ϵ *at every clean example.*

Proof. Set $b = 1/5$ and $c = \sqrt{24/25}$. Assign class 1 to $(1/4, 2/5), (1/4, 3/5)$ and class 2 to $(3/4, 2/5), (3/4, 3/5)$. For $0 < \eta < 1/8$, define the ReLU triangle

$$T_{t,\eta}(u) = \frac{[u-t]_+ - 2[u-t-\eta]_+ + [u-t-2\eta]_+}{\eta}.$$

It vanishes outside $[t, t+2\eta]$, is zero at t, and is one at $t+\eta$. Define the finite ReLU map

$$a_\eta(u) = c(2-4u) - 2cT_{1/4,\eta}(u) + 2cT_{3/4,\eta}(u),$$
$$z_\eta(u,v) = (a_\eta(u), b(10v-5)),$$

and fixed classifier weights $w_1 = (1,0)$, $w_2 = (-1,0)$, with logits $f_j(x) = w_j^\top z_\eta(x)$. The triangles vanish at every clean example. Their features are $(c,-b), (c,b), (-c,-b), (-c,b)$ for every η, each of norm one. Normalization therefore leaves them unchanged, and equal class weighting gives

$$W = b^2 = 1/25, \qquad B = c^2 = 24/25, \qquad W + B = 1.$$

The stabilized log ratio, frozen-target penalties, clean margin $2c$, and cross-entropy $\log(1+e^{-2c})$ are independent of η.

Perturb each clean input by $(\eta, 0)$. The inputs remain in $[0,1]^2$ and the triangle supports do not overlap. The first feature coordinate becomes $-c(1+4\eta) < 0$ for class 1 and $c(1-4\eta) > 0$ for class 2, so all four predictions flip. The second coordinate remains $\pm b$, keeping the raw norm above the normalization floor; normalization does not change the decision. Choose $0 < \eta < \min\{\epsilon, 1/8\}$ to obtain the stated perturbation. □

Removing the two triangles preserves the same clean features and logits but places the decision boundary at $u = 1/2$, at input distance $1/4$ from every clean example. The constructed family has no uniform weight or input-sensitivity bound as η decreases. It proves that the measured clean geometry and logits alone cannot supply a uniform positive lower bound on robust radius across such feature maps; it does not identify the mechanism in the trained models. Matched measurements of margins, perturbation-induced feature displacement, and attack success are still required for that empirical question.

6 Experimental design

The evaluation follows the distinction between representation objectives and coefficient control developed in Sections 4 and 5. Classification studies compare acquisition rules and regularization treatments at matched annotation budgets, then test feedback against independently selected constant and scheduled strengths. The final classification study additionally holds acquired labels fixed to isolate training components and uses an inventory filtered for exact content identities before any of its training. Separate language, diffusion, and audio studies evaluate the spectral variants under their task-specific protocols.

6.1 Study organization

Throughout this paper, a *condition* is a method configuration, a *run* is one condition under one seed, and a *phase* is one acquisition/training stage within a run.

The initial ImageNet-100 comparison evaluated nine conditions across six paired seeds. After reviewing those results, we evaluated two additional conditions: entropy acquisition with aligned DCR and MACC-Lite, and a matched-protocol UHerding-Margin comparator. Each used the original six seed-specific initializations, targets, splits, annotation budgets, and evaluation endpoint. Their comparison lists were fixed before the additional runs, and their statistical families are reported separately. A third follow-up added entropy with fixed aligned DCR, retaining the same per-seed initial model and optimizer, splits, targets, floor penalty, acquisition protocol, and optimization schedule. Its two contrasts were declared before the new fixed-condition outcomes were inspected, after the entropy and Lite outcomes were already known. This later control adds six runs and 30 phases; it is not fresh-seed confirmation. The resulting classification record for these studies comprises 72 method–seed runs and 360 post-initialization phases across 12 conditions, with six paired seeds.

A separate study then tunes fixed and preset-schedule controls on three new ResNet-50 seeds, confirms the selected comparisons on ten further ResNet-50 seeds, and transfers all settings to ten ResNet-18 seeds. Its 104 method runs, 23 shared initializations, and 520 post-initialization phases are additional experiments. The tuning, confirmation, and transfer seed sets are disjoint. They are not pooled with the earlier six-seed comparisons, and fresh seeds do not make the previously studied dataset an untouched benchmark.

A third classification study uses a prospectively content-filtered inventory and disjoint tuning seeds 61001–61003, ResNet-50 confirmation seeds 62001–62010, and ResNet-18 transfer seeds 63001–63010. It comprises 36 tuning runs, 110 confirmation runs (60 end-to-end and 50 common-label replays), and 60 transfer runs: 206 method runs, 23 shared initializations, 1,030 post-initialization phases, and 8,141,080 optimizer steps including initializations. Its 17-contrast inferential family is separate from both earlier classification studies. These three records retain their original observations rather than pooling seeds across changing designs. The 206-run study and its separately specified robustness and practical-cost diagnostics are complete. The latter evaluates 30 frozen ResNet-50 checkpoints

and 150 attack suites, without retraining or changing the original 17-contrast analysis. Its three primary comparisons form a separate inferential family.

Engineering checks verify implementation behavior and are not counted as predictive-performance experiments. The original ImageNet-100 spectral pilot uses seeds 601 and 602. Saved-checkpoint attack diagnostics and a bounded aligned-objective development screen precede the final seeds 801–806. The aligned and spectral development protocols also differ in sampling, BatchNorm treatment, and integration choices; their absolute accuracies cannot isolate the effect of changing DCR.

C4 uses six separate training seeds and an additional unused-shard evaluation of the same saved checkpoints. Diffusion development uses seeds 1001 and 1002, followed by confirmation seeds 1101–1106. Audio development uses seeds 1201 and 1202 for two epochs per condition; confirmation uses seeds 1301–1306 for ten epochs. Confirmation starts from the original pretrained assets and new seed-specific task components, not development-final checkpoints. Only audio confirmation evaluates caption-model predictions on the retained official test split. Development results are not pooled into confirmation inference.

A subsequent raw-source AudioCaps study comprises 18 fixed-strength selection runs on seeds 81001–81003 and 40 confirmation runs on seeds 82001–82010. It re-extracts audio features from the newly audited recordings and retains a separate five-contrast family. The earlier 24 audio runs, their available subset, and their outcomes are preserved; neither seed counts nor results are pooled across the two inventories.

6.2 ImageNet-100 data and label accounting

We use the fixed 100-class subset associated with Contrastive Multiview Coding [27], drawn from ImageNet [28]. The inventory used by the earlier 72- and 104-run studies contains 126,689 training images and 5,000 validation images, with 50 validation images per class. Original image bytes and dimensions are retained. A manifest records class mappings, source revision, paths, and SHA-256 values, which are checked before execution. The evaluated benchmark is ImageNet-100, not ImageNet-1K.

Stratified sampling supplies 12,669 initial annotations per seed. The controller holdout contains 633 of these, leaving 12,036 gradient-training images. Five rounds each acquire 12,669 additional examples. Annotation totals are 12,669, 25,338, 38,007, 50,676, 63,345, and 76,014. The final gradient-training set contains 75,381 images. The final annotation fraction is approximately 60.0005%, reported as 60%. Holdout labels are included in every method's annotation cost. Pool labels are consulted only after selection.

The split procedure excludes holdout indices from gradient training and acquisition. A subsequent content-hash audit identifies five training–validation image matches and duplicate content within the training inventory. The frozen benchmark results are retained, with post-hoc duplicate-exclusion sensitivity reported in Section 7.5. Index-disjoint partitions in those earlier studies are therefore not asserted to be content-disjoint. The subsequent 206-run study applies content exclusions before training and has different absolute counts (Section 6.5).

6.3 Classification optimization and matching

The initial study and later confirmation use a randomly initialized ResNet-50 with a 100-class head; the transfer stage uses ResNet-18 [29]. Geometry is measured at the spatially pooled outputs of layer3 and layer4. Each seed has one shared 40-nominal-epoch cross-entropy initialization, reused by the original methods and later matched conditions within that seed. The new control study uses independently initialized models for its own seed sets. A physical minibatch contains 16 uniformly selected represented classes and eight distinct acquired examples per class. A nominal epoch has $\lceil |\mathcal{T}_r|/128 \rceil$ sampled minibatches rather than a permutation pass through the labeled set.

Training uses random resized 224×224 crops and horizontal flips. Scoring, feedback, and evaluation resize the shorter side to 256 and take a 224-pixel center crop. ImageNet normalization is inside the model, so attacks act in pixel space. Augmentation and sampling streams are keyed by seed, phase, step, and example index. Common random-number conventions do not make sampled images identical once acquired sets differ.

Optimization uses SGD with momentum 0.9, Nesterov acceleration, and weight decay 10^{-4} on multidimensional parameters. One-dimensional parameters receive no decay. Parameters remain FP32 with BF16 autocast, and gradients are clipped at norm 5. Initialization uses learning rate 0.05, three nominal warm-up epochs, and cosine decay. Each subsequent 20-epoch phase starts at rate 0.025 with one warm-up epoch and cosine decay; momentum is retained. BatchNorm running statistics are updated during initialization and frozen afterward; affine parameters remain trainable.

In the earlier studies, the five acquired-label phases contain 3,880, 5,840, 7,820, 9,800, and 11,780 optimizer steps. Initialization contains 3,800 steps. Each conceptual condition run therefore includes 42,920 steps and 5,493,760 sampled example presentations. Initialization is executed once per seed and allocated once per conceptual condition when reporting comparable costs.

6.4 Earlier tuned controls and backbone transfer

All four treatments share entropy acquisition, annotation budgets, and the optimization protocol. The three regularized treatments share the aligned loss, target calibration, and fixed separation penalty; entropy alone omits both representation penalties. The experiment therefore compares the complete regularized-feedback treatment with entropy alone, and compares coefficient

Table 2: Initial ImageNet-100 comparison. All nine conditions use seeds 801–806 and the same 10–60% annotation schedule. "None" disables both aligned penalties; all conditions record geometry. The three follow-up conditions are reported separately.

Condition	Acquisition	Representation control
Random	Uniform random selection	None
Entropy	Highest predictive entropy	None
BADGE	Exact factorized weight-gradient distances	None; independent implementation, bias excluded
Noise Stability	Perturbed-output deviations with greedy coverage	None; five noise draws, independent implementation
CÆNL acquisition only	Predicted-class diagonal Mahalanobis score	None
Random + aligned fixed DCR	Uniform random selection	Aligned objective, $\lambda_l = 0.03$
CÆNL + aligned fixed DCR	Mahalanobis score	Aligned objective, $\lambda_l = 0.03$
CÆNL + aligned MACC-Lite	Mahalanobis score	Bounded, smoothed coefficient feedback
CÆNL + aligned Full MACC	Mahalanobis score	Learned coefficient increments and activation flags

rules within the same regularized treatment. Four constant strengths, 0.003, 0.01, 0.03, and 0.10, and four linear schedules are evaluated on ResNet-50 seeds 41001–41003, giving 24 tuning runs. Each schedule is either increasing or decreasing, with peak $a \in \{0.03, 0.10\}$:

$$\lambda_l(p) = ap \quad \text{or} \quad a(1-p), \qquad p \in [0,1]. \tag{23}$$

Progress p is the cumulative fraction of the 39,120 post-initialization optimizer steps across all five rounds. Strengths refresh at phase starts and nonterminal 100-step control boundaries, using the same event cadence as Lite. Schedules receive no observation feedback. Both layers use the same fixed or scheduled strength.

Within each family, selection maximizes mean final accuracy on the 633-example training holdout across the three tuning seeds. Accuracy ties within 10^{-12} are broken by lower mean holdout cross-entropy, then candidate identifier. Official validation is not evaluated during tuning. The selected constant and schedule are frozen before confirmation begins. Lite retains $\lambda_{\max} = 0.10$, initial strength 0.03, smoothing 0.20, and Huber scale 0.25 from the earlier development, without additional tuning in this study.

Confirmation uses ResNet-50 seeds 42001–42010; transfer uses ResNet-18 seeds 43001–43010. Each stage compares entropy alone, entropy with Lite, entropy with the selected constant, and entropy with the selected schedule, for 40 runs per backbone. Every seed has a shared initial model, targets, holdout and first acquisition. Later labeled sets can diverge with the model predictions. Each backbone is initialized from scratch, and ResNet-18 receives no retuning. This is architecture transfer on one dataset, not transfer of trained weights or evaluation on a new image benchmark.

6.5 Content-filtered replication and component isolation

Inventory and annotation accounting. Before this study's training, decoded-RGB identities (dimensions and pixel bytes) are grouped alongside original file hashes. Every member of a conflicting-label or cross-split identity group is excluded; within one split and label, only the lexicographically first duplicate is kept. The audit identifies 56 duplicate groups, including 48 conflicting-label and five cross-split groups. The policy excludes 104 images: 96 for label conflicts and eight same-label duplicates. It retains 126,590 training and 4,995 validation images. Original source images remain unchanged. This addresses exact-content overlap for the new study, without certifying absence of near-duplicates or authenticating upstream labels.

Each seed begins with 12,659 annotations and acquires another 12,659 in each of five rounds. The totals are 12,659, 25,318, 37,977, 50,636, 63,295, and 75,954 (exactly 60% of the retained training inventory). The 632 feedback-holdout labels are included in these counts and excluded from gradient training. Benchmark labels are used to curate content groups; this simulation's annotation budget does not include the cost of obtaining or curating the benchmark itself.

Control selection and matched references. The sampler, layers, optimization, augmentation, target calibration, and frozen post-initialization BatchNorm follow the earlier protocol. Each method has 39,100 post-initialization optimizer steps: 3,860, 5,840, 7,820, 9,800, and 11,780 across its five phases. The smaller inventory slightly changes the earlier step count. The shared 40-epoch initialization has 3,760 steps and is executed once per seed. Four fixed DCR strengths $\{0.003, 0.01, 0.03, 0.10\}$, four linear schedules (increasing or decreasing, peak 0.03 or 0.10), and four contrastive-feature strengths in the same fixed grid give 12 candidates on each of three tuning seeds. Each family is selected by mean final training-holdout accuracy, then lower cross-entropy for accuracy ties within 10^{-12}, then identifier. Official validation is not used for selection. Lite retains its earlier settings without a new search. All settings are frozen for transfer.

The contrastive comparator uses the positive-pair log-probability form of supervised contrastive learning [15], at temperature 0.1, averaged over normalized layer3/layer4 features. With eight examples per class, each anchor has seven positive partners. It uses one view and no projection head; it is a matched feature-

penalty comparison, not the original two-view SupCon training recipe. Each backbone compares entropy, Lite, selected fixed DCR, selected schedule, selected contrastive regularization, and a full-label reference. The latter uses all 126,590 training annotations from round 1, excluding the same feedback holdout from gradients, with the same optimizer-step and learning-rate schedule as active learning. It is a compute-matched reference rather than a separately optimized full-supervision solution trained to demonstrated convergence.

Common-label component replay. For each ResNet-50 confirmation seed, the entropy-only run supplies its complete acquired-label sequence to five additional training conditions: ratio-only fixed DCR; separation-only regularization; both penalties with fixed strength; both with Lite; and both with the selected schedule. The ratio strength is the selected fixed value; separation weight is 0.01, as in the full objective. The entropy, ratio-only, separation-only, and fixed-both cells form a 2×2 component comparison. Reduced objectives are not independently retuned, so removal effects are conditional on strengths selected for the complete objective. Replay shares initialization, targets, holdout, acquired labels, minibatch identities, and augmentation seeds; batch-chain parity is checked. It reuses labels, not features or gradients. These contrasts isolate training interventions conditional on entropy's label sequence. They complement the end-to-end comparisons, whose acquired sets may diverge, rather than decomposing those pipeline effects automatically.

6.6 Acquisition baselines and matched comparisons

Table 2 separates acquisition, regularization, and adaptation. BADGE uses exact factorized last-linear-layer weight-gradient inner products rather than the pilot's random projection; bias terms are excluded. Noise Stability uses five globally normalized Gaussian perturbations of the trainable parameter vector at relative scale 0.001, normalized softmax-output deviations, concatenation without projection, and greedy farthest-first selection. Model parameters and buffers are restored after noisy forwards.

Nonrandom methods cache required pool outputs at the current checkpoint. Mahalanobis acquisition additionally estimates class statistics from acquired examples. Equal annotation budgets and pass sizes do not equalize acquisition time, especially for Noise Stability. Random plus aligned fixed DCR tests whether the objective contributes without geometric acquisition. The final matrix has no new full-label reference; the pilot's different-budget supervised reference is not mixed into the matched final comparison.

6.7 UHerding-Margin: matched-protocol implementation

The recent comparator implements the uncertainty-weighted coverage mechanism of Uncertainty Herding [19], as `uherding_margin_initial_r50`. Each seed's shared-initial supervised ResNet-50 supplies unit-normalized layer4 features g_i, cached for the full training pool and fixed across rounds. This differs from a separately pretrained self-supervised representation. The task classifier continues under the common warm-start schedule.

Before acquisition, temperature T minimizes 15-bin ECE over 81 logarithmically spaced values in $[0.1, 10]$ using the budgeted 633-image holdout. Temperature-scaled probabilities define $\omega_i = 1 - (p_{i,(1)} - p_{i,(2)})$. The bandwidth σ is the smallest positive Euclidean distance between fixed features of currently annotated examples, including held-out annotations. Squared distances at most 10^{-12} are numerical zero. The kernel is $k_\sigma(g_i, g_j) = \exp(-\|g_i - g_j\|_2^2/\sigma^2)$.

For candidate pass P, greedy selection maximizes

$$G_P(j \mid \mathcal{C}) = \sum_{i \in P} \omega_i \left[k_\sigma(g_i, g_j) - \max_{a \in \mathcal{C}} k_\sigma(g_i, g_a) \right]_+ , \tag{24}$$

where $\mathcal{C}$ contains annotated centers and examples already selected in the round. The 10,000-candidate and 10%-per-pass convention is retained until 12,669 new examples are acquired. Distances, kernels, and gains use float64. Lazy greedy is checked against exhaustive greedy on deterministic CPU and CUDA fixtures.

No DCR is applied. Initializations, optimizer states, label accounting, balanced batches, post-initialization BatchNorm treatment, and the final attack endpoint match the task protocol. Feature extraction, acquisition including calibration, and training are timed separately. This comparison evaluates the disclosed adaptation, not the original authors' representation, unrestricted selection domain, or published benchmark scores.

6.8 Classification evaluation and statistical analysis

Every post-initialization phase in the earlier 72- and 104-run studies' confirmation/transfer stages is evaluated on all 5,000 official validation images in full precision. Shared initializations and tuning in the later control study do not evaluate official validation. Reported outcomes include pooled cross-entropy, top-1 and top-5 accuracy, 15-bin ECE, and equal-class normalized moments. The selected checkpoint is the final scheduled checkpoint, not the best validation epoch. Logits, predictions, labels, and compact class-moment sufficient statistics are retained.

The final checkpoint in the original six-seed studies receives standard AutoAttack [25] at L_∞ radius 8/255 on all 5,000 images. Its components are APGD-CE, targeted APGD, targeted FAB, and Square Attack, without shortened prescribed iterations. Attack minibatches contain 64 images; restartable storage chunks

contain 256. Pixel bounds, perturbation norms, and survival masks are recorded. No adversarial training is used. The 104- and 206-run control studies add no adversarial evaluations. A compatibility repair changed `view` to `reshape` in the library's zero-gradient diagnostic, retaining the diagnostic, attack configuration, and completed training.

For the original classification studies, the paired experimental units are six seeds, conditional on the fixed dataset and protocol. Seven contrasts are specified separately for final clean and robust accuracy. For target-minus-reference differences d_i, the individual paired-t interval is

$$\bar{d} \pm t_{0.975,5} \frac{s_d}{\sqrt{6}}.$$

Both paired-t tests and enumeration of all 2^6 sign flips are retained. Holm correction [30] is applied within each declared endpoint and test family. Intervals are individual, unadjusted intervals, not simultaneous confidence bands. The paired-t analysis relies on its distributional approximation for the differences. Interpreting sign-flip enumeration as an exact null test requires the corresponding sign-exchangeability or symmetry assumption; enumeration alone does not make it assumption-free.

Resolution of the six-seed exact tests. For the two-sided absolute-mean statistic, a sign pattern and its negation are equally extreme. Hence

$$p_{\text{exact}} \geq 2/2^6 = 0.03125. \tag{25}$$

For a Holm family of $m \geq 2$ comparisons at $\alpha = 0.05$, the first threshold is $0.05/m \leq 0.025$. No rejection is therefore attainable with these six-seed, two-sided, Holm-adjusted exact tests. This is a resolution limit, not evidence that the effects are absent. Raw and adjusted values are retained alongside the paired-t analysis.

The entropy-plus-Lite, UHerding, and entropy-plus-fixed follow-ups each use a separate two-contrast clean-accuracy family, specified before the respective additional runs but after earlier outcomes were available. For UHerding, the contrasts are UHerding minus entropy and entropy-plus-Lite minus UHerding. For the fixed-control follow-up, they are entropy-plus-fixed minus entropy and entropy-plus-Lite minus entropy-plus-fixed. New fixed-condition robustness is descriptive; the original robustness family is unchanged. These within-family corrections do not provide one family-wise error guarantee for the full sequential research programme. Geometry, calibration, top-5 accuracy, and intermediate budgets are descriptive secondary outcomes.

Ten-seed control study. The later study declares six final-top-1 contrasts: Lite minus entropy, tuned fixed, and tuned schedule, separately on each backbone. These form one Holm family across both backbones, with separate corrections for paired-t and exact sign-flip tests. Each contrast uses ten paired seeds; its individual interval uses $t_{0.975,9} s_d/\sqrt{10}$ and its exact test enumerates 2^{10} signs. The smallest attainable two-sided raw value is $2/2^{10} = 0.001953125$, so adjusted rejection is possible in this family. The six-seed resolution restriction above does not apply to this study. Sample size and comparisons were fixed before execution, without outcome-dependent stopping or seed replacement. A 0.5-percentage-point practical reference is descriptive, not an equivalence margin or stopping rule. Intermediate budgets, controller dynamics, secondary metrics and duplicate-exclusion analyses do not redefine the primary endpoint.

Content-filtered 17-contrast family. The 206-run study evaluates final checkpoints on all 4,995 retained validation images. Ten paired seeds are the experimental units for each contrast. Five Lite-minus-reference contrasts on each backbone compare entropy, fixed DCR, schedule, contrastive regularization, and full labels. Seven ResNet-50 common-label contrasts compare ratio minus entropy, separation minus entropy, fixed-both minus ratio, fixed-both minus separation, Lite minus fixed-both, schedule minus fixed-both, and the factorial interaction (fixed-both minus ratio minus separation plus entropy). These 17 comparisons form one prespecified Holm family, separately for paired-t and all 2^{10} sign-flip tests. Intervals are individual, unadjusted 95% paired-t intervals, not simultaneous bands; the assumptions stated above still apply. Intermediate budgets, cross-entropy, calibration, geometry, and timings are secondary. Ten seeds allow adjusted exact-test rejection in this family but do not guarantee power. Neither a nonsignificant full-label difference nor the descriptive 0.5-point reference defines equivalence or label savings. Appendix C retains every planned contrast, selection result, and paired final accuracy.

6.9 Cross-task experimental protocols

These protocols concern Equation (16). The fixed, Lite, and Full names identify task-specific implementations; the aligned classification loss is not used in these studies.

6.9.1 Language modeling on C4

Pretrained GPT-2 [31], with 124,439,808 parameters, is fine-tuned on a deterministic C4 subset [32]. Each of seeds 301–306 processes 99,999,744 tokens in 12,207 steps, with sequence length 1,024 and eight sequences per step. Baseline, fixed spectral DCR, spectral MACC-Lite, and spectral Full MACC share the budget and training configuration. The spectral implementation monitors h6 and h12 with target ranks 16 and 4, retained spectrum size 128, and a 2,048-token DCR subsample. Fixed strength is 0.01.

Training uses the first English C4 training shard, with the first validation shard used for feedback and development (Appendix F). All four conditions therefore evaluate the same single-shard training distribution. The original reporting set has 1,889 validation

Table 3: Task-specific regularization and feedback conventions. "Lite" and "Full" denote different task-specific implementations, not an unchanged controller transferred across all tasks. Task protocols provide the remaining settings.

Task	Loss and observed statistic	Fixed/adaptive coefficients	Update convention
ImageNet-100 classification	Aligned normalized within/between log ratio; fixed separation safeguard	Fixed 0.03 in the initial study and selected 0.10 in the later study; adaptive initialization 0.03; bounds $[0, 0.10]$	Aligned Lite uses the leaky rule; Full uses the aligned reward; windows of at most 100 steps and phase-boundary handling
GPT-2/C4	Spectral loss with $k = 128$, ranks 16 and 4 at h6/h12; separate covariance-spectrum feedback	Fixed 0.01; adaptive initialization 0.005; bounds $[0, 0.08]$	Spectral control every 25 steps; Lite step size 0.002 and deadband 0.1
CIFAR-10 diffusion	Spectral loss with $k = 64$, rank-16 target; full EMA-covariance monitor	Fixed 0.01; adaptive initialization 0.01; bounds $[0, 0.05]$	Spectral control every 50 steps; Lite step size 0.001 and deadband 0.1
AudioCaps captioning	Spectral loss with $k = 64$, rank-16 targets at prefix/encoder; different sample axes and zero padding	Fixed 0.01; adaptive initialization 0.01; bounds $[0, 0.05]$	Spectral control every 50 steps and final step; terminal reward credited without an unused next action

sequences, with 64 additional sequences reserved for controller feedback.

After checkpoint selection and training were frozen, all 24 checkpoints were evaluated on validation shard 00002, which had not been used for local training, feedback, or configuration development. Each checkpoint makes 1,999,872 next-token predictions. NLL and accuracy are pooled over tokens; ECE is computed from pooled confidence-bin statistics. Perplexity is the exponential of NLL, not an independent endpoint. "Untouched" concerns local study use, not verified absence of GPT-2 pretraining overlap.

Optimization uses learning rate 5×10^{-5}, weight decay 0.01, and 200 warm-up steps. AdamW has $\beta_1 = 0.9$, $\beta_2 = 0.95$; one-dimensional, bias, and normalization parameters are exempt from decay. Norm-1 clipping and linear warm-up followed by cosine decay toward 10% of the base rate are used. Parameters remain FP32 with BF16 CUDA autocast.

Spectral controllers operate every 25 steps. Both adaptive methods initialize their coefficients at 0.005 within $[0, 0.08]$. Spectral Lite uses step size 0.002 and deadband 0.1. Full MACC uses two 128-unit hidden layers, policy learning rate 3×10^{-4}, exploration probability 0.15 decreasing to 0.02, reward-baseline weight 0.1, and entropy coefficient 0.001. Its coefficient increments are $\{-0.005, -0.0025, 0, 0.0025, 0.005\}$. Its language-specific reward settings include validation-delta scale 5.0, value-delta clip 0.1, target-deviation weight 0.005, and zero query penalty. These settings are not replaced by the later diffusion/audio or aligned-classification controller settings.

The unused-shard evaluator compares fixed DCR, Lite, and Full separately with baseline, giving three contrasts per endpoint. It reports individual paired-t intervals and raw paired-t and exact sign-flip values. Holm adjustment is applied separately to the paired-t and exact values within each three-contrast endpoint family. NLL is primary; perplexity is its monotone transformation, while token accuracy and ECE remain secondary.

6.9.2 CIFAR-10 diffusion

We fine-tune the 35,746,307-parameter pretrained `google/ddpm-cifar10-32` U-Net with the noise-prediction objective [33]. Seeds 1101–1106 evaluate baseline fine-tuning, fixed spectral DCR, spectral MACC-Lite, and spectral Full MACC. The unchanged pretrained model is also sampled with each seed's generation-noise stream. Those six references measure sampling variation for one model, not six pretrained-model replications.

A fixed partition of the original 50,000 CIFAR-10 training images [34] supplies 45,000 fine-tuning images and 5,000 internal holdout images. The holdout contains 256 controller images and 1,024 disjoint denoising-assessment images with fixed noise and timesteps. These images may have been seen during pretraining; the holdout is not pristine relative to the pretrained model. Official test outcomes do not enter control or checkpoint selection.

Each trained condition performs 10,000 updates with batch size 128, matched per-seed image order and noise streams, horizontal flips, AdamW learning rate 2×10^{-5}, zero weight decay, and norm-1 clipping. Parameters are FP32 with BF16 autocast. EMA weights use decay 0.9995. The final scheduled checkpoint is retained independently of performance. No active label acquisition occurs.

DCR uses 256-dimensional spatially pooled mid-block features across timesteps 0–999, retained spectrum size 64, and target rank 16. The full EMA-covariance monitor uses decay 0.99 and cadence 50 steps. Fixed strength is 0.01. Adaptive methods start at 0.01 within $[0, 0.05]$. For finite monitored rank $r_{\mathrm{eff},l,t}$, spectral Lite uses

$$e^{\mathrm{sp}}_{l,t} = \left|\log\left(\frac{\max\{r_{\mathrm{eff},l,t}, 10^{-12}\}}{16}\right)\right| - 0.1,$$
$$\lambda_{l,t+1} = \Pi_{[0,0.05]}\left(\lambda_{l,t} + 0.001 e^{\mathrm{sp}}_{l,t}\right), \tag{26}$$

where $\Pi_{[a,b]}(v) = \min\{b, \max\{a, v\}\}$. The error is not replaced by its positive part: the coefficient increases outside the log-rank deadband, decreases inside it, and

is unchanged at its boundary, subject to projection. Missing or nonfinite rank observations leave the corresponding coefficient unchanged. This accumulating projected rule differs from the leaky classification rule in Equation (13).

Spectral Full MACC uses the retained two-layer 128-unit policy, increments $\{-0.005, -0.0025, 0, 0.0025, 0.005\}$, activation flags, policy learning rate 3×10^{-4}, exploration decreasing from 0.10 to 0.02, baseline weight 0.05, and entropy coefficient 0.001. Its reward combines the scaled held-out loss improvement (scale 10, clipped to $[-2, 2]$) and a $-0.1|\log(r_{\text{eff}}/16)|$ target-deviation term. The query head is disabled, and uniformly explored actions are excluded from on-policy gradient updates. At termination, the pending action is credited without sampling an unused next action.

Each final EMA model generates 50,000 uint8 images with DDIM [35], 100 sampling steps, $\eta = 0$, and batches of 128. Noise streams are matched within a seed. The primary endpoint is FID [36] against the complete original 50,000-image training distribution, using `torch-fidelity` 0.4.0. KID [37] and Inception score [38] are secondary. A final EMA denoising assessment uses all 10,000 official test images and fixed noise/timestep seed 541211. Per-image errors are retained.

Four FID contrasts are frozen: baseline minus pretrained reference and each regularized condition minus baseline. Negative differences favor the target. Paired-t and all 64 sign flips are reported with separate Holm correction across these four contrasts. The exact-resolution limitation in Section 6.8 applies. Development uses two different seeds, 2,000 updates, and 5,000 generated samples; its absolute FID values are not pooled with confirmation.

6.9.3 AudioCaps captioning

The audio study uses original AudioCaps annotations [39] and a documented available-recording subset from `OpenSound/AudioCaps`. All 473 hosted Parquet shards were checked against published digests. IDs, start times, and captions were matched to pinned original annotations. The retained counts are 44,466 training, 440 validation, and 866 test clips, versus original inventories of 49,838, 495, and 975 unique clips. Coverage is 89.22%, 88.89%, and 88.82%. Each retained validation/test clip has five references.

Preparation excludes recordings outside eight to twelve seconds, one test identity with conflicting decoded variants, and both members of a nonzero waveform duplicate spanning train and test. Sources, variants, and exclusion records are preserved. Four all-zero training waveforms remain flagged and retained. Checks found no cross-split video IDs or exact nonzero waveform duplicates, including after actual mono/resampling preprocessing. These checks do not establish acoustic authenticity, absence of near-duplicates, or equivalence to an official raw-audio distribution. The mirror's provenance and licensing limitations remain explicit; raw audio is not redistributed.

The frozen Audio Spectrogram Transformer (AST) [40] checkpoint `MIT/ast-finetuned-audioset-10-10-0.4593` supplies shared features. Channels are averaged, resampled to 16 kHz, and limited to the first 10.24 seconds for a single-window representation. Original waveforms are not rewritten. FP32 AST computation supplies 64 pooled patch tokens of dimension 768, stored in FP16 and reused read-only. The test features use the frozen backbone only. AudioSet [41] pretraining may include underlying recordings, so no absence of pretraining overlap is claimed.

The task model has a two-layer trainable audio-token encoder, an eight-token learned prefix, and trainable pretrained GPT-2, totaling 141,774,336 parameters. Baseline, fixed spectral DCR, spectral Lite, and spectral Full use matched accepted initial task weights, epoch permutations, per-step dropout seeds, and optimizer budgets. Seeds 1301–1306 start from original GPT-2 assets and new seed-specific encoder/prefix parameters, not pilot-final weights. The initialization recovery is documented in Appendix F.

AdamW uses encoder/prefix rate 10^{-4}, decoder rate 5×10^{-5}, decay 0.01, and norm-1 clipping. Parameters remain FP32 with BF16 autocast. Each epoch permutes 44,466 training pairs and uses 1,389 batches of 32, dropping the last 18 pairs. A new deterministic permutation is used every epoch. Each of the 24 runs performs ten epochs, or 13,890 updates.

For zero-based optimizer-step index $s = 0, \ldots, S-1$, with $S = 13{,}890$ and $w = 100$ warm-up steps, every parameter group's base rate $\eta_{g,0}$ is multiplied by

$$a(s) = \begin{cases} (s+1)/w, & 0 \le s < w, \\ \dfrac{1 + \cos\big(\pi p(s)\big)}{2}, & w \le s, \end{cases} \tag{27}$$

where $p(s) = \min\{1, (s-w)/(S-w)\}$ and $\eta_g(s) = \eta_{g,0} a(s)$. The schedule has terminal factor zero and no positive learning-rate floor. Its zero endpoint is $s = S$; the last executed update is $s = S - 1$, so its final applied factor is slightly above zero. The final scheduled checkpoint is evaluated without best-epoch selection or early stopping.

Spectral sample axes and rank. The prefix representation reshapes eight embeddings per clip into $32 \times 8 = 256$ vectors of dimension 768. The pooled encoder representation has 32 vectors of dimension 768. Thus their centered sample-rank bounds are 255 and 31, respectively. The configured 64-entry spectrum is a storage and truncation dimension, not a claim of 64 nonzero eigenvalues for both representations. The encoder Gram spectrum is zero-padded to 64 entries as in Equation (17). A rank-16 template is feasible for both full training batches. The EMA-covariance monitor is distinct from either current-batch spectrum and can have a larger effective rank.

Control operates every 50 steps and at the final step, yielding 278 adaptive events. The fixed coefficient is

0.01. Adaptive coefficients start at 0.01 within $[0, 0.05]$. Spectral Lite uses Equation (26) at both monitored layers, with step size 0.001, deadband 0.1, and target rank 16. Full retains the diffusion spectral-policy settings, replacing denoising feedback loss with caption loss on the audio holdout. Its terminal action is credited without an unused next action, and the acquisition head is disabled.

The first 64 validation IDs in frozen manifest order provide controller feedback; the remaining 376 provide development assessment. Neither is used for task-model gradient updates. Final confirmation checkpoints evaluate all 866 retained test recordings. Beam size is three, decoding batch size is 32, and at most 30 new tokens are generated.

The primary endpoint is CIDEr-D [42] from the retained `pycocoevalcap`/PTB path. BLEU-1–4 [43], ROUGE-L [44], METEOR [45], SPICE [46], SPIDEr [47], and pooled token NLL are secondary outcomes. SPIDEr combines CIDEr-D and SPICE on their recorded scales. NLL is weighted by actual next-token target counts. A separate development diagnostic mismatches audio and captions deterministically without training updates.

The three primary CIDEr-D contrasts compare fixed, Lite, and Full with baseline. Positive differences favor the regularized condition. Both paired-t and all 64 sign flips are retained with separate Holm corrections across these three comparisons and individual unadjusted intervals. Secondary outcomes do not replace the primary endpoint. The six paired seeds remain the experimental units; repeated evaluation of 866 clips does not create extra training replications.

6.10 Completed checkpoint and raw-source follow-ups

Frozen-checkpoint robustness. The separately fixed diagnostic protocol evaluates the final ResNet-50 entropy, tuned-fixed, and Lite checkpoints for seeds 62001–62010 on all 4,995 cleaned validation images. Standard AutoAttack retains APGD-CE, APGD-T, FAB-T, and Square at L_∞ radii $\{0.5, 1, 2, 4, 8\}/255$, using FP32 inputs/parameters without autocast, attack batches of 32 and committed chunks of 128. All 30 models and 150 suites are required regardless of outcomes. Clean-logit replay precedes attacks. Robust top-1 at 2/255 is primary: fixed minus entropy, Lite minus entropy, and Lite minus fixed form one three-contrast Holm family, separately for paired-t and exhaustive 1,024-sign tests. Intervals are individual unadjusted 95% paired-t intervals. Other radii, shared-clean-correct intersections, geometry, displacement, margins, and local margin-gradient proxies are descriptive. The seed, rather than an image or attack chunk, is the experimental unit.

Practical training paths. From the seed-62001 entropy checkpoint, six paths cross streamlined versus fully instrumented execution with entropy, fixed DCR, and Lite. Each has 20 warm-up and 200 measured steps, eight cached real batches, fresh optimizer state and five order-varied timing repetitions. Streamlined fixed execution removes statistics and feedback unused by its loss; Lite retains its prescribed controller and holdout operations. Instrumented paths all include monitoring. Timing includes measured training, transfers, and applicable feedback, but excludes cache setup, checkpoint I/O, acquisition, attacks, and exports. Modified model and optimizer states are discarded. This benchmark measures short-path overhead, not convergence, full-job runtime, or a new trained method.

Raw-source AudioCaps. The separate follow-up uses the original encoded WAV archive and pinned official memberships/captions. Before model outcomes, the fixed audit requires finite decoding and durations from 8 to 12 seconds and removes identified cross-split canonical-input/YouTube-ID conflicts. Of 51,308 audited records, 840 are excluded: 49,025 training, 487 validation and 956 test clips remain, with about 98% coverage in each official split and no retained all-zero clips. These checks improve the earlier available-subset coverage without authenticating every recording or excluding all acoustic near-duplicates. The benchmark has been studied before; AST pretraining may include underlying AudioCaps recordings.

Frozen AST features are newly extracted in FP32 from the first 10.24 seconds, pooled to 64 tokens of dimension 768 and cached in FP16. The retained captioner trains its two-layer audio encoder, eight-query prefix and GPT-2 decoder. Its rank-16 spectral DCR and controller settings remain distinct from supervised classification scatter DCR. Every method loads the same saved complete initial model for its seed, with paired data order/dropout and the same ten-epoch optimizer budget. Training uses batch size 32 and final scheduled checkpoints only. The first 64 lexicographically ordered validation IDs supply controller feedback; the remaining validation IDs supply development selection.

The fixed-strength candidates $\{0, 0.001, 0.003, 0.01, 0.03, 0.05\}$ each use three tuning seeds. Highest mean final-development CIDEr-D selects the strength (exact tie: smaller strength) before confirmation training; test captions and test NLL do not enter selection. Ten fresh paired seeds compare baseline, selected fixed DCR, Lite, and Full MACC on all retained official-test clips, with five references per clip and beam size three. All outcomes are retained. The five planned CIDEr-D contrasts are fixed minus baseline, Lite minus baseline, Full minus baseline, Lite minus fixed, and Full minus fixed. The primary procedure is the two-sided exhaustive 1,024-sign test with Holm correction across those five at familywise $\alpha = 0.05$. A separately corrected paired-t family is sensitivity analysis, never an alternative route to declaring success. Intervals are individual unadjusted 95% paired-t intervals. Other caption metrics, token-weighted NLL and timings are descriptive. The small-sample and sign-exchangeability assumptions in Section 6.8 still apply. None of these family-wise

Table 4: Content-filtered study: final top-1 accuracy (%), mean ± sample SD over ten paired seeds per backbone. All active-learning conditions use entropy acquisition and 75,954 annotations; the full-label reference uses 126,590. All share 39,100 post-initialization optimizer steps. Validation contains 4,995 images.

Treatment	ResNet-50	ResNet-18
Entropy	71.493 ± 0.542	71.540 ± 0.567
Entropy + Lite	73.393 ± 0.707	73.151 ± 0.557
Entropy + tuned fixed	73.546 ± 0.747	73.243 ± 0.410
Entropy + schedule	72.763 ± 0.808	72.871 ± 0.595
Entropy + contrastive	72.300 ± 0.777	72.282 ± 0.484
Full labels, matched steps	73.524 ± 0.606	74.819 ± 0.490

corrections controls error across the entire sequential research programme.

6.11 Execution environment and reproducibility

Experiments run sequentially on an IBM Cloud instance with one NVIDIA L40S GPU, PyTorch 2.6.0, torchvision 0.21.0, CUDA runtime 12.4, and Python 3.11.7. Datasets and checkpoints reside on persistent storage. Protocols, source hashes, splits, acquisitions, controller records, summaries, and attack-chunk digests are retained. Integrity checks establish consistency of recorded artifacts, not third-party replication of GPU execution.

7 Results and analysis

We first report the completed 206-run content-filtered study, separating end-to-end treatment effects from common-label component effects. The earlier 104-run control comparisons, coefficient dynamics, and duplicate sensitivity follow, then the six-seed acquisition, control, and adversarial studies. Spectral task results and computational measurements complete the evaluation. Each study retains its own experimental units, inventory, and statistical family.

7.1 Content-filtered confirmation and backbone transfer

Selection on seeds 61001–61003 chooses fixed DCR strength 0.10, the decreasing schedule with peak 0.10, and contrastive strength 0.10. Their mean holdout accuracies are 77.795%, 77.162%, and 76.213%, respectively. The contrastive choice uses the cross-entropy tie-break. These are selection scores within finite grids, not independent performance estimates or globally optimal strengths. Full candidate results appear in Appendix C.

On the retained validation set, Lite improves on entropy alone by 1.900 points on ResNet-50 (individual 95% CI $[1.536, 2.264]$) and 1.612 points on ResNet-18 ($[1.247, 1.977]$). Every paired seed difference is positive. Both survive the 17-comparison Holm correction under paired-t and exact tests. Lite also exceeds the matched one-view contrastive-feature regularizer by 1.093 and 0.869 points, respectively; both pass both corrected procedures. This comparison supports the evaluated DCR treatment relative to that implementation, without establishing superiority over the original SupCon recipe.

Table 5: Common-label ResNet-50 replay. Mean ± sample SD final accuracy (%) over seeds 62001–62010. Every condition shares entropy's acquired labels, batches, and augmentation streams. The entropy row is the donor condition from Table 4, not a new run.

Training intervention	Top-1 (%)
No representation penalty	71.493 ± 0.542
Ratio only, fixed 0.10	73.614 ± 0.703
Separation only, weight 0.01	71.421 ± 0.587
Ratio + separation, fixed	73.481 ± 0.812
Ratio + separation, Lite	73.542 ± 0.785
Ratio + separation, schedule	73.037 ± 0.828

The tuned fixed treatment has slightly higher mean accuracy than Lite on both backbones (Table 4). Lite minus fixed is -0.152 points on ResNet-50, CI $[-0.593, +0.289]$, and -0.092 points on ResNet-18, CI $[-0.403, +0.219]$. The comparisons establish neither superiority, equivalence, nor non-inferiority. Lite minus the selected schedule is $+0.631$ and $+0.280$ points; neither survives the full correction under either procedure. The positive treatment effect over entropy therefore persists after prospective exact-content filtering, while an independent feedback advantage remains unestablished.

Matched-step full-label reference. Lite uses 75,954 annotations versus 126,590 for the full-label reference. Its accuracy difference is -0.130 points on ResNet-50, CI $[-0.531, +0.270]$, an inconclusive comparison. On ResNet-18 the difference is -1.668 points, CI $[-2.248, -1.087]$, favoring full labels under both adjusted tests. Equal optimizer-step budgets do not imply equally optimized methods. In particular, the nonsignificant ResNet-50 difference does not establish equal-performance label savings.

7.2 Attribution with common acquired labels

Table 5 holds entropy's annotation sequence fixed across the six training treatments. The ratio-only intervention improves final accuracy by 2.120 points over unregularized training, CI $[1.783, 2.457]$, and fixed-both exceeds separation-only by 2.060 points, CI $[1.672, 2.448]$. Both survive both adjusted procedures. These contrasts identify a beneficial scatter-penalty intervention conditional on common labels and the selected strength.

Separation-only minus entropy is -0.072 points, CI $[-0.328, +0.183]$. Adding separation to the fixed ratio penalty changes accuracy by -0.132 points, CI $[-0.437, +0.173]$. Adding feedback to fixed-both changes accuracy by $+0.060$ points, CI $[-0.199, +0.320]$. The factorial interaction is -0.060 points, CI $[-0.524, +0.404]$. All are inconclusive; they

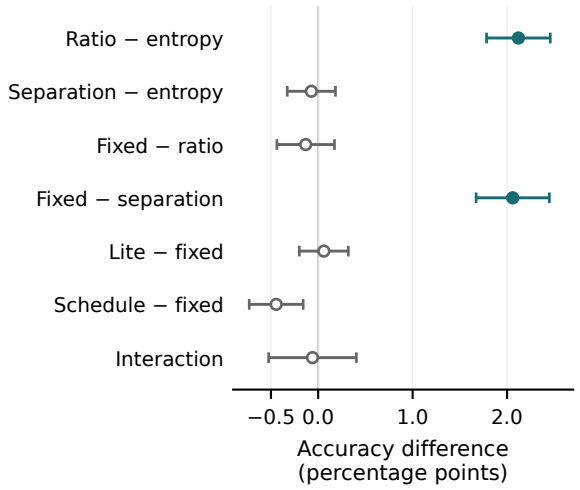


Figure 2: Common-label ResNet-50 effects in the 206-run study. Points and bars show mean accuracy differences and individual, unadjusted 95% paired-t intervals over ten seeds. Filled teal points survive both 17-comparison Holm procedures; open gray points do not. Interval exclusion alone does not imply adjusted significance. Fixed denotes both penalties at fixed strength; interaction is fixed minus ratio minus separation plus entropy.

do not show that separation or feedback is ineffective or unnecessary. The replay schedule-minus-fixed difference of −0.444 points also fails the family correction. Figure 2 displays these conditional effects; all 17 planned contrasts, including nonsignificant comparisons, are retained in Table 33.

The intervention evidence is more specific than the end-to-end comparison: acquisition cannot explain the replay difference because the acquired sets are shared. It does not establish mediation by the measured scatter statistic, isolate normalization, or assign all gains in freely evolving acquisition pipelines to one component. Reduced objectives were evaluated at strengths selected for the complete objective, and these component effects have not been repeated on a second dataset.

7.3 Earlier 104-run tuned-control and transfer study

The training-holdout search selects constant strength 0.10 and the linearly decreasing schedule with peak 0.10. Their mean tuning accuracies are 76.567% and 76.461%, respectively; these are selection scores rather than independent performance estimates. The complete candidate ranking is shown in Table 6, including the 0.03 constant used in the earlier comparisons. Selection occurs within the specified grid; it does not establish a globally optimal constant or schedule.

Table 7 reports the subsequent ten-seed evaluations. Lite improves final accuracy over entropy by 1.742 percentage points on ResNet-50 and 1.834 points on ResNet-18. Every paired seed difference is positive on each backbone. Both contrasts pass the planned six-comparison Holm correction under paired-t and exact sign-flip tests (Table 8). The same feedback settings therefore retain a classification benefit on the second architecture without retuning.

Table 6: All candidate controls on the three ResNet-50 tuning seeds 41001–41003. Values are final training-holdout mean ± sample SD; each seed has 633 holdout images. Selection maximizes mean accuracy within each family, with cross-entropy and identifier tie-breaks. These are selection scores, not independent performance estimates. Official validation is not evaluated during tuning.

Control	Holdout top-1 (%)	Holdout cross-entropy	Selected
Constant 0.003	75.408 ± 2.519	0.976 ± 0.066	—
Constant 0.010	75.619 ± 1.003	0.955 ± 0.050	—
Constant 0.030	75.197 ± 0.790	0.963 ± 0.056	—
Constant 0.100	76.567 ± 1.121	0.886 ± 0.058	Yes
Decreasing, peak 0.03	74.829 ± 0.638	0.969 ± 0.047	—
Decreasing, peak 0.10	76.461 ± 0.632	0.948 ± 0.045	Yes
Increasing, peak 0.03	75.566 ± 1.227	0.980 ± 0.052	—
Increasing, peak 0.10	75.935 ± 1.164	0.936 ± 0.023	—

The tuned fixed condition has the highest final mean on both backbones. Lite minus fixed is −0.188 points on ResNet-50 and −0.334 on ResNet-18, with individual intervals spanning zero. Neither comparison establishes superiority in either direction, equivalence, or non-inferiority. Lite exceeds the selected schedule by 0.398 and 0.324 points; only the ResNet-18 comparison survives the family correction. The final-budget gain over entropy is therefore supported, but these paired comparisons do not establish an advantage over the independently tuned constant.

Figure 3 shows the measured trajectories and primary effect intervals. Benefits are budget dependent: on ResNet-50, Lite trails entropy by 0.362 and 0.108 points at 20% and 30% annotations, then leads by 0.850, 1.552 and 1.742 points at 40%, 50% and 60%. Corresponding ResNet-18 gains are 0.150, 0.882, 1.750, 2.040 and 1.834 points. These intermediate outcomes are descriptive; they do not establish a universal label-saving rate.

7.4 Feedback dynamics and representation outcomes

The recorded coefficients follow the prescribed recurrence throughout training. Step-weighted mean layer3/layer4 strengths are 0.0450/0.0866 on ResNet-50 and 0.0537/0.0789 on ResNet-18. Their layer dependence and within-phase changes are visible in Figure 4. This verifies that feedback acts on the measured signal; it does not isolate the causal value of adaptation. Subsequent acquired sets also depend on the resulting model predictions.

Final mean log scatter ratios for entropy, Lite, fixed and schedule are 1.387, 0.471, 0.322 and 0.991 on ResNet-50, and 1.162, 0.528, 0.406 and 0.802 on ResNet-18. These descriptive changes accompany lower cross-entropy and ECE under the regularized treatments, but do not establish geometry as a mediator or a robustness certificate. The fixed separation safeguard is active on only 5 layer3 and 102 layer4 minibatch records across the regularized post-initialization traces; its independent contribution is not identified by this experiment.

Table 7: Earlier 104-run control-study outcomes at 60% annotations on ImageNet-100. Values are mean ± sample SD over ten paired seeds per backbone. Every treatment uses entropy acquisition. Fixed and scheduled strengths are selected using separate ResNet-50 tuning seeds; all settings are frozen for confirmation and transfer.

Backbone	Treatment	Clean top-1 (%)	Cross-entropy	ECE
ResNet-50	Entropy	71.678 ± 0.808	1.137 ± 0.024	0.079 ± 0.007
ResNet-50	Entropy + MACC-Lite	73.420 ± 0.512	1.020 ± 0.014	0.058 ± 0.006
ResNet-50	Entropy + tuned fixed	73.608 ± 0.474	1.010 ± 0.019	0.054 ± 0.006
ResNet-50	Entropy + tuned schedule	73.022 ± 0.371	1.052 ± 0.019	0.064 ± 0.004
ResNet-18	Entropy	71.252 ± 0.397	1.197 ± 0.022	0.095 ± 0.004
ResNet-18	Entropy + MACC-Lite	73.086 ± 0.397	1.091 ± 0.018	0.082 ± 0.003
ResNet-18	Entropy + tuned fixed	73.420 ± 0.673	1.086 ± 0.024	0.079 ± 0.006
ResNet-18	Entropy + tuned schedule	72.762 ± 0.483	1.113 ± 0.025	0.084 ± 0.004

Table 8: All six predeclared final-top-1 contrasts in the control study. Effects and individual, unadjusted 95% paired-t intervals are in percentage points. Each test column is separately Holm-adjusted across the full six-comparison family.

Backbone	Contrast	Difference	95% CI	Adjusted t p	Adjusted exact p
ResNet-50	Lite − entropy	$+1.742$	$[+1.354, +2.130]$	1.57×10^{-5}	0.01172
ResNet-50	Lite − tuned fixed	-0.188	$[-0.529, +0.153]$	0.2436	0.25781
ResNet-50	Lite − tuned schedule	$+0.398$	$[+0.073, +0.723]$	0.065	0.08203
ResNet-18	Lite − entropy	$+1.834$	$[+1.558, +2.110]$	6.74×10^{-7}	0.01172
ResNet-18	Lite − tuned fixed	-0.334	$[-0.741, +0.073]$	0.1929	0.20312
ResNet-18	Lite − tuned schedule	$+0.324$	$[+0.155, +0.493]$	0.007541	0.03125

7.5 Content-duplicate sensitivity

The dataset manifest contains five exact SHA-256 matches between training and validation images, all with conflicting class labels, and 51 duplicate-content pairs within training, 43 with conflicting labels. Seven of the 23 new seed holdouts contain one or two examples duplicated elsewhere in training; their counterpart enters the final gradient-training set in 24 method runs. None of the three tuning holdouts is affected. The exclusions by index operate as specified, but do not eliminate overlap by content.

In a post-hoc sensitivity analysis, the same five validation indices are excluded for every method, retaining 4,995 examples. Lite-minus-entropy effects become 1.742 and 1.836 points on ResNet-50 and ResNet-18, while Lite-minus-fixed becomes -0.188 and -0.328 points. All six adjusted test decisions remain unchanged; excluding duplicate-content holdout items also leaves the selected controls unchanged. The planned 5,000-image outcomes remain primary. Evaluation-time exclusion does not undo possible training influence, and hash equality does not diagnose the source of the conflicting labels. Appendix B provides the sensitivity contrasts and seed-level endpoints.

7.6 Earlier six-seed classification comparisons

The earlier studies complement the tuned-control comparison by testing acquisition rules and learned control. Table 9 consolidates their 12 conditions. Entropy acquisition with aligned DCR and MACC-Lite has the highest mean within these six-seed studies, $73.360 \pm 0.208\%$ clean accuracy, compared with $71.463 \pm 0.320\%$ for entropy alone and $71.747 \pm 0.626\%$ for UHerding-Margin. The later entropy-plus-fixed condition reaches $72.263 \pm 0.381\%$. The table distinguishes the initial comparison from the three follow-up conditions; their inferential families remain separate. Within the initial nine-condition comparison, geometric acquisition with aligned MACC-Lite attains $72.120 \pm 0.490\%$. Adding fixed aligned DCR to geometric acquisition gives $+0.850$ percentage points, and replacing fixed strength with Lite gives another $+1.080$ points. Both component differences are positive in all six paired seeds.

Figure 5 shows the available per-phase clean-accuracy measurements, separating acquisition rules from aligned regularization. The budget trajectories are descriptive; the final 60% endpoint remains the basis of the declared clean-accuracy comparisons. The entropy-plus-Lite and entropy-plus-fixed trajectories are shown separately in Figure 6, together with paired gains over entropy. All displayed points are measured phase endpoints.

The two component contrasts pass the planned adjusted paired-t threshold. Geometric acquisition with Lite versus unregularized entropy does not: its raw paired-t value is 0.0466, but its adjusted value is 0.0931. This contrast differs from the entropy-plus-Lite treatment evaluated in Section 7.3. The individual interval narrowly excludes zero but is not a simultaneous interval. Table 10 preserves all original effects and intervals. The nonrejection of the adjusted exact tests must be interpreted with their design-imposed resolution limit, not as an independent demonstration that the effects are absent.

Acquisition alone improves over random by $+0.897$ points, while random plus fixed aligned DCR improves by $+0.820$. These comparisons show favorable directions for both components within the declared regime. Acquisition alone remains below entropy in mean ac-

Table 9: The 12 earlier six-seed ImageNet-100 conditions at 60% annotations. Values are mean ± sample SD over six paired seeds; clean and standard AutoAttack accuracy use all 5,000 validation images. Study I denotes the initial nine-condition comparison, A the later acquisition-rule ablation, B the later UHerding comparison, and C the later entropy-plus-fixed control. The table is descriptive: study-specific inferential families remain separate. Bold identifies the highest observed mean clean accuracy, not a universal superiority claim.

Condition	Study	Clean top-1 (%)	AutoAttack 8/255 (%)	Cross-entropy	ECE
Random	I	69.293 ± 0.334	0.013 ± 0.010	1.272712 ± 0.021145	0.094945 ± 0.002917
Entropy	I	71.463 ± 0.320	0.010 ± 0.011	1.150591 ± 0.020293	0.079698 ± 0.003849
BADGE	I	71.110 ± 0.467	0.003 ± 0.008	1.149581 ± 0.007022	0.077088 ± 0.004058
Noise Stability	I	71.023 ± 0.249	0.003 ± 0.008	1.162821 ± 0.006167	0.078703 ± 0.002375
CÆNL acquisition only	I	70.190 ± 0.240	0.013 ± 0.010	1.210545 ± 0.010404	0.085552 ± 0.004501
Random + aligned fixed DCR	I	70.113 ± 0.452	0.017 ± 0.008	1.217939 ± 0.021315	0.085964 ± 0.006228
CÆNL + aligned fixed DCR	I	71.040 ± 0.427	0.017 ± 0.008	1.155351 ± 0.015230	0.078428 ± 0.003667
CÆNL + aligned MACC-Lite	I	72.120 ± 0.490	0.007 ± 0.010	1.078885 ± 0.008014	0.062948 ± 0.004965
CÆNL + aligned Full MACC	I	70.803 ± 0.318	0.007 ± 0.010	1.171322 ± 0.023746	0.081639 ± 0.004671
Entropy + aligned MACC-Lite	A	$\mathbf{73.360 \pm 0.208}$	0.003 ± 0.008	1.028967 ± 0.010101	0.055426 ± 0.001256
UHerding-Margin	B	71.747 ± 0.626	0.007 ± 0.010	1.127922 ± 0.021039	0.073977 ± 0.003294
Entropy + aligned fixed DCR	C	72.263 ± 0.381	0.000 ± 0.000	1.092014 ± 0.013386	0.071037 ± 0.004253

curacy. The acquisition-rule ablation shows that geometric acquisition is not necessary for the favorable entropy-plus-Lite treatment. No equal-performance full-supervision comparison establishes a 40% annotation saving.

7.7 Acquisition-rule ablation under aligned MACC-Lite

The acquisition-rule condition replaces geometric acquisition with entropy while retaining the aligned objective, Lite settings, six shared initializations, annotation and feedback splits, targets, and five 20-epoch phases. Its first acquisition is checked against the original entropy selection. Thirty additional phases and full-set final attacks are completed. These are paired acquisition-rule runs, not new independent seeds for the original matrix.

Entropy with aligned DCR and Lite reaches $73.360 \pm 0.208\%$ and exceeds unregularized entropy and geometry-plus-Lite in every seed. Table 11 gives the paired comparisons; Appendix A.1 provides the seed-level outcomes (Table 27).

The gain over entropy is +1.897 percentage points, with individual 95% paired-t interval $[+1.625, +2.169]$. The gain over geometry plus Lite is +1.240 points, with interval $[+0.707, +1.773]$. Both comparisons pass the acquisition-rule adjusted paired-t analysis. Their raw exact value is 0.03125 and adjusted value 0.0625, consistent with the resolution limit. The evidence supports the *combined aligned-loss-and-Lite treatment within entropy acquisition*. That acquisition-rule study did not include a fixed-strength entropy control. The later matched follow-up in Section 7.8 provides that comparison while preserving the acquisition-rule study's original two-contrast family.

One image survives the attack for seed 802 and none for the other seeds, giving mean robust accuracy 0.0033%. This descriptive result does not change the robustness conclusion or the original robustness family.

7.8 Matched fixed-strength control within entropy acquisition

The third follow-up keeps entropy acquisition and replaces Lite with the fixed aligned coefficient $\lambda_l = 0.03$ at layer3 and layer4. The separation penalty remains $\beta_B = 0.01$ with the original per-seed targets and floors. All six initial model and optimizer states, label/holdout splits, training settings, and the five-round schedule are reused. The first acquisition matches original entropy exactly; later acquired sets are allowed to change with the trained model. This tests the coefficient-control treatment within the active-learning pipeline, including its subsequent effect on selection.

Entropy plus fixed DCR reaches $72.263 \pm 0.381\%$, improving over entropy by +0.800 percentage points with individual 95% interval $[+0.498, +1.102]$. Lite improves over fixed by +1.097 points with interval $[+0.714, +1.479]$. Both differences are positive in all six seeds. Within this six-seed two-contrast family, both Holm-adjusted paired-t values are 0.001442. Both raw exact values are 0.03125 and adjusted values are 0.0625. Table 12 reports the full statistics, and Table 27 gives the paired outcomes. This follow-up favors Lite over the prespecified 0.03 coefficient under the paired-t analysis. The independently tuned constant and schedule are evaluated with new seeds in Section 7.3; those comparisons address different controls.

Table 10: All seven original clean-accuracy contrasts in percentage points. Intervals are individual unadjusted 95% paired-t intervals. Each p-value column is separately Holm-adjusted across seven contrasts. The exact tests have the resolution limit described in Section 6.8.

Target minus reference	Difference (pp)	95% CI (pp)	Adjusted t p	Adjusted exact p
Aligned fixed DCR − acquisition only	+0.850	[+0.430, +1.270]	0.0208	0.2188
Aligned MACC-Lite − aligned fixed DCR	+1.080	[+0.447, +1.713]	0.0309	0.2188
Aligned Full MACC − aligned fixed DCR	−0.237	[−0.714, +0.241]	0.2589	0.2813
Aligned MACC-Lite − entropy	+0.657	[+0.015, +1.299]	0.0931	0.2188
Aligned Full MACC − entropy	−0.660	[−0.869, −0.451]	0.0032	0.2188
Acquisition only − random	+0.897	[+0.389, +1.405]	0.0309	0.2188
Random + aligned fixed DCR − random	+0.820	[+0.297, +1.343]	0.0309	0.2188

Table 11: Two acquisition-rule clean-accuracy contrasts. Positive values favor entropy with aligned DCR and Lite. Intervals are individual unadjusted paired-t intervals; each test family is Holm-adjusted across these two contrasts. This family was specified after inspection of the initial comparison.

Reference condition	Gain (pp)	95% CI (pp)	Adjusted t p	Adjusted exact p
Entropy without aligned regularization	+1.897	[+1.625, +2.169]	0.0000198	0.0625
Geometry acquisition + aligned MACC-Lite	+1.240	[+0.707, +1.773]	0.001875	0.0625

Table 12: Six-seed fixed-control clean-accuracy contrasts (percentage points). Individual 95% intervals are unadjusted; Holm is separate for the two paired-t and two exact tests.

Contrast	Difference	95% interval
Fixed − entropy	+0.800	[+0.498, +1.102]
Lite − fixed	+1.097	[+0.714, +1.479]
Test	**Fixed − entropy**	**Lite − fixed**
Raw t p	0.001047	0.000721
Adjusted t p	0.001442	0.001442
Exact raw p	0.03125	0.03125
Exact adjusted p	0.0625	0.0625

Figure 6 includes the Lite trajectory and the matched fixed trajectory. At 20% and 30% annotations, fixed DCR has slightly higher mean accuracy than Lite; the ordering reverses at later budgets. Thus the final-budget advantage should not be described as uniform throughout acquisition. Fixed DCR has zero robust-correct images in every seed under standard AutoAttack 8/255 on all 5,000 validation images. This fixed-strength condition also has negligible adversarial accuracy under the specified threat model.

7.9 Comparison with UHerding-Margin

Under the matched protocol in Section 6.7, UHerding-Margin attains mean clean accuracy $71.747 \pm 0.626\%$, versus $71.463 \pm 0.320\%$ for entropy and $73.360 \pm 0.208\%$ for entropy with aligned DCR and Lite. UHerding minus entropy is +0.283 points with interval $[-0.430, +0.997]$, adjusted paired-t value 0.3540, and adjusted exact value 0.3125. This does not establish superiority or equivalence.

Entropy with the aligned Lite treatment exceeds the UHerding adaptation in every seed. The mean difference is +1.613 points, with interval $[+0.957, +2.270]$. The adjusted paired-t value is 0.002923; the adjusted exact value is 0.0625. This compares complete pipelines and does not establish superiority of entropy alone, or of the combined treatment over the authors' unmodified UHerding benchmark implementation.

UHerding retains one correct adversarial image for each of seeds 802 and 805, and none otherwise: mean robust accuracy 0.0067%. The comparison adds an executed recent mechanism while retaining its representation and selection-domain adaptations.

7.10 Earlier representation and learned-control results

Within the original nine-condition matrix, geometric acquisition with aligned Lite has the lowest mean log trace ratio, cross-entropy, and ECE. Table 14 reports descriptive geometry separately from task performance. These statistics are not a causal mediation analysis, and the all-class assessment ratio does not establish attainment of the 16-class anchor target.

Full MACC executes 370, 368, 368, 374, 365, and 373 learned-policy updates for seeds 801–806. Its mean accuracy is 0.237 points below aligned fixed DCR, with an interval spanning zero, and 0.660 points below entropy. The latter difference passes the adjusted paired-t test. The learned controller is active, but its additional capacity does not establish benefit under the present settings. Feedback horizon, changing acquisition phases, and imperfect credit assignment are possible explanations rather than identified causes.

Short development continuations found Lite coefficients below the 0.10 cap. That observation does not substitute for a long-run stability analysis. Proposition 4 bounds coefficients, not the network's geometric state or task error.

7.11 Adversarial evaluation

Across the earlier six-seed classification studies, final checkpoints retain zero or one correct adversarial prediction out of 5,000. Mean robust accuracy is approximately 0.000%–0.017%. No planned robustness com-

Table 13: The UHerding comparison's two clean-accuracy contrasts, in percentage points. Intervals are individual unadjusted 95% paired-t intervals; each test column is Holm-adjusted across these two comparisons only.

Target minus reference	Difference	95% CI	Adjusted t p	Adjusted exact p
UHerding-Margin − entropy	+0.283	[−0.430, +0.997]	0.3540	0.3125
Entropy + aligned Lite − UHerding-Margin	+1.613	[+0.957, +2.270]	0.002923	0.0625

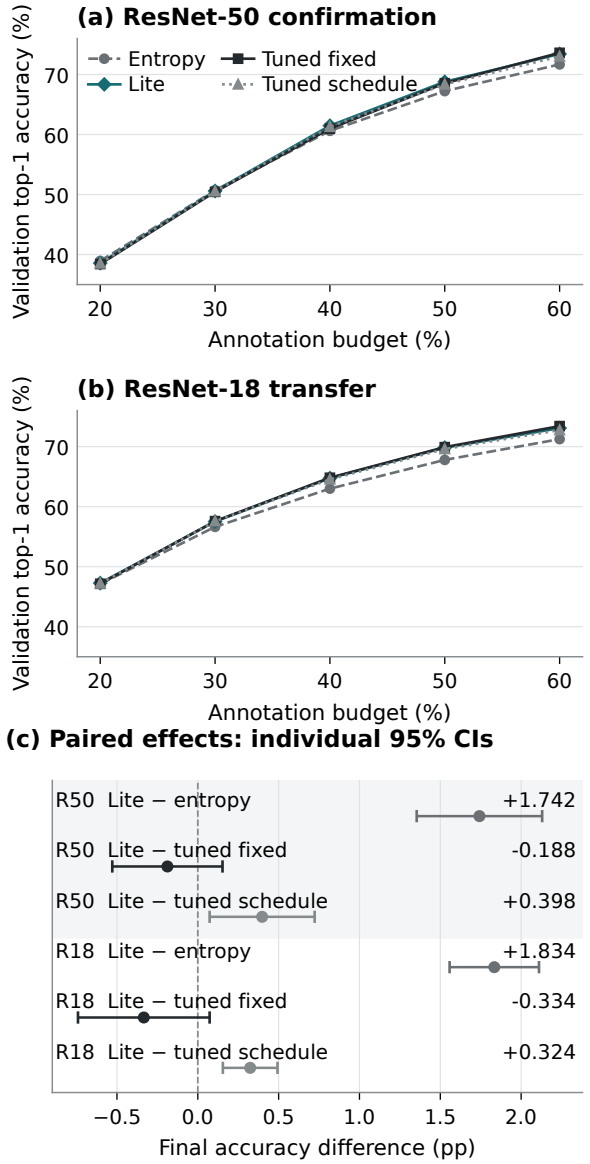


Figure 3: Earlier 104-run control comparison. Panels (a) and (b) show mean official validation accuracy over ten seeds at measured acquisition budgets. Panel (c) shows all six final Lite-minus-control effects and individual 95% paired-t intervals. The zero line denotes no difference; interval exclusion does not by itself imply multiplicity-adjusted significance. The test results are given in Table 8.

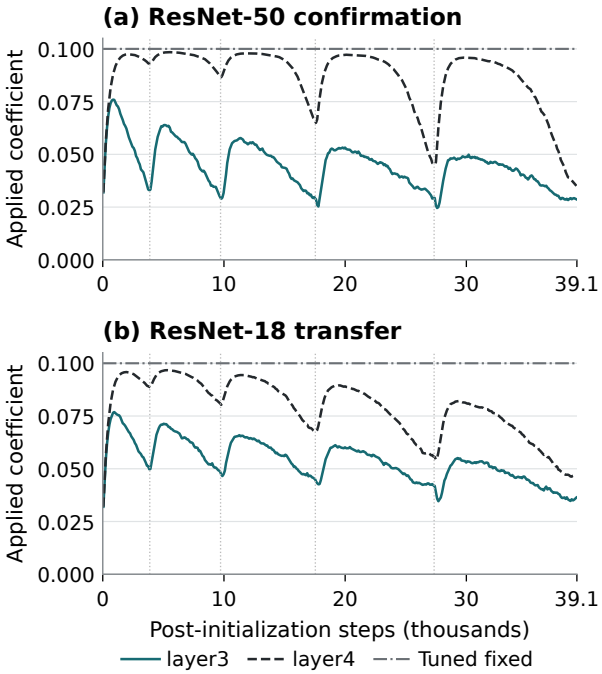


Figure 4: Applied Lite strengths from native training traces. Each curve is the ten-seed mean coefficient in successive windows of at most 100 steps within each phase; separate line styles identify layer3 and layer4. Vertical separators mark acquisition-round boundaries, and the horizontal reference is the tuned constant strength 0.10. Coefficients carry state across phases while responding to their respective moving geometry residuals.

Table 14: Final all-class normalized geometry, mean ± seed SD. The statistic averages layer3 and layer4 $\log((W+\epsilon)/(B+\epsilon))$.

Condition	Mean log trace ratio
Random	1.3881 ± 0.0197
Entropy	1.3916 ± 0.0170
Acquisition only	1.3864 ± 0.0151
Random + aligned fixed	1.0607 ± 0.0218
CÆNL + aligned fixed	0.9767 ± 0.0200
CÆNL + aligned Lite	0.4937 ± 0.0390
CÆNL + aligned Full	1.0825 ± 0.0982
Entropy + aligned fixed	1.0027 ± 0.0125
Entropy + aligned Lite	0.4725 ± 0.0434

parison establishes an improvement. Table 29 reports all seven original paired contrasts and their uncertainty, with the original test families retained. These near-zero empirical outcomes are not a population certificate, and repeated predictions on the same images do not become independent observations when pooled across seeds.

The original spectral pilot also had zero robustness at 8/255; extending its selected-checkpoint diagnostic from 1,000 to all 5,000 images reproduced that result. Smaller-radius post-hoc diagnostics did not establish an overall advantage over random acquisition. The final study retains the original attack radius. Improved clean accuracy and supervised geometry therefore coexist with negligible adversarial accuracy under this threat model.

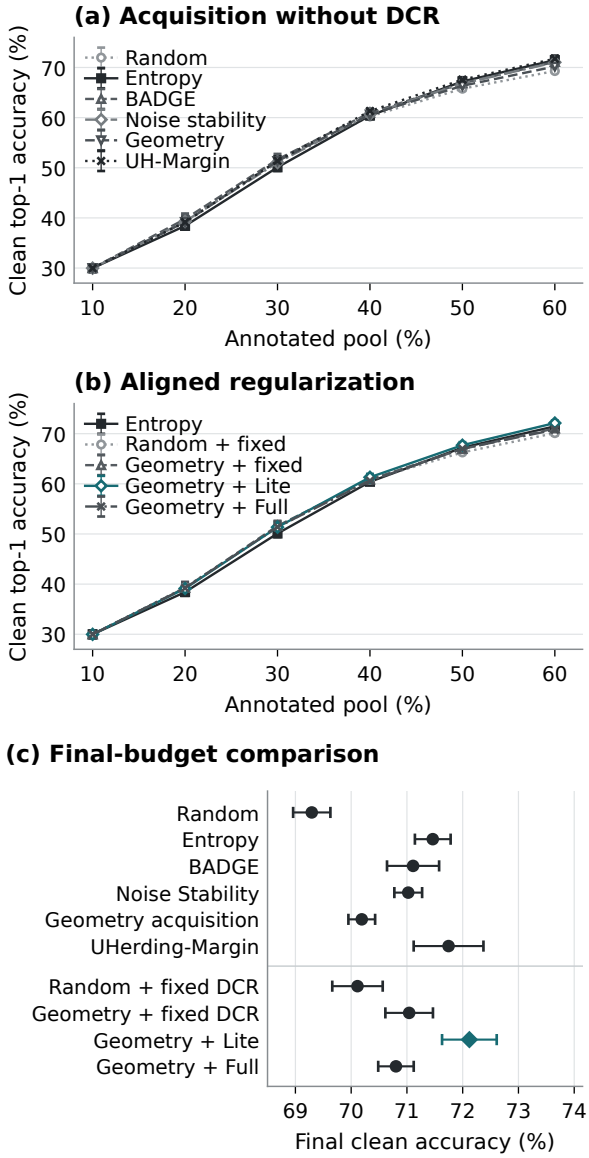


Figure 5: Acquisition and aligned regularization under the common ImageNet-100 protocol. Panels (a) and (b) show measured accuracy trajectories on identical axes; panel (c) enlarges the final 60% endpoint for the original nine conditions and the UHerding-Margin follow-up. Markers and bars represent six-seed means and sample SDs, not confidence intervals. All trajectories start from the same 10% checkpoint. Fixed, Lite, and Full denote aligned DCR conditions; entropy regularization follow-ups appear in Figure 6. Connecting lines guide the eye, and no new area-under-curve test or inferential family is introduced. Geometry and UH-Margin abbreviate geometry acquisition and UHerding-Margin.

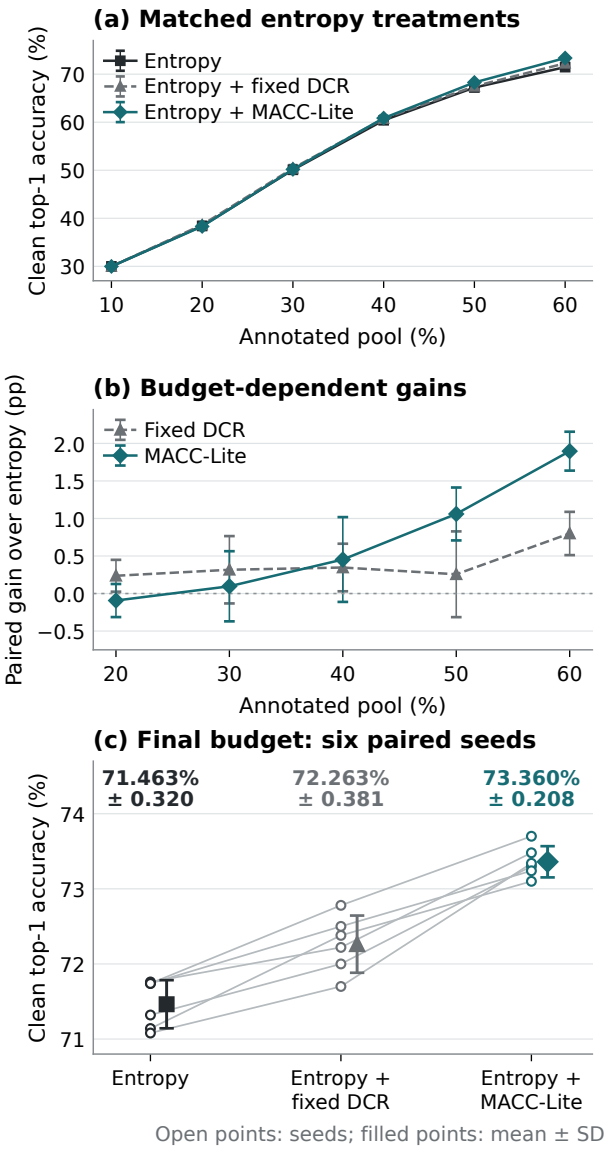


Figure 6: Matched entropy treatments across annotation budgets and seeds. (a) Mean clean accuracy and sample SD across six seeds. (b) Mean and sample SD of within-seed differences from entropy; the common 10% initialization has zero gain and is omitted. (c) Individual final-budget accuracies, with lines joining the same seed across conditions; adjacent filled markers and bars show mean ± sample SD. Overlapping seed values are retained. These bars are not confidence intervals. Fixed DCR uses $\lambda_l = 0.03$ and the same separation penalty as Lite. Lite is not uniformly better at early budgets, but exceeds both comparators in all six final-budget seeds. Only the final 60% comparisons enter the declared follow-up tests. All points derive from recorded measurements.

Completed content-filtered checkpoint diagnostics. All 150 specified suites are complete. Table 15 retains every radius; Figure 7 displays the same means. At the primary 2/255 radius, Lite minus entropy is -0.0100 percentage points with interval $[-0.0380, +0.0180]$; fixed minus entropy is $+0.0080$ with $[-0.0364, +0.0524]$; Lite minus fixed is -0.0180 with $[-0.0553, +0.0192]$. All three adjusted paired-t values are 0.9070 and all three adjusted exact values are 1. These comparisons establish no robust advantage or equivalence. At the smallest radius, mean robust accuracy is 22.300% for entropy, 22.280% for fixed and 22.194% for Lite; thus the result cannot be attributed solely to reporting an 8/255 floor. These smaller-radius observations remain descriptive.

Conditioning on the same clean-correct intersection for each seed gives 15 additional descriptive contrasts (three at each radius); every individual interval spans zero. This conditioning does not replace the all-validation primary endpoint. Clean scatter is lower under fixed and Lite, but AutoAttack returned-

Table 15: Completed frozen-checkpoint AutoAttack: mean robust top-1 accuracy (%) across ten paired seeds on all 4,995 cleaned validation images. The primary radius is 2/255; other radii are descriptive.

Radius	Entropy	Fixed	Lite
0.5/255	22.300	22.280	22.194
1/255	4.040	3.982	3.956
2/255	0.312	0.320	0.302
4/255	0.058	0.066	0.060
8/255	0.014	0.008	0.010

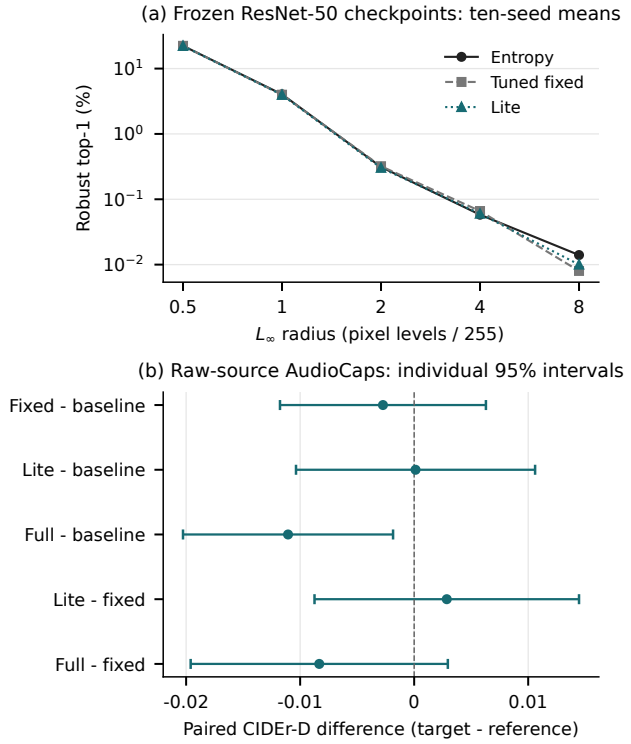


Figure 7: Completed diagnostics and raw-source audio follow-up. (a) All five specified attack radii; logarithmic axes display the low robust accuracies without omitting the smallest-radius measurements. Lines join descriptive means, not confidence bands. (b) All five planned audio effects with individual unadjusted paired-t intervals. Full minus baseline excludes zero individually, but none of the five primary Holm-adjusted exact tests rejects. No favorable radius, seed, or contrast is selected for this figure.

output geometry mixes unchanged clean errors and survivors with successful perturbations retained from different component attacks. It is not a uniform perturbation treatment. Outcome-conditioned margins/displacements and local gradient proxies therefore describe sensitivity without establishing mediation, certification, or the cause of the geometry–robustness dissociation. Appendix E retains all primary comparisons and the complete radius-by-seed numerical release.

7.12 Language modeling on C4

The following three task studies evaluate the spectral formulation in Section 4.5, with protocols given in Section 6.9. For C4, Table 16 reports the unused local validation shard, and Table 17 gives all three NLL contrasts against baseline. Fixed spectral DCR lowers NLL in every paired seed, with mean difference -0.00029250 nats/token and individual interval $[-0.00032813, -0.00025688]$. Its raw paired-$t$ value is 4.42×10^{-6} and its archived within-endpoint Holm-adjusted value is 8.85×10^{-6}. The raw exact value is 0.03125; applying Holm correction within the same three-contrast family gives 0.09375.

The effect is close to the original-shard difference of -0.00028734. The second shard checks sensitivity to local data reuse but adds no independent training seeds. The practical change is approximately 0.0086% in NLL and 0.0292% in perplexity. Fixed DCR does not establish a token-accuracy or pooled-ECE benefit. Spectral Lite increases NLL in all six seeds by a mean of 0.00194729 nats/token, with interval $[0.00187898, 0.00201560]$. Full is inconclusive relative to baseline: $+0.00023624$, interval $[-0.00018778, +0.00066025]$, raw paired-$t$ value 0.2115. This is not an equivalence result. The secondary comparison outputs are retained in Appendix D.

The language controller records 488 actions per seed. Lite spends much of training at the 0.08 coefficient cap; Full records 428–445 learned-policy updates. The adaptive results therefore do not reflect an inactive controller. Effective ranks increase under regularization, so the mechanism is spectral shaping rather than uniformly stronger low-rank collapse. The small fixed-DCR NLL result does not establish broader calibration, memorization, or cross-modal gains.

7.13 CIFAR-10 diffusion

Unregularized fine-tuning lowers mean FID from 10.512 to 8.057. Fixed spectral DCR, Lite, and Full instead increase FID relative to that baseline by 0.156, 0.415, and 0.154. Each increase occurs in every seed. Improvement over the unchanged pretrained model therefore does not establish added value from DCR.

All four contrasts pass their adjusted paired-t tests. Each raw exact value is 0.03125, giving adjusted value 0.125. The consistent adverse direction is retained with the small-sample assumptions and resolution qualification. It is not a claim that every diffusion regularizer is harmful.

KID is also higher in mean under every regularized condition. Slightly higher Inception scores remain secondary and do not replace FID. Mean official-test EMA denoising MSE is 0.02918539 for baseline, 0.02918606 for fixed DCR, 0.02918634 for Lite, and 0.02918632 for Full, compared with 0.02935395 for the reference. These small descriptive differences are not formal equivalence or significant secondary effects.

Full executes 188, 186, 189, 185, 183, and 186 learned updates, with 200 controller events per seed. Lite also executes 200 events and ends at the 0.05 cap. The final value does not specify its occupancy over training. Final monitored ranks are approximately 23–24 for baseline, 49 for fixed DCR, 38–39 for Lite, and 42–60 for Full. Active control and representation changes do not establish target attainment or better generation.

The two-seed pilot favored fixed DCR and Full relative to its baseline. That direction does not persist in

Table 16: Unused-local-shard C4 evaluation of the same 24 final GPT-2 checkpoints. Values are mean ± sample SD over six training seeds, with 1,999,872 next-token predictions per checkpoint. The formulations are spectral, not aligned classification DCR.

Condition	NLL (nats/token)	Perplexity	Pooled token ECE	Token accuracy
Baseline	3.384212 ± 0.000239	29.49473 ± 0.00705	0.012324 ± 0.000426	0.380067 ± 0.000123
Fixed spectral DCR	3.383919 ± 0.000253	29.48611 ± 0.00746	0.012216 ± 0.000438	0.380029 ± 0.000083
Spectral MACC-Lite	3.386159 ± 0.000220	29.55222 ± 0.00651	0.013131 ± 0.000396	0.379920 ± 0.000098
Spectral Full MACC	3.384448 ± 0.000462	29.50170 ± 0.01362	0.012438 ± 0.000584	0.380040 ± 0.000111

Table 17: All three unused-shard C4 NLL contrasts. Differences and individual 95% paired-t intervals are in 10^{-4} nats/token; negative differences favor regularization. Holm adjustment is within the three NLL contrasts, separately for each test. Both test families use the same three contrasts; intervals are unadjusted.

Condition minus baseline	Difference	95% CI	Raw t p	Holm t p	Raw exact p	Holm exact p
Fixed spectral DCR	−2.92505	$[-3.28129, -2.56881]$	4.4239×10^{-6}	8.8477×10^{-6}	0.03125	0.09375
Spectral MACC-Lite	+19.47289	$[+18.78979, +20.15599]$	8.9651×10^{-9}	2.6895×10^{-8}	0.03125	0.09375
Spectral Full MACC	+2.36236	$[-1.87778, +6.60250]$	0.21152	0.21152	0.21875	0.21875

Table 18: Diffusion confirmation with 50,000 generated images per condition and seed. Values are mean ± sample SD over six seed streams. FID is primary. Reference variation is sampling variation for one unchanged pretrained model.

Condition	FID ↓	KID $\times 10^3$ ↓	Inception score ↑
Pretrained reference	10.512 ± 0.100	9.153 ± 0.124	8.479 ± 0.039
Unregularized fine-tuning	$\mathbf{8.057 \pm 0.048}$	6.605 ± 0.084	8.692 ± 0.043
Fixed spectral DCR	8.213 ± 0.070	6.726 ± 0.105	8.724 ± 0.049
Spectral MACC-Lite	8.472 ± 0.100	6.939 ± 0.115	8.728 ± 0.055
Spectral Full MACC	8.211 ± 0.116	6.744 ± 0.088	8.729 ± 0.050

confirmation. Training horizon and generated-sample count both changed, so the present comparison does not isolate the reason for the reversal.

7.14 AudioCaps captioning

Fixed DCR has a slightly higher mean CIDEr-D than baseline; Lite and Full have lower means. Every primary interval spans zero, and both adjusted test families give $p = 1.0000$ for all three contrasts. This is an inconclusive comparison, not proof of equivalence, zero population effect, or universal ineffectiveness.

Seed-level CIDEr-D values are reported in Appendix A.2.

Raw paired-t values are 0.8971, 0.4673, and 0.5072; raw exact values are 0.84375, 0.50000, and 0.53125. Unlike a case resting only on the exact-test resolution limit, these raw results also provide no strong evidence of a difference. Mean test NLL is 2.026833 for baseline, 2.029867 for fixed, 2.033260 for Lite, and 2.030362 for Full. These remain descriptive secondary results.

Full records 260, 253, 260, 266, 259, and 254 learned-policy updates. Both adaptive conditions record 278 events. All 24 runs produce 866 nonempty captions. The mismatched-audio development diagnostic has higher NLL than correctly paired audio in every run. That establishes sensitivity to the supplied conditioning under this diagnostic, not added benefit from DCR or a comprehensive grounding guarantee.

The pilot evaluates 376 development clips after two epochs; confirmation evaluates 866 test clips after ten epochs. Their absolute scores cannot attribute a difference solely to training duration. The study evaluates audio–text captioning; audio retrieval, audio classification, other architectures, and full-original-benchmark comparability remain outside its scope.

7.15 Separate raw-source AudioCaps confirmation

All 18 tuning and 40 confirmation runs completed. Development selection chooses fixed spectral strength 0.001; every candidate score is retained in Appendix E. Table 22 reports the newly extracted raw-source results, without replacing the earlier 866-clip comparison. Baseline and Lite have similar descriptive CIDEr-D means (0.689205 and 0.689326); their paired difference is +0.000121, interval $[-0.010363, +0.010606]$. Lite minus tuned fixed is +0.002859, interval $[-0.008740, +0.014458]$. Both primary adjusted exact values are 1. Fixed minus baseline is −0.002738 with interval $[-0.011768, +0.006293]$ and adjusted exact value 1.

Full MACC is numerically lower than baseline: difference −0.011062, individual interval $[-0.020274, -0.001851]$, raw exact value 0.02734, but adjusted exact value 0.13672. Its separately adjusted paired-t value is 0.11868. Full minus fixed is −0.008324, interval $[-0.019602, +0.002953]$, adjusted exact value 0.5. Thus none of the five prespecified primary comparisons rejects. The adverse Full direction is retained, without upgrading a raw test or individual interval into an adjusted-significant harm claim. Neither adaptive benefit, fixed benefit, nor equivalence is established. Descriptive SPIDEr and token NLL do not replace CIDEr-D. The wider source coverage and independently selected fixed comparator address limitations of the earlier design; they do not make AudioCaps an untouched benchmark or establish cross-task transfer of

Table 19: All four diffusion FID contrasts. Positive differences indicate worse FID for the target. Intervals are individual unadjusted paired-t intervals. Each test family is Holm-adjusted over these four comparisons.

Target minus reference	Difference	95% interval	Adjusted t p	Adjusted exact p
Baseline − pretrained reference	−2.45468	$[-2.58346, -2.32589]$	2.6769×10^{-7}	0.125
Fixed DCR − baseline	+0.15601	$[+0.11357, +0.19845]$	0.0004481	0.125
MACC-Lite − baseline	+0.41472	$[+0.31931, +0.51012]$	0.0003004	0.125
Full MACC − baseline	+0.15431	$[+0.06205, +0.24658]$	0.0077204	0.125

Table 20: AudioCaps results on the fixed available subset: mean ± sample SD over seeds 1301–1306. Each run uses 866 test clips with five references each. Metrics retain their native scales; CIDEr-D is primary.

Condition	CIDEr-D	BLEU-4	METEOR	SPICE	SPIDEr
Baseline	0.6736 ± 0.0220	0.2401 ± 0.0137	0.2328 ± 0.0043	0.1664 ± 0.0046	0.4200 ± 0.0132
Fixed spectral DCR	0.6745 ± 0.0185	0.2384 ± 0.0114	0.2335 ± 0.0022	0.1650 ± 0.0041	0.4198 ± 0.0110
Spectral MACC-Lite	0.6668 ± 0.0212	0.2378 ± 0.0108	0.2320 ± 0.0035	0.1664 ± 0.0020	0.4166 ± 0.0112
Spectral Full MACC	0.6703 ± 0.0175	0.2382 ± 0.0094	0.2319 ± 0.0023	0.1659 ± 0.0041	0.4181 ± 0.0100

Table 21: All three primary audio CIDEr-D contrasts. Positive differences favor the target. Intervals are individual unadjusted paired-t intervals; each test family is Holm-adjusted across three comparisons.

Target minus baseline	Difference	95% interval	Adjusted t p	Adjusted exact p
Fixed spectral DCR	+0.000896	$[-0.016040, +0.017833]$	1.0000	1.0000
Spectral MACC-Lite	−0.006781	$[-0.028950, +0.015389]$	1.0000	1.0000
Spectral Full MACC	−0.003329	$[-0.015313, +0.008656]$	1.0000	1.0000

Table 22: Separate raw-source AudioCaps confirmation, ten paired seeds, 956 official-test clips with five references each. CIDEr-D is primary (mean ± seed SD); SPIDEr is descriptive.

Condition	CIDEr-D	SPIDEr
Baseline	0.6892 ± 0.0109	0.4304
Tuned fixed	0.6865 ± 0.0123	0.4286
MACC-Lite	0.6893 ± 0.0125	0.4306
Full MACC	0.6781 ± 0.0163	0.4245

the different supervised classification objective.

7.16 Computational cost and memory

The implementation costs follow the components being used. Fixed DCR still computes differentiable layer statistics because they define the regularization loss. MACC-Lite adds moving measurements, coefficient states, and bounded update arithmetic. The evaluated Lite path also records held-out cross-entropy at control boundaries, although this metric does not enter its coefficient action. Full MACC uses held-out task feedback in its state and reward and additionally optimizes a policy. These are training operations and can be omitted during ordinary logits-based classifier inference. Geometric acquisition between training rounds still requires feature extraction and fitted class statistics.

Table 23 reports timing components for the initial nine-condition ImageNet-100 study. The synchronized training-compute timer includes forward, loss, backward, clipping, and optimizer execution, but excludes subsequent controller observation and policy operations. Feedback and acquisition have separate timers. Loading, serialization, evaluation, and orchestration are not fully represented by summing these columns. All references record geometry, and initialization is allocated once per conceptual condition. Thus, these comparisons do not estimate overhead against the leanest unregularized or fixed-weight implementation. Coefficient selection and earlier method development incur costs beyond one selected run; the recorded selection costs below are reported separately rather than absorbed into a per-run overhead percentage.

The maximum PyTorch-allocated memory is 6,075,090,432 bytes for every conceptual run in the original nine-condition matrix. Shared initialization may dominate that maximum; equal values do not establish zero additional controller memory. The phase-specific comparison below is narrower but still does not isolate a complete controller footprint.

The duration in Table 24 and the synchronized training-compute component in Table 23 have different measurement boundaries. The former is a phase-duration record; the latter excludes controller observation and policy operations and includes an allocated initialization component. They are not additive or directly subtractable. These timers do not define a common end-to-end boundary.

The entropy-plus-Lite mean phase duration is 2.59% higher, a difference of 136.29 s. These are sequential campaigns, not interleaved timing trials. The peak difference is 262,144 bytes, or 0.25 MiB. Neither quantity is complete application overhead or an isolated controller measurement.

The entropy-plus-fixed follow-up separately records means of 4,837.85 s for synchronized training compute, 182.55 s for feedback inference, 483.99 s for acquisition, and 782.66 s for final AutoAttack. The training and feedback components exclude the reused initialization; they therefore cannot be inserted directly into Table 23 or compared as complete application time. The fixed-control records do not contain a phase-duration or

Table 23: Recorded ImageNet-100 timing components in seconds per conceptual condition run, averaged over six seeds. Training and feedback include their allocated shared-initialization costs. These are component, not complete application, timings.

Method	Training compute	Feedback inference	Acquisition	Final AutoAttack
Random	5093.24	182.25	0.85	792.52
Entropy	5092.81	182.45	490.83	800.95
BADGE	5093.97	182.55	523.26	775.80
Noise Stability	5093.61	182.52	2942.89	777.36
Acquisition only	5093.23	182.38	739.69	798.50
Random + aligned fixed	5240.51	182.30	0.86	811.51
CÆNL + aligned fixed	5238.40	182.28	740.18	820.73
CÆNL + aligned Lite	5238.85	182.27	739.58	804.70
CÆNL + aligned Full	5240.13	182.30	738.97	783.91

Table 24: Post-initialization training time and peak allocation. Duration is the mean sum over five post-initialization training phases. Peaks describe PyTorch allocation in those phases, not isolated controller memory.

Condition	Five-phase duration (s)	Recorded phase peak (bytes)
Entropy	5,255.02	6,066,439,680
Entropy + aligned MACC-Lite	5,391.31	6,066,701,824
Difference	+136.29	+262,144

phase-peak memory aggregate for this new condition, so no new overhead ratio is inferred.

Earlier 104-run control study. The additional 24 tuning runs take 39.66 hours of recorded post-initialization phase training plus acquisition, or 40.07 hours including the three shared initializations. These totals exclude final reporting and other unrecorded overhead. Mean post-initialization training-plus-acquisition time per run is 1.626, 1.669, 1.674 and 1.674 hours for entropy, Lite, fixed and schedule on ResNet-50; corresponding ResNet-18 times are 0.759, 0.760, 0.764 and 0.762 hours. The small per-run difference between Lite and fixed does not account for the separate tuning cost or Lite's earlier development cost. Hence, unchanged feedback in this campaign is not synonymous with an absence of prior tuning. All four conditions compute the monitored holdout metric at matched control windows, including those whose actions do not consume it.

Content-filtered study timers. Across the ten confirmation seeds, mean recorded post-initialization training times are 5,268.5 s for entropy, 5,440.5 s for Lite, and 5,440.1 s for tuned fixed DCR. Corresponding ResNet-18 means are 2,187.4, 2,196.6, and 2,187.3 s. Acquisition adds approximately 482 s per end-to-end run. These secondary timers exclude shared initialization and do not constitute a randomized lean-implementation benchmark or a complete application-cost comparison. Replay records zero acquisition time because they reuse donor labels; the donor's acquisition and annotation costs are not thereby eliminated. The completed six-path diagnostic below provides a more direct comparison of streamlined and fully instrumented training paths.

Table 25: Mean elapsed seconds for 200 measured cached-data steps after 20 warm-up steps, five timing repetitions per path. Fixed strength is 0.10. These are timing repeats, not five training seeds.

Path	Entropy	Fixed	Lite
Streamlined	23.73	24.76	26.12
Instrumented	25.58	26.28	26.29

Completed cached-data path benchmark. Table 25 reports the six prescribed paths. Across paired timing repetitions, streamlined fixed adds 4.34% over streamlined entropy; streamlined Lite adds 5.47% over fixed and 10.04% over entropy. Fully instrumented Lite adds 0.05% over instrumented fixed, illustrating how shared monitoring changes the comparison boundary. These percentages average paired per-repetition overheads; they are not inferential estimates from independent training seeds. Lite's serialized controller state is 2,036 JSON bytes in this measurement. No whole-application speed or hardware-independent saving follows from these short cached-data paths.

The completed 12-run fixed grid records 19.57 hours of training plus acquisition; schedule and contrastive grids record 19.60 and 19.48 hours. These boundaries exclude shared initialization, setup and some evaluation/checkpoint work. Lite's earlier development cost is unknown, not zero, and engineering effort is unmeasured. The new audio study records about 55.01, 57.39, 57.34 and 57.97 minutes mean job wall time for baseline, fixed, Lite and Full, respectively, plus 16.71 hours for its 18 tuning jobs. Feature extraction and initialization are outside those job means; clean-audit phase costs are separate. Neither study establishes a total-cost advantage of feedback.

Spectral-task costs. A five-pair GPT-2/C4 microbenchmark times 250 optimizer steps. Plain execution takes 33.5293 ± 0.0429 s; the sampled-monitoring path takes 33.7633 ± 0.0731 s. The paired time increase is 0.698% with SD 0.142 percentage points; peak allocation increases about 0.089 GiB. Its BF16 parameters, microbatches of two sequences with four accumulation

passes, ten monitor events, optimizer initialization, and timing boundary differ from production.

Production C4 time increases by 7.88%, 7.68%, and 16.62% for fixed, Lite, and Full relative to its monitored baseline. These ratios cannot be combined with the microbenchmark to obtain total overhead relative to plain training.

Diffusion mean task-job wall times are 1,616.81 s for reference sampling and 2,919.80, 2,974.48, 2,978.48, and 2,977.43 s for baseline, fixed, Lite, and Full. They include task-driver training, sampling, and metrics, but exclude the wrapper's later official-test EMA assessment and study preparation. The reference has no training updates and is not an overhead denominator.

Audio successful training-loop times span 1,175.53–1,201.80 s for baseline, 1,308.80–1,333.21 s for fixed, 1,309.24–1,337.74 s for Lite, and 1,336.39–1,375.73 s for Full. These include controller feedback and checkpoints but exclude setup, final decoding, metrics, and saved-result audits. Feature extraction is reused from the pilot. Rejected initialization attempts and verification are separate. No universal sub-10% or end-to-end overhead claim follows.

8 Discussion and limitations

The central question is whether representation geometry can guide regularization during active learning, and whether adapting its coefficient improves on a separately selected constant. The combined aligned-DCR and MACC-Lite treatment improves final-budget classification over entropy acquisition alone. The tuned constant attains slightly higher mean accuracy, with inconclusive paired differences. These findings support geometry-informed regularization under the evaluated protocol and motivate a separate assessment of the coefficient policy. Their interpretation connects the measured gains to the formulation in Sections 3–4, the analysis in Section 5, and the distinct task protocols.

8.1 Implications for representation control

The content-filtered study provides the most direct component evidence. Its final-budget Lite-minus-entropy gains are 1.900 percentage points on ResNet-50 and 1.612 points on ResNet-18, surviving both prespecified adjusted procedures. The earlier 104-run gains of 1.742 and 1.834 points and the six-seed acquisition-rule ablation provide complementary evidence under their original inventories. The feedback settings are unchanged across backbones. These results support a representation treatment that complements entropy acquisition; the benefit does not require the proposed geometric acquisition score. The new training inventory addresses exact-content overlap prospectively, rather than relying only on post-hoc exclusion at evaluation.

The comparison with fixed strength addresses a more specific question. Earlier six-seed studies favor Lite over the prespecified coefficient 0.03 under adjusted paired-t inference. Both subsequent studies select 0.10 independently, with slightly higher fixed-treatment means than Lite on both backbones. In the content-filtered study, Lite-minus-fixed effects are -0.152 and -0.092 points, with individual intervals $[-0.593, +0.289]$ and $[-0.403, +0.219]$. Neither establishes superiority, equivalence, or non-inferiority. The earlier result against 0.03 and the later tuned-control results concern different comparators. Together they show why coefficient selection must be controlled when attributing a benefit to adaptation. Lite's advantage over the selected schedule is also not established in the new 17-contrast family.

The formulation contributes an explicit connection between the measured quantity and the intervention: the supervised scatter residual enters both the loss and the coefficient update. Section 5 characterizes this connection through an isolated local descent result and bounded coefficient responses. The recorded layer-dependent trajectories confirm that the implemented controller responds to the signal. These analytical and implementation properties make the mechanism interpretable, but do not establish that alignment is necessary for the predictive gain. Common-label replay now identifies a beneficial scatter-ratio intervention, with a 2.120-point gain over unregularized training. It finds no established increment from separation or feedback at the selected strengths. The latter comparisons are inconclusive, rather than evidence that those components are ineffective. The design does not isolate normalization or prove that changes in the measured geometry mediate prediction. Notably, the matched-step full-label ResNet-18 reference has higher mean log scatter than Lite (0.997 versus 0.539), yet higher accuracy (74.819% versus 73.151%; Table 37). This descriptive reversal illustrates why smaller aggregate scatter is not a universal predictive ranking. Nor does clean scatter control input sensitivity. The completed five-radius diagnostic strengthens this distinction: no primary 2/255 advantage is established, and descriptive 0.5/255 accuracy remains similar despite reduced clean scatter. The empirical cause of the geometry–robustness dissociation remains open.

Full MACC adds learned actions but provides no established task benefit over the relevant controls. Its result is conditional on the specified state, reward, action space, and horizon; the present comparisons do not identify which of these choices limits performance. For practical evaluation, the tuned constant is therefore an essential reference alongside Lite. Both transfer here without backbone-specific retuning, so the transfer result alone cannot show that feedback saves tuning effort. The measured computation in Section 7.16 informs this comparison, but its protocol-specific timing boundaries do not support a universal overhead or deployment-efficiency claim. When representative labeled feedback data permit coefficient selection, tuned fixed DCR is the simplest regularized reference supported by these comparisons. MACC-Lite is an explicit feedback alternative, but the completed evidence does not establish a predictive or total-cost advantage that would justify

recommending it generally. Full MACC incurs policy optimization without an established additional benefit under the tested configuration.

Table 26 summarizes the findings and the scope of the corresponding inferences.

8.2 Limits of classification and statistical inference

The classification evidence covers two ResNet backbones on one 100-class subset, with a class-balanced sampler and frozen post-initialization BatchNorm statistics. Matching these choices supports comparison within the protocol, but does not establish the best training recipe for every method. In particular, UHerding-Margin uses fixed shared-initial supervised features, capped candidate passes, holdout temperature calibration, and warm-start training. Its ranking cannot be generalized to uncertainty-weighted coverage methods under their original or independently optimized protocols. ImageNet-1K and ViT performance remain unevaluated.

The endpoint also matters. Mixed early-budget ordering on ResNet-50 shows that a final-budget gain is not a uniform improvement throughout acquisition. The new matched-step full-label reference gives an inconclusive ResNet-50 comparison and a ResNet-18 advantage over Lite; neither establishes an equal-performance annotation saving. The full-label reference is not separately optimized to convergence. Geometric atypicality may select unhelpful examples, while aggregate scatter can conceal a poorly separated class pair. Training anchors and all-class assessment differ in class support and transforms. Consequently, the measured geometry should be interpreted alongside predictive outcomes at each budget.

New seeds quantify variability from initialization and data partitioning within this benchmark. They do not create an untouched image benchmark: official validation images were used for reporting during development. The earlier studies contain five exact training–validation content matches with conflicting labels. Excluding them at evaluation leaves their comparison decisions unchanged but cannot undo possible training exposure. Their within-training duplicates and affected holdouts (Section 7.5) remain part of those studies' provenance. The new 206-run study instead excludes all identified conflicting and cross-split exact identities before training, with one retained copy for same-label within-split duplicates. That correction does not authenticate upstream labels or certify perceptual near-duplicate absence, and it does not create a previously untouched benchmark.

Statistical interpretation must retain the distinction between study families. The original six-seed exact tests have the resolution limit derived in Section 6.8; their nonrejection does not independently demonstrate absent effects. The two ten-seed families permit rejection and retain both test types. Paired-t inference remains conditional on small-sample assumptions, and neither procedure provides study-wide error control across the full sequence of experiments. The later comparisons were specified before execution, but were informed by earlier findings. Keeping the original and later families separate preserves this chronology; it does not eliminate the influence of sequential development on the research questions selected.

8.3 Task dependence and generalization

The language, diffusion, and audio experiments evaluate spectral regularization, whereas classification uses a supervised class-scatter objective. Their joint value is to test the scope of geometry-based interventions across tasks, with outcomes judged by the relevant task metrics. They do not constitute a transfer test of one unchanged loss or controller across modalities.

On C4, fixed spectral DCR produces a small loss reduction on an unused local validation shard. This is evidence for that objective and protocol, rather than for adaptive control. The result concerns a locally selected single-shard training subset and pretrained GPT-2; local non-use of the evaluation shard does not establish absence of pretraining overlap. In CIFAR-10 diffusion, every spectral condition worsens FID relative to matched unregularized fine-tuning. One possible explanation is that the imposed spectral target restricts variation needed by the denoising objective, but the present experiments do not test that mechanism. The result covers one pretrained U-Net, sampler, and training schedule, without determining whether a separately tuned spectral objective could behave differently.

Neither AudioCaps study establishes a CIDEr-D benefit or statistical equivalence. The earlier study retains an approximately 89%-coverage third-party subset and four flagged zero training clips under its frozen policy. The separate raw-source study raises retained coverage to about 98%, has no retained zero clips and adds development-selected fixed strength with ten paired seeds, yet still establishes no adjusted primary gain. This addresses the narrow comparator and source-coverage objections without converting a null comparison into equivalence. Possible AST pretraining overlap, prior benchmark exposure, residual near-duplicates and upstream authenticity remain limitations in both studies. Higher coverage alone does not authenticate recordings or justify extrapolation to other audio tasks.

Across these studies, a geometric target is a modeling choice whose usefulness depends on its relationship to the task loss. The positive classification result, small fixed C4 effect, adverse diffusion result, and inconclusive audio contrasts motivate task-specific validation of that relationship. None supplies evidence of a universal cross-task gain, and the spectral outcomes cannot directly validate or invalidate the different aligned classification objective.

8.4 Directions for further research

Common-label component replay now narrows the source of the classification benefit to a supported scatter

Table 26: Interpretation and scope of the reported evidence. Analytical statements retain their stated assumptions; empirical findings apply to the evaluated protocols.

Question	Supported finding	Limit of the inference
Objective–feedback alignment	Shared supervised scatter residual, local isolated descent, and bounded coefficients	Necessity of alignment, global optimization, and closed-loop stability are not established
Active learning	Matched acquisition comparisons; entropy can be combined with aligned regularization	No universal acquisition ranking or equal-performance 40% annotation saving
Scatter components	Ratio-only benefit under common entropy labels; complete fixed DCR exceeds separation-only	Separation increment, feedback increment, interaction, normalization, and geometry mediation remain unresolved
MACC-Lite	Final-budget gains over entropy on two backbones; earlier gains over specified 0.03 controls	No established superiority to the tuned constant, equivalence, or isolated causal benefit of adaptation
Full MACC	Policy updates execute, without an established added classification benefit	Results concern the evaluated controller; learned query thresholds are not tested
Adversarial robustness	Completed five-radius evaluation; no established 2/255 advantage	No meaningful robustness gain or certification follows from the geometric objectives
Language modeling	Small fixed spectral-DCR loss reduction on an unused local C4 shard	Effect is specific to this spectral objective and language-modeling protocol
Diffusion	All spectral conditions have higher FID than matched fine-tuning	The result does not evaluate aligned classification DCR or other tuned generative objectives
Audio captioning	Separate six- and ten-seed studies with no established adjusted CIDEr-D gain	No equivalence or full-original-benchmark comparability is established

intervention under this protocol. Further mechanistic work should distinguish the role of normalization and target choice, retune reduced objectives under matched search budgets, and test whether measured geometric changes mediate predictive outcomes. Replication on independent image datasets is needed to assess whether the present component effects generalize beyond ImageNet-100.

The practical value of feedback also warrants a test beyond backbone transfer on the same subset. Predefined comparisons on new datasets, architectures, class-imbalance regimes, and annotation budgets could measure whether unchanged feedback settings reduce control-selection effort while maintaining predictive performance. Such a study should specify the primary endpoint and a meaningful effect size in advance, use a sample size appropriate to that question, and account for both tuning and training cost. The present prospective exact-content grouping addresses the identified byte/RGB identities; perceptual near-duplicate auditing and upstream provenance checks remain useful extensions.

For the spectral variants, varying target rank and regularization strength would help test whether the observed task dependence reflects the imposed geometry or its magnitude. Any robustness extension would require an explicit link to input sensitivity or robust optimization, followed by a dedicated adversarial evaluation. These are directions for testing mechanisms and scope; their outcomes are not implied by the present clean-accuracy or representation measurements.

8.5 Reproducibility and data availability

Code, experiment configurations, numeric result records, and CPU statistical reconstruction scripts are available at https://github.com/alqithami/CANL. The earlier code, numerical records, and reconstruction tools are pinned to revision `caf2c05`; each protocol identifies its frozen training launcher. The subsequent classification and diagnostic campaign code is available at `ad830e5`. The completed 206-run endpoint records and CPU reconstruction are pinned to `8b64be0` and accompany this version as a numerical supplement. Their binding and verification scope are given in Appendix C. The completed diagnostic and raw-source audio records, all eight new planned contrasts, original audio study code, and CPU reconstruction are pinned to `51412af`. They accompany this version in the numerical supplement; its release manifest records the immutable revision and public-file checksums. Historical studies and later variants remain separate; licensed raw data and large checkpoints are excluded. Appendix F documents the retained evidence, the audio initialization workaround that limits seed-only byte-identical reproduction, and the verification boundaries. Statistical reconstruction from released records is distinct from independently repeating training or checkpoint inference.

9 Conclusion

This study introduced CÆNL, an active-learning framework that connects representation measurements to an explicit regularization objective and its coefficient control. For classification, the same normalized within-class to between-class scatter residual enters the loss

and feedback, alongside a safeguard against loss of class separation. Separating acquisition, regularization, and coefficient selection makes their roles open to controlled comparison. The supporting analysis characterizes the geometry measurements, conditional local regularizer descent, and bounded coefficient responses, while identifying the additional conditions needed to relate feature geometry to prediction and robustness. The full-update condition makes gradient interaction explicit, and the constructive counterexample shows that identical clean geometry and logits can coexist with arbitrarily small robust radii when input sensitivity is unconstrained.

The principal empirical finding is a classification benefit from the evaluated scatter intervention, supported by complementary acquisition, tuned-control, and common-label experiments. The completed 206-run study retains the original formulation while filtering exact-content duplicates before training. At 60% annotations, Lite with DCR exceeds entropy alone by 1.900 percentage points on ResNet-50 and 1.612 points on ResNet-18; both survive the prespecified 17-comparison corrections. Common-label replay further identifies a 2.120-point benefit from the fixed scatter-ratio penalty without differences in acquired examples. This supplies conditional component evidence beyond the earlier end-to-end improvements. Lite also outperforms the tested matched one-view contrastive-feature regularizer, with the scope of that implementation made explicit.

The coefficient-policy comparisons delimit this contribution. Independently tuned fixed DCR has slightly higher mean accuracy than Lite on both backbones. The end-to-end and common-label comparisons establish no independent feedback advantage, and they do not establish equivalence or absence of a feedback effect. The separation increment and factorial interaction are also inconclusive at the selected strengths. The full-label, matched-step comparison favors full labels on ResNet-18 and is inconclusive on ResNet-50, preventing an equal-performance annotation-saving claim. Thus the evidence supports the specified representation intervention more directly than a general claim of superior adaptive control.

The complementary experiments define the scope of this finding. Earlier acquisition comparisons show that the benefit can be obtained without geometric acquisition, whereas learned Full MACC adds no established classification advantage. Completed attacks at five radii establish no primary 2/255 robust-accuracy advantage, while retaining near-zero 8/255 accuracy. Smaller-radius and conditional geometry summaries do not identify a causal robustness mechanism. The separate spectral formulations yield a small fixed-DCR loss reduction on C4, worse CIFAR-10 diffusion FID, and no established AudioCaps captioning benefit in either the earlier subset study or the new 58-run raw-source study. Short path timings quantify feedback overhead without establishing total-cost savings. These task-dependent outcomes emphasize that a geometric objective must be assessed through the predictive or generative task it is intended to improve.

The study provides a specified representation-control mechanism, conditional analytical results, and a controlled empirical decomposition of its classification benefit. Tuned fixed DCR is the simplest supported regularized reference under the tested conditions; bounded feedback is an interpretable alternative whose distinct predictive and total-cost utility remains to be established. Independent datasets, broader architectures, and matched tuning budgets are needed to test that utility beyond the present benchmark. Robustness mechanisms require input-sensitivity evidence beyond clean scatter measurements. The released implementations and numerical records support scrutiny of both the positive findings and the negative or inconclusive comparisons, and provide a basis for further study of geometry-informed learning.

Funding

This research received no external funding or in-kind support.

Declaration of competing interest

The author declares no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

CRediT authorship contribution statement

Saad Alqithami: Conceptualization, Methodology, Data curation, Formal analysis, Investigation, Project administration, Resources, Software, Validation, Visualization, Writing – original draft, Writing – review and editing.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

ChatGPT (OpenAI) was used to assist with manuscript organization and editing, mathematical exposition, programming and debugging, interpretation of saved experimental artifacts, reference checking, and preparation of the LaTeX/TikZ and Matplotlib figures. Numerical figures use recorded experimental measurements or stated statistical summaries; schematics describe the implemented workflow. The scope of saved-result verification is documented in Appendix F. The author is responsible for the scientific design, interpretation, and content of this work.

A Supplementary results and development analysis

This appendix provides paired observations for the six-seed classification follow-ups and AudioCaps evaluation,

followed by the development screen that motivated the aligned classification objective. The separate ten-seed control study is reported in Appendix B, and secondary C4 comparisons in Appendix D. Experimental units and comparison families follow Section 6.8.

A.1 Paired classification outcomes

Table 27 gives the observations underlying the entropy-treatment comparisons in Sections 7.7 and 7.8. Each row pairs methods under the same seed and shared initialization. The fixed condition uses aligned DCR with coefficient 0.03; it is distinct from the independently tuned constant in Appendix B. Table 28 adds the UHerding-Margin comparison from Section 7.9. The repeated entropy and Lite columns are shared reference runs, not additional independent observations. The tables retain the separate inferential families specified for these follow-ups.

Table 29 completes the original robustness family from Section 6.8. The same six paired seeds and full-set attack endpoint are used for each contrast. All Holm-adjusted paired-t and exact values exceed 0.05. The individual intervals are reported under the stated small-sample assumptions; they are not population robustness certificates.

A.2 Paired AudioCaps outcomes

Table 30 reports CIDEr-D for the six paired training seeds in the AudioCaps study. Baseline, fixed, Lite, and Full follow the task-specific spectral protocol in Section 6.9.3; the latter three names do not denote the aligned classification loss. All conditions use the same 866 retained test recordings and reference captions. These seed-level observations underlie the means and primary contrasts in Tables 20 and 21. The training seed remains the replication unit; the test recordings are evaluation examples within each run.

A.3 Development screen for the classification objective

The aligned formulation was selected using a bounded development screen before the end-to-end classification comparisons. Each of two source seeds had a shared 200-step stabilization and five matched 2,000-step continuations on a fixed labeled set. The selection criterion required a favorable geometric direction and an assessment-loss increase no larger than 0.02 nats in either seed. This tolerance guided development; it was not a statistical noninferiority test.

Relative to no regularization, spectral DCR changed the geometric-mean trace ratio by $+0.077\%$ and $+0.054\%$ in the two seeds. The corresponding changes were -2.009% and -2.365% for aligned fixed DCR, -4.426% and -5.015% for aligned Lite, and -1.691% and -1.300% for aligned Full. Loss changes were mixed but remained within the screen's tolerance. These observations motivated testing the aligned objective under the full acquisition and training protocol; they do not themselves establish annotation efficiency or robustness. The development observations remain separate from the six-seed and ten-seed performance comparisons.

B Control-study seed outcomes and sensitivity

The control study separates coefficient selection from paired confirmation and backbone transfer. This appendix supplements Section 7.3 with the individual final outcomes and the content-duplicate sensitivity analysis described in Section 7.5.

B.1 Confirmation and transfer outcomes

Table 31 reports final official-validation accuracy for every confirmation and transfer seed. All four treatments use entropy acquisition. Constant and scheduled strengths were selected using the distinct ResNet-50 tuning seeds 41001–41003 (Table 6) and then held fixed for both backbones. Those tuning runs do not enter the paired confirmation or transfer comparisons. Within each backbone, rows retain the seed pairing used for the effects in Table 8.

B.2 Sensitivity to content-duplicate exclusion

Table 32 excludes the same five content-overlapping validation examples from every treatment, retaining 4,995 images. It preserves the six primary contrasts and applies Holm correction separately to the paired-t and exact test values across that family. All adjusted decisions agree with the primary analysis. This post-hoc check assesses the sensitivity of the reported comparisons; it does not reverse possible training exposure or reproduce training on a deduplicated dataset. The planned 5,000-image endpoint remains primary.

C Content-filtered classification replication

This appendix retains the complete inferential record of the 206-run study in Sections 6.5 and 7.1. Table 33 reports all 17 planned contrasts, including raw and adjusted tests. Table 34 records every selection candidate; Tables 35 and 36 retain the paired final accuracies. Secondary predictive and geometric outcomes appear in Table 37. Tuning seeds are excluded from confirmation inference. The 60 end-to-end confirmation and 60 transfer runs, plus 50 component replays and 36 tuning runs, exhaust the declared 206 methods; replay donor rows are not counted twice.

The content audit excludes 99 training and five validation images before this study's partitions and training. Validation has 50 examples in 96 classes, 49 in three classes, and 48 in one class. Top-1 accuracy is

Table 27: Paired entropy-acquisition treatments at 60% annotations on ImageNet-100. Clean accuracies are percentages and differences are percentage points; fixed aligned DCR uses coefficient 0.03. The entropy and Lite reference runs are shared with the corresponding follow-up comparisons. Robust-correct counts use standard AutoAttack at L_∞ radius 8/255 on all 5,000 validation images. Fixed DCR has 0/5,000 robust-correct images in every seed.

Seed	Entropy	Entropy + fixed	Entropy + Lite	Fixed − entropy	Lite − fixed	Lite robust correct
801	71.14	72.38	73.10	+1.24	+0.72	0/5,000
802	71.76	72.22	73.48	+0.46	+1.26	1/5,000
803	71.32	72.00	73.30	+0.68	+1.30	0/5,000
804	71.74	72.78	73.70	+1.04	+0.92	0/5,000
805	71.08	71.70	73.34	+0.62	+1.64	0/5,000
806	71.74	72.50	73.24	+0.76	+0.74	0/5,000
Mean	71.463	72.263	73.360	+0.800	+1.097	—
Sample SD	0.320	0.381	0.208	0.288	0.364	—

Table 28: Paired UHerding-Margin comparison at 60% annotations on ImageNet-100. Clean accuracies are percentages; the summary row reports mean ± sample SD. Entropy and entropy-plus-Lite are the same reference runs as in Table 27. UHerding-Margin uses the fixed shared-initial ResNet-50 features specified in Section 6.7. Robust-correct counts concern UHerding-Margin under standard AutoAttack at L_∞ radius 8/255 on all 5,000 validation images.

Seed	Entropy	UHerding-Margin	Entropy + aligned Lite	UHerding robust correct
801	71.14	71.58	73.10	0/5,000
802	71.76	72.64	73.48	1/5,000
803	71.32	72.38	73.30	0/5,000
804	71.74	71.46	73.70	0/5,000
805	71.08	71.40	73.34	1/5,000
806	71.74	71.02	73.24	0/5,000
Mean ± SD	71.463 ± 0.320	71.747 ± 0.626	73.360 ± 0.208	—

Table 29: All seven original AutoAttack accuracy contrasts at L_∞ radius 8/255 on 5,000 ImageNet-100 validation images. Effects and individual, unadjusted 95% paired-t intervals are in percentage points over six paired seeds. Each test column is separately Holm-adjusted across these seven comparisons. Aligned conditions use geometric acquisition unless random acquisition is named explicitly.

Target minus reference	Difference (pp)	95% CI (pp)	Adjusted t p	Adjusted exact p
Aligned fixed DCR − acquisition only	+0.0033	$[-0.0125, +0.0191]$	1.0000	1.0000
Aligned MACC-Lite − aligned fixed DCR	−0.0100	$[-0.0215, +0.0015]$	0.5291	1.0000
Aligned Full MACC − aligned fixed DCR	−0.0100	$[-0.0215, +0.0015]$	0.5291	1.0000
Aligned MACC-Lite − entropy	−0.0033	$[-0.0119, +0.0052]$	1.0000	1.0000
Aligned Full MACC − entropy	−0.0033	$[-0.0191, +0.0125]$	1.0000	1.0000
Acquisition only − random	0.0000	$[-0.0133, +0.0133]$	1.0000	1.0000
Random + aligned fixed DCR − random	+0.0033	$[-0.0052, +0.0119]$	1.0000	1.0000

Table 30: Paired AudioCaps CIDEr-D outcomes on the 866 retained test recordings, on the native metric scale. Fixed, Lite, and Full denote the task-specific spectral conditions; larger values indicate better caption agreement. Each row corresponds to one shared training seed.

Seed	Baseline	Fixed	Lite	Full
1301	0.699902	0.706662	0.695297	0.697474
1302	0.683473	0.656485	0.683658	0.675072
1303	0.634166	0.656615	0.658886	0.647493
1304	0.681798	0.678712	0.663435	0.671009
1305	0.672890	0.676179	0.633839	0.655220
1306	0.669446	0.672400	0.665875	0.675435

pooled over the 4,995 retained images; geometry retains equal-class weighting. The final active-learning gradient set has 75,322 images after removing the 632 feedback examples; the full-label gradient pool has 125,958. These counts differ from the earlier unfiltered studies. The duplicate policy checks exact file and decoded-RGB identities; filename/synset authentication and perceptual near-duplicate absence are not established.

The experiment's frozen input binding is

```
18a7164a39e295a27fe864a1d858fae3069a25a9f4997
f6435d8d048f870e65e
```

The completed compact review archive has SHA-256

```
37b1d7d55f2bab65582f2109583381cf83433686c6a05
17fd24f7364483565db
```

An inherited machine-readable selection description still says "24 tuning runs"; the frozen protocol, candidate grid, expected counts, and completed records

Table 31: Control-study final top-1 accuracy (%) at 60% annotations on all 5,000 ImageNet-100 validation images. Each row pairs four entropy-acquisition treatments under one seed. Fixed and scheduled strengths were selected before these runs and transferred unchanged to ResNet-18.

Backbone	Seed	Entropy	Lite	Tuned fixed	Tuned schedule
ResNet-50	42001	70.78	73.16	73.10	72.40
ResNet-50	42002	71.22	72.68	73.32	72.64
ResNet-50	42003	72.04	73.56	73.30	72.98
ResNet-50	42004	72.76	73.66	74.68	73.40
ResNet-50	42005	70.28	72.84	73.52	72.62
ResNet-50	42006	71.50	73.62	73.58	72.92
ResNet-50	42007	72.56	74.14	74.06	73.34
ResNet-50	42008	72.12	73.16	73.60	73.44
ResNet-50	42009	72.30	74.22	73.76	73.14
ResNet-50	42010	71.22	73.16	73.16	73.34
ResNet-18	43001	71.66	73.02	73.06	72.56
ResNet-18	43002	70.48	72.96	72.90	72.36
ResNet-18	43003	71.04	73.28	73.60	72.86
ResNet-18	43004	71.28	73.02	74.34	73.14
ResNet-18	43005	71.40	73.18	72.54	72.68
ResNet-18	43006	70.72	72.52	73.14	72.32
ResNet-18	43007	71.72	73.88	74.12	73.88
ResNet-18	43008	71.38	72.88	73.82	72.60
ResNet-18	43009	71.50	73.48	74.14	72.94
ResNet-18	43010	71.34	72.64	72.54	72.28

Table 32: Content-duplicate sensitivity on 4,995 validation images. Effects and individual, unadjusted 95% paired-t intervals are in percentage points. Fixed and schedule denote the selected controls. Each test column is separately Holm-adjusted across the same six contrasts as in the primary analysis; all adjusted decisions are unchanged. Positive effects favor Lite.

Backbone	Contrast	Difference	95% CI	Adjusted t p	Adjusted exact p
ResNet-50	Lite − entropy	+1.742	[+1.358, +2.125]	1.42×10^{-5}	0.01172
ResNet-50	Lite − fixed	−0.188	[−0.527, +0.151]	0.2406	0.25391
ResNet-50	Lite − schedule	+0.398	[+0.075, +0.722]	0.06367	0.08789
ResNet-18	Lite − entropy	+1.836	[+1.564, +2.108]	5.86×10^{-7}	0.01172
ResNet-18	Lite − fixed	−0.328	[−0.732, +0.075]	0.1978	0.20703
ResNet-18	Lite − schedule	+0.322	[+0.155, +0.490]	0.00744	0.03125

specify 36. All 36 records enter selection; the frozen inputs are preserved rather than silently rewriting that documentation discrepancy. This has no effect on the reconstructed selection or contrasts.

D Secondary C4 outcomes

The C4 results in Section 7.12 use NLL as the primary endpoint. Table 38 provides the accompanying perplexity, pooled token ECE, and token-accuracy contrasts on the unused local validation shard. They evaluate the same final checkpoints from six training seeds; the additional shard does not introduce new training replications. Perplexity is the exponential of NLL and is reported as a descriptive transformation, not an independent primary outcome.

Each endpoint has three regularized-condition-minus-baseline contrasts. Individual paired-t confidence intervals are unadjusted; Holm correction is applied separately to the paired-t and exact sign-flip values within each endpoint's three-contrast family. The adjusted exact values are reconstructed from the archived raw values. This correction does not control multiplicity jointly across endpoints. Positive differences indicate worse perplexity or ECE, whereas negative differences indicate worse token accuracy. The scaling shown in each panel applies to both effects and interval limits; test p-values are not rescaled.

E Completed checkpoint diagnostics and raw-source audio

The diagnostic and audio families are separate from the 17 classification contrasts. All new planned contrasts, not only favorable comparisons, are retained below. Confidence intervals are individual unadjusted 95% paired-t intervals, not simultaneous bands. Public JSON additionally retains all 150 per-radius suite summaries, shared-clean-correct conditioning, geometry/displacement groups, all cost repetitions, all 18 audio selection runs and 40 confirmation endpoints and secondary metrics. CPU reconstruction verifies counts, seed pairing, estimates, intervals, exact tests, Holm

Table 33: All 17 prespecified contrasts in the content-filtered study. Effects and individual unadjusted 95% paired-t intervals are in percentage points. Raw tests and separate Holm adjustments over the entire 17-comparison family are retained. Replay uses common labels on ResNet-50; end-to-end comparisons follow each condition's prescribed annotation protocol. The interaction is fixed minus ratio minus separation plus entropy.

Design	Contrast	Effect	95% CI	Raw t p	Holm t p	Raw exact p	Holm exact p
ResNet-50	Lite − entropy	+1.900	[+1.536, +2.264]	8.77×10^{-7}	1.31×10^{-5}	0.00195	0.0332
ResNet-50	Lite − fixed	−0.152	[−0.593, +0.289]	0.455	1	0.45703	1
ResNet-50	Lite − schedule	+0.631	[+0.196, +1.065]	0.00949	0.08543	0.01758	0.1582
ResNet-50	Lite − contrastive	+1.093	[+0.691, +1.495]	0.00017	0.00203	0.00391	0.04688
ResNet-50	Lite − full labels	−0.130	[−0.531, +0.270]	0.48098	1	0.53125	1
ResNet-18	Lite − entropy	+1.612	[+1.247, +1.977]	3.61×10^{-6}	5.06×10^{-5}	0.00195	0.0332
ResNet-18	Lite − fixed	−0.092	[−0.403, +0.219]	0.51958	1	0.5332	1
ResNet-18	Lite − schedule	+0.280	[−0.152, +0.712]	0.17627	1	0.18359	1
ResNet-18	Lite − contrastive	+0.869	[+0.428, +1.309]	0.00157	0.01731	0.00391	0.04688
ResNet-18	Lite − full labels	−1.668	[−2.248, −1.087]	0.00011	0.00146	0.00195	0.0332
R50 replay	ratio − entropy	+2.120	[+1.783, +2.457]	1.79×10^{-7}	3.05×10^{-6}	0.00195	0.0332
R50 replay	separation − entropy	−0.072	[−0.328, +0.183]	0.53929	1	0.53906	1
R50 replay	fixed − ratio	−0.132	[−0.437, +0.173]	0.35271	1	0.36328	1
R50 replay	fixed − separation	+2.060	[+1.672, +2.448]	7.61×10^{-7}	1.22×10^{-5}	0.00195	0.0332
R50 replay	Lite − fixed	+0.060	[−0.199, +0.320]	0.61317	1	0.63086	1
R50 replay	schedule − fixed	−0.444	[−0.730, −0.159]	0.00651	0.06514	0.01172	0.11719
R50 replay	Factorial interaction	−0.060	[−0.524, +0.404]	0.7765	1	0.78711	1

Table 34: All 12 candidates in the 36-run content-filtered search. Values are training-holdout means over three seeds (632 images each). These are selection scores, not independent validation estimates. An asterisk denotes the selected candidate in its family.

Candidate	Top-1 (%)	Cross-entropy
Fixed 0.003	76.688	0.9178
Fixed 0.01	76.055	0.9265
Fixed 0.03	76.055	0.8861
Fixed 0.1*	77.795	0.8338
Increasing 0.03	75.633	0.9201
Increasing 0.1	76.530	0.8696
Decreasing 0.03	75.949	0.9028
Decreasing 0.1*	77.162	0.8484
Contrastive 0.003	76.002	0.9431
Contrastive 0.01	75.422	0.9306
Contrastive 0.03	76.213	0.9209
Contrastive 0.1*	76.213	0.8866

decisions and selection means without GPU execution or raw-data access.

F Data provenance and reproducibility

Reproducibility depends on the data sources, accepted model states, and records available for verification. This appendix documents those boundaries for the protocols in Section 6 and the materials released in Section 8.5.

F.1 Data sources and evaluation partitions

ImageNet-100 uses the CMC class selection from the original ImageNet images, preserved without reencoding. The earlier selected inventory contains 126,689 training and 5,000 validation images. Content-hash overlaps and label conflicts are described in Section 7.5. The recorded checksums identify the copies used and support integrity checks; they do not independently certify upstream data provenance. The 206-run study applies the prospective identity-group exclusions in Section 6.5, retaining 126,590 training and 4,995 validation images. Its exact-content separation does not retrospectively change the earlier inventories or certify upstream authenticity.

C4 training uses `en/c4-train.00000-of-01024.json.gz`, and controller feedback and development use `en/c4-validation.00000-of-00008.json.gz`. Final unused-shard evaluation uses `en/c4-validation.00002-of-00008.json.gz`. The latter shard was not used in local training or development; absence of GPT-2 pretraining overlap is not established.

The CIFAR-10 diffusion study partitions the original training images for fine-tuning and internal assessment, while FID uses the complete original training distribution as its reference. The partition and official-test denoising assessment are specified in Section 6.9.2; the internal holdout is not claimed to be unseen during model pretraining. AudioCaps uses original annotations matched to the available `OpenSound/AudioCaps` recordings. The coverage, exclusions, waveform checks, and possible AST pretraining overlap are documented in Section 6.9.3. These data-specific conditions continue to apply when reproducing the reported comparisons.

F.2 Verification of saved results and execution records

Verification of the initial classification study reconstructs metrics from saved validation, feedback, and attack arrays. Checks of the exported numerical records also reproduce paired effects and tests across classification, C4, diffusion, and audio, including the entropy-treatment budget curves and layer-geometry summaries.

Table 35: Every end-to-end confirmation and transfer seed in the content-filtered study: final top-1 accuracy (%) on 4,995 images. The six conditions share each seed's initialization. "Full" uses all training labels with matched optimizer steps; the other five use 60%.

Backbone	Seed	Entropy	Lite	Fixed	Schedule	Contrastive	Full
ResNet-50	62001	71.451	72.392	73.373	72.753	72.392	73.734
ResNet-50	62002	71.331	73.173	73.153	72.252	72.052	73.033
ResNet-50	62003	71.151	73.814	73.393	72.633	72.653	73.493
ResNet-50	62004	71.471	73.694	74.715	73.293	73.093	73.874
ResNet-50	62005	71.231	73.213	72.933	71.772	71.532	72.653
ResNet-50	62006	72.573	74.274	74.274	73.253	73.013	74.555
ResNet-50	62007	71.632	72.953	73.453	73.073	71.992	73.453
ResNet-50	62008	71.652	74.074	73.153	73.233	72.132	73.614
ResNet-50	62009	70.490	72.272	72.412	71.311	70.791	72.693
ResNet-50	62010	71.952	74.074	74.595	74.054	73.353	74.134
ResNet-18	63001	71.211	72.713	72.933	72.713	71.732	74.454
ResNet-18	63002	72.172	73.514	73.153	72.973	72.252	74.154
ResNet-18	63003	71.792	73.814	73.834	73.393	72.633	75.035
ResNet-18	63004	71.612	74.154	73.413	72.973	72.052	74.334
ResNet-18	63005	71.552	73.353	73.534	72.392	72.713	74.695
ResNet-18	63006	70.811	72.693	73.413	72.873	72.492	74.955
ResNet-18	63007	71.892	72.893	73.113	72.593	71.852	75.495
ResNet-18	63008	72.212	73.173	73.714	74.134	72.653	75.696
ResNet-18	63009	70.450	72.372	72.713	71.892	71.491	74.615
ResNet-18	63010	71.692	72.833	72.613	72.773	72.953	74.755

Table 36: All common-label component outcomes: final top-1 accuracy (%) on 4,995 images. Entropy is the common-label donor, reused from Table 35; all other columns are separately trained replays of that seed's acquisition sequence.

Seed	Entropy	Ratio	Separation	Fixed both	Lite both	Schedule both
62001	71.451	73.433	71.752	72.693	73.053	73.053
62002	71.331	73.654	71.471	72.853	72.913	72.212
62003	71.151	73.574	71.171	73.794	73.473	73.373
62004	71.471	74.394	71.752	74.034	74.154	73.413
62005	71.231	72.573	70.551	72.873	73.013	72.232
62006	72.573	74.134	71.892	74.034	74.074	74.214
62007	71.632	73.754	71.772	74.194	73.413	73.273
62008	71.652	74.174	71.552	74.014	74.414	73.413
62009	70.490	72.272	70.270	71.972	72.192	71.451
62010	71.952	74.174	72.032	74.354	74.715	73.734

Table 37: Secondary final outcomes in the content-filtered study, mean ± seed SD over ten paired seeds. Log scatter is the equal-layer mean of the reported all-class normalized log scatter ratios. These descriptive metrics do not define additional primary test families or establish geometric mediation of accuracy.

Backbone	Treatment	Cross-entropy	ECE	Top-5 (%)	Mean log scatter
ResNet-50	Entropy	1.144 ± 0.016	0.079 ± 0.003	90.875 ± 0.216	1.389 ± 0.011
ResNet-50	Entropy + Lite	1.021 ± 0.021	0.055 ± 0.006	91.982 ± 0.333	0.499 ± 0.060
ResNet-50	Entropy + tuned fixed	1.012 ± 0.018	0.055 ± 0.006	92.130 ± 0.208	0.355 ± 0.066
ResNet-50	Entropy + schedule	1.056 ± 0.022	0.066 ± 0.006	91.784 ± 0.310	1.006 ± 0.033
ResNet-50	Entropy + contrastive	1.093 ± 0.024	0.074 ± 0.005	91.369 ± 0.314	1.007 ± 0.031
ResNet-50	Full labels, matched steps	0.992 ± 0.017	0.039 ± 0.004	92.012 ± 0.395	1.197 ± 0.017
ResNet-18	Entropy	1.196 ± 0.023	0.095 ± 0.004	90.454 ± 0.409	1.165 ± 0.014
ResNet-18	Entropy + Lite	1.090 ± 0.022	0.083 ± 0.005	91.508 ± 0.427	0.539 ± 0.026
ResNet-18	Entropy + tuned fixed	1.086 ± 0.017	0.081 ± 0.004	91.495 ± 0.237	0.424 ± 0.023
ResNet-18	Entropy + schedule	1.119 ± 0.021	0.084 ± 0.005	91.183 ± 0.385	0.816 ± 0.020
ResNet-18	Entropy + contrastive	1.144 ± 0.020	0.090 ± 0.004	91.023 ± 0.192	0.839 ± 0.021
ResNet-18	Full labels, matched steps	0.970 ± 0.018	0.052 ± 0.004	92.276 ± 0.319	0.997 ± 0.015

Table 38: Secondary C4 comparisons on the unused local validation shard. Each panel reports three regularized-condition-minus-baseline contrasts with individual, unadjusted 95% paired-t intervals. The paired-t and exact sign-flip p-values are separately Holm-adjusted within each panel. Multiply effects and interval limits by the scale shown in the panel heading to recover their original units; token ECE and accuracy use proportions rather than percentages.

Condition minus baseline	Difference	95% CI	Raw t p	Holm t p	Raw exact p	Holm exact p
Perplexity: differences in 10^{-2}						
Fixed spectral DCR	−0.86260	$[-0.96757, -0.75763]$	4.4061×10^{-6}	8.8122×10^{-6}	0.03125	0.09375
Spectral MACC-Lite	+5.74906	$[+5.54778, +5.95034]$	8.8781×10^{-9}	2.6634×10^{-8}	0.03125	0.09375
Spectral Full MACC	+0.69705	$[-0.55390, +1.94800]$	0.21147	0.21147	0.21875	0.21875
Pooled token ECE: differences in 10^{-4}						
Fixed spectral DCR	−1.07333	$[-4.14827, +2.00161]$	0.41069	0.78043	0.37500	0.68750
Spectral MACC-Lite	+8.07857	$[+6.41836, +9.73877]$	5.7942×10^{-5}	0.00017	0.03125	0.09375
Spectral Full MACC	+1.14058	$[-1.97737, +4.25853]$	0.39021	0.78043	0.34375	0.68750
Token accuracy: differences in 10^{-4}						
Fixed spectral DCR	−0.38336	$[-1.21617, +0.44945]$	0.28989	0.28989	0.28125	0.28125
Spectral MACC-Lite	−1.47509	$[-1.89933, -1.05086]$	0.00029	0.00088	0.03125	0.09375
Spectral Full MACC	−0.26752	$[-0.63924, +0.10420]$	0.12355	0.2471	0.09375	0.18750

Table 39: All three primary robust-top-1 effects at 2/255 (percentage points), ten paired seeds. Adjusted t p = 0.90699 and adjusted exact $p = 1$ for every contrast.

Contrast	Difference	95% interval
Lite − entropy	-0.01001	$[-0.03805, +0.01803]$
fixed − entropy	+0.00801	$[-0.03637, +0.05238]$
Lite − fixed	-0.01802	$[-0.05527, +0.01924]$

Table 40: All five raw-source CIDEr-D contrasts. The last column is the primary Holm-adjusted exact p; paired-t results are sensitivity analysis only.

Contrast	Difference	Exact p
Fixed – baseline	-0.002738	1.00000
95% interval: $[-0.011768, +0.006293]$; adjusted t $p = 1.00000$		
Lite – baseline	+0.000121	1.00000
95% interval: $[-0.010363, +0.010606]$; adjusted t $p = 1.00000$		
Full – baseline	-0.011062	0.13672
95% interval: $[-0.020274, -0.001851]$; adjusted t $p = 0.11868$		
Lite – fixed	+0.002859	1.00000
95% interval: $[-0.008740, +0.014458]$; adjusted t $p = 1.00000$		
Full – fixed	-0.008324	0.50000
95% interval: $[-0.019602, +0.002953]$; adjusted t $p = 0.51717$		

Table 41: All ten paired raw-source AudioCaps CIDEr-D outcomes. B: baseline; F: tuned fixed; L: Lite; P: Full MACC.

Seed	B	F	L	P
82001	0.694819	0.689906	0.696341	0.702507
82002	0.696669	0.697879	0.665054	0.676051
82003	0.699352	0.711596	0.698589	0.693431
82004	0.682921	0.689995	0.701035	0.671198
82005	0.679101	0.675012	0.681540	0.645812
82006	0.688257	0.683582	0.689939	0.688774
82007	0.679004	0.686247	0.679320	0.661486
82008	0.699766	0.669267	0.684731	0.680135
82009	0.670182	0.674752	0.688223	0.675300
82010	0.701978	0.686436	0.708491	0.686731

Table 42: All fixed-strength candidate means across tuning seeds 81001–81003. Selection uses development CIDEr-D only; the selected strength is 0.001. All 18 individual scores accompany the numerical release.

Strength	Mean development CIDEr-D
0	0.752072
0.001	0.773986
0.003	0.752452
0.01	0.766636
0.03	0.762990
0.05	0.753626

Table 43: Descriptive raw-source test BLEU-4 and pooled token NLL (ten-seed means). Neither defines an additional success criterion.

Condition	BLEU-4	NLL
Baseline	0.249729	1.994744
Tuned fixed	0.250383	1.994845
MACC-Lite	0.248169	1.999819
Full MACC	0.244954	1.997075

The server-recorded audio verification additionally recomputes caption metrics from saved predictions and original references. These checks establish the arithmetic of the retained records; they do not independently regenerate predictions, adversarial examples, generated images, or captions from raw inputs and checkpoints.

For the control study, an independent NumPy/SciPy reconstruction checks all 104 final holdout arrays and 80 final official-validation arrays, including labels, logits, top-1/top-5 accuracy, cross-entropy, ECE, and class-scatter sufficient statistics. It also checks 543 training traces containing 4,155,880 recorded optimizer steps, 520 acquisition selections, and 40,872 control windows. Source and payload digests, split reconstruction, target calibration, learning-rate schedules, regularizer arithmetic, moving residuals, and applied coefficients are consistent with the recorded protocol.

This reconstruction does not rerun checkpoint inference, recover individual feature vectors, or independently recompute every earlier-round prediction

and acquisition-candidate score. Earlier-round learning curves use hash-verified summaries. The server exporter records verification of 543 checkpoint hashes; the checkpoint bytes were not included in the independent reconstruction. This distinction separates verification of saved evidence from a new execution of the experiment.

For the 206-run study, the complete compact review was retrieved only after worker completion and checked against the authoritative archive checksum, checksum file, ZIP integrity, and payload digests. A compact export supplies all 206 endpoint records, the planned 17-contrast analysis, tuning rankings, inventory audit, and protocol. An independent NumPy/SciPy reconstruction reproduces every effect, paired SD, confidence interval, raw and Holm-adjusted paired-t and exhaustive sign-flip value, and all candidate selection scores. It also checks unique method counts, shared initial-state groups, annotations, and optimizer steps. This is an independent arithmetic audit of the exported records, not local regeneration of raw predictions or an independent repeat of training, acquisition, or image auditing.

The server separately records re-execution of the earlier 104 final checkpoints (104 feedback sets and 80 official-validation sets), using the original archived inference implementation and independent NumPy metric arithmetic. It also records reload checks for all 206 new final models. These two scopes are distinct. The compact export retains every new reload flag and one detailed replay example, rather than all checkpoint binaries or raw inference arrays. Accordingly, neither server replay nor local statistical reconstruction is represented as a third-party implementation replication. The completed diagnostic review was subsequently retrieved after worker exit and verified against its authoritative archive hash and every one of its 7,455 payload digests. All three primary 2/255 contrasts were independently reconstructed from the 150 suite summaries, including exact sign enumeration and individual confidence intervals. This is arithmetic verification, not independent generation of adversarial inputs. Server clean-logit replay records exact agreement for all 30 checkpoints.

The raw-source audio review likewise passed archive, checksum, ZIP and all 920 payload checks. Its clean child-process audits strictly reload final models and reproduce saved development-probe logits and NLL bitwise; saved-caption metrics and original references are checked. The audits do not regenerate all test captions or establish independent external reproduction. The released primary CIDEr-D uses the retained recomputed pycocoevalcap/PTB scores; native-versus-recomputed discrepancies are preserved, with a maximum absolute CIDEr-D discrepancy below 0.001. Every seed remains in the analysis. Hashes identify retained evidence; they do not prove scientific validity by themselves.

F.3 Audio model initialization

In the earlier 24-run audio study, initialization required a bounded retry after a pre-training mismatch in the encoder positional parameter. A construction was accepted only when its whole-model digest matched the original seed-specific record. No training outcome or evaluation score entered that acceptance rule. The completed outputs, seeds, and acceptance records were preserved.

The acceptance procedure enforces a match to the recorded model state, but does not resolve the underlying initialization nondeterminism. Accepted-state hashes alone cannot reconstruct their tensors, and the availability of all corresponding initial-state binaries is not established. Exact regeneration from seeds alone is therefore not claimed. This limitation concerns reconstruction of the initial state; it does not change the paired comparisons computed from the retained completed runs. The separate raw-source follow-up saves one complete CPU initial state per seed and enforces its full-state tensor hash across methods. Those saved binaries, rather than seeds alone, define pairing; the cause of the historical mismatch is not assigned to an uninitialized parameter. Raw-source checkpoints are retained outside the compact review.

F.4 Released materials and access boundaries

The repository identifies dataset revisions, source-file digests, experiment configurations, and exported per-seed measurements, together with the numerical reconstruction workflows. Raw licensed images, corpus text, audio recordings, and large model checkpoints are not redistributed. Original data-access conditions and the audio subset's coverage and provenance limitations continue to apply. The release supports inspection and statistical reconstruction within the evidence boundaries described above; independent training or inference also requires the relevant data and model assets.

References

[1] B. Settles, *Active Learning Literature Survey*, Computer Sciences Technical Report 1648, University of Wisconsin–Madison, 2009; updated 26 January 2010. https://burrsettles.com/pub/settles.activelearning.pdf

[2] O. Sener, S. Savarese, Active learning for convolutional neural networks: A core-set approach, in: *International Conference on Learning Representations*, 2018. https://arxiv.org/abs/1708.00489

[3] J. T. Ash, C. Zhang, A. Krishnamurthy, J. Langford, A. Agarwal, Deep batch active learning by diverse, uncertain gradient lower bounds, in: *International Conference on Learning Representations*, 2020. https://openreview.net/forum?id=ryghZJBKPS

[4] V. Papyan, X. Y. Han, D. L. Donoho, Prevalence of neural collapse during the terminal phase of deep learning training, *Proceedings of the National Academy*

of Sciences 117(40) (2020) 24652–24663. doi:10.1073/pnas.2015509117

[5] Z. Chen, V. Badrinarayanan, C.-Y. Lee, A. Rabinovich, GradNorm: Gradient normalization for adaptive loss balancing in deep multitask networks, in: *Proceedings of the 35th International Conference on Machine Learning*, PMLR 80, 2018, pp. 794–803. https://proceedings.mlr.press/v80/chen18a.html

[6] Z. Zhu, T. Ding, J. Zhou, X. Li, C. You, J. Sulam, Q. Qu, A geometric analysis of neural collapse with unconstrained features, in: *Advances in Neural Information Processing Systems* 34, 2021, pp. 29820–29834. https://arxiv.org/abs/2105.02375

[7] P. Súkeník, C. Lampert, M. Mondelli, Neural collapse vs. low-rank bias: Is deep neural collapse really optimal?, in: *Advances in Neural Information Processing Systems* 37, 2024, pp. 138250–138288. doi:10.52202/079017-4388

[8] P. Súkeník, C. H. Lampert, M. Mondelli, Neural collapse is globally optimal in deep regularized ResNets and transformers, in: *Advances in Neural Information Processing Systems* 38, 2025, pp. 43573–43604. doi:10.52202/085713-1450 NeurIPS 2025 proceedings.

[9] W. Hong, S. Ling, Beyond unconstrained features: Neural collapse for shallow neural networks with general data, *Journal of Machine Learning Research* 27(82) (2026) 1–46. https://jmlr.org/papers/v27/24-1429.html

[10] J. Wang, Z. Gao, Z. Xie, Space alignment matters: The missing piece for inducing neural collapse in long-tailed learning, *Proceedings of the AAAI Conference on Artificial Intelligence* 40(31) (2026) 26302–26309. doi:10.1609/aaai.v40i31.39835

[11] J. Alcala, R. Andreeva, V. A. Kobzar, D. G. Mixon, S. Na, S. Sule, Y. Xie, Neural collapse in the orthoplex regime, arXiv preprint arXiv:2603.20587, 2026. https://arxiv.org/abs/2603.20587

[12] R. A. Fisher, The use of multiple measurements in taxonomic problems, *Annals of Eugenics* 7(2) (1936) 179–188. doi:10.1111/j.1469-1809.1936.tb02137.x

[13] M. Dorfer, R. Kelz, G. Widmer, Deep linear discriminant analysis, in: *International Conference on Learning Representations*, 2016. https://arxiv.org/abs/1511.04707

[14] Y. Wen, K. Zhang, Z. Li, Y. Qiao, A discriminative feature learning approach for deep face recognition, in: *European Conference on Computer Vision*, LNCS 9911, 2016, pp. 499–515. doi:10.1007/978-3-319-46478-7_31

[15] P. Khosla, P. Teterwak, C. Wang, A. Sarna, Y. Tian, P. Isola, A. Maschinot, C. Liu, D. Krishnan, Supervised contrastive learning, in: *Advances in Neural Information Processing Systems* 33, 2020, pp. 18661–18673. NeurIPS 2020 proceedings.

[16] A. Bardes, J. Ponce, Y. LeCun, VICReg: Variance-invariance-covariance regularization for self-supervised learning, in: *International Conference on Learning Representations*, 2022. https://arxiv.org/abs/2105.04906

[17] Y. Gal, R. Islam, Z. Ghahramani, Deep Bayesian active learning with image data, in: *Proceedings of the 34th International Conference on Machine Learning*, PMLR 70, 2017, pp. 1183–1192. https://proceedings.mlr.press/v70/gal17a.html

[18] X. Li, P. Yang, Y. Gu, X. Zhan, T. Wang, M. Xu, C. Xu, Deep active learning with noise stability, *Proceedings of the AAAI Conference on Artificial Intelligence* 38(12) (2024) 13655–13663. doi:10.1609/aaai.v38i12.29270

[19] W. Bae, D. J. Sutherland, G. L. Oliveira, Uncertainty herding: One active learning method for all label budgets, in: *International Conference on Learning Representations*, 2025, pp. 788–805. ICLR 2025 proceedings.

[20] C. Jung, S. Salehi, A. Schmeink, Active learning using aggregated acquisition functions: Accuracy and sustainability analysis, *IEEE Journal of Selected Topics in Signal Processing*, 2026, pp. 1–15. doi:10.1109/JSTSP.2026.3673105 Author version: https://arxiv.org/abs/2602.07440

[21] G. Cui, Y. Liu, P. Wang, H. Hsu, X. Sun, X. Hao, Z. Liu, Labeled TrustSet guided: Batch active learning with reinforcement learning, in: *International Joint Conference on Neural Networks*, 2026, author manuscript. https://arxiv.org/abs/2604.12303

[22] A. Kendall, Y. Gal, R. Cipolla, Multi-task learning using uncertainty to weigh losses for scene geometry and semantics, in: *IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 7482–7491. doi:10.1109/CVPR.2018.00781

[23] R. J. Williams, Simple statistical gradient-following algorithms for connectionist reinforcement learning, *Machine Learning* 8(3–4) (1992) 229–256. doi:10.1007/BF00992696

[24] J. Su, Y. S. Zhang, N. Tsilivis, J. Kempe, On the robustness of neural collapse and the neural collapse of robustness, *Transactions on Machine Learning Research*, 2024. https://arxiv.org/abs/2311.07444

[25] F. Croce, M. Hein, Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks, in: *Proceedings of the 37th International Conference on Machine Learning*, PMLR 119, 2020, pp. 2206–2216. https://proceedings.mlr.press/v119/croce20b.html

[26] C. Guo, G. Pleiss, Y. Sun, K. Q. Weinberger, On calibration of modern neural networks, in: *Proceedings of the 34th International Conference on Machine Learning*, PMLR 70, 2017, pp. 1321–1330. https://proceedings.mlr.press/v70/guo17a.html

[27] Y. Tian, D. Krishnan, P. Isola, Contrastive multiview coding, in: *European Conference on Computer Vision*, LNCS 12356, 2020, pp. 776–794. doi:10.1007/978-3-030-58621-8_45

[28] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, L. Fei-Fei, ImageNet: A large-scale hierarchical image database, in: *IEEE Conference on Computer Vision and Pattern Recognition*, 2009, pp. 248–255. doi:10.1109/CVPR.2009.5206848

[29] K. He, X. Zhang, S. Ren, J. Sun, Deep residual learning for image recognition, in: *IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770–778. doi:10.1109/CVPR.2016.90

[30] S. Holm, A simple sequentially rejective multiple test procedure, *Scandinavian Journal of Statistics* 6(2) (1979) 65–70. https://www.jstor.org/stable/4615733

[31] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, I. Sutskever, *Language Models are Unsupervised Multitask Learners*, OpenAI technical report, 2019. Technical report.

[32] C. Raffel, N. Shazeer, A. Roberts, K. Lee, S. Narang, M. Matena, Y. Zhou, W. Li, P. J. Liu, Exploring the limits of transfer learning with a unified text-to-text transformer, *Journal of Machine Learning Research* 21(140) (2020) 1–67. https://jmlr.org/papers/v21/20-074.html

[33] J. Ho, A. Jain, P. Abbeel, Denoising diffusion probabilistic models, in: *Advances in Neural Information Processing Systems* 33, 2020, pp. 6840–6851. https://arxiv.org/abs/2006.11239

[34] A. Krizhevsky, *Learning Multiple Layers of Features from Tiny Images*, technical report, University of Toronto, 2009. https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf

[35] J. Song, C. Meng, S. Ermon, Denoising diffusion implicit models, in: *International Conference on Learning Representations*, 2021. https://arxiv.org/abs/2010.02502

[36] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, S. Hochreiter, GANs trained by a two time-scale update rule converge to a local Nash equilibrium, in: *Advances in Neural Information Processing Systems* 30, 2017, pp. 6626–6637. https://arxiv.org/abs/1706.08500

[37] M. Bińkowski, D. J. Sutherland, M. Arbel, A. Gretton, Demystifying MMD GANs, in: *International Conference on Learning Representations*, 2018. https://arxiv.org/abs/1801.01401

[38] T. Salimans, I. Goodfellow, W. Zaremba, V. Cheung, A. Radford, X. Chen, Improved techniques for training GANs, in: *Advances in Neural Information Processing Systems* 29, 2016, pp. 2234–2242. https://arxiv.org/abs/1606.03498

[39] C. D. Kim, B. Kim, H. Lee, G. Kim, AudioCaps: Generating captions for audios in the wild, in: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Volume 1 (Long and Short Papers), 2019, pp. 119–132. doi:10.18653/v1/N19-1011

[40] Y. Gong, Y.-A. Chung, J. Glass, AST: Audio Spectrogram Transformer, in: *Interspeech*, 2021, pp. 571–575. doi:10.21437/Interspeech.2021-698 https://www.isca-archive.org/interspeech_2021/gong21b_interspeech.html

[41] J. F. Gemmeke, D. P. W. Ellis, D. Freedman, A. Jansen, W. Lawrence, R. C. Moore, M. Plakal, M. Ritter, Audio Set: An ontology and human-labeled dataset for audio events, in: *IEEE International Conference on Acoustics, Speech and Signal Processing*, 2017, pp. 776–780. doi:10.1109/ICASSP.2017.7952261 Author publication record.

[42] R. Vedantam, C. L. Zitnick, D. Parikh, CIDEr: Consensus-based image description evaluation, in: *IEEE Conference on Computer Vision and Pattern Recognition*, 2015, pp. 4566–4575. doi:10.1109/CVPR.2015.7299087 CVPR open-access paper.

[43] K. Papineni, S. Roukos, T. Ward, W.-J. Zhu, Bleu: A method for automatic evaluation of machine translation, in: *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, 2002, pp. 311–318. doi:10.3115/1073083.1073135

[44] C.-Y. Lin, ROUGE: A package for automatic evaluation of summaries, in: *Text Summarization Branches Out*, 2004, pp. 74–81. https://aclanthology.org/W04-1013/

[45] M. Denkowski, A. Lavie, Meteor Universal: Language specific translation evaluation for any target language, in: *Proceedings of the Ninth Workshop on Statistical Machine Translation*, 2014, pp. 376–380. doi:10.3115/v1/W14-3348

[46] P. Anderson, B. Fernando, M. Johnson, S. Gould, SPICE: Semantic propositional image caption evaluation, in: *European Conference on Computer Vision*, LNCS 9909, 2016, pp. 382–398. doi:10.1007/978-3-319-46454-1_24

[47] S. Liu, Z. Zhu, N. Ye, S. Guadarrama, K. Murphy, Improved image captioning via policy gradient optimization of SPIDEr, in: *IEEE International Conference on Computer Vision*, 2017, pp. 873–881. doi:10.1109/ICCV.2017.100 ICCV open-access paper.

[48] M. Hein, M. Andriushchenko, Formal guarantees on the robustness of a classifier against adversarial manipulation, in: *Advances in Neural Information Processing Systems*, 30, 2017, pp. 2266–2276. Conference publication record.